



Jupiter Horizon Investigation Image Analysis 30330

AUTHORS

Andrey Galkin - s252898
Vojtech De Coninck - s252918

December 19, 2025

Contents

1	Introduction	1
1.1	JUNO Mission	1
1.2	Problem statement	2
1.3	Research goals	2
2	Background and Data	4
2.1	Juno ASC's camera	4
2.2	Dataset description	5
2.3	SPICE toolkit	7
3	Image Processing	9
3.1	Radiometric Calibration	9
3.1.1	Dark frame subtraction	9
3.1.2	Residual sensor artifacts	10
3.2	External Noise Mitigation	12
3.2.1	Radiation noise characterization	12
3.2.2	Gradient-based filtering	13
3.3	Despinning	14
3.3.1	Odd/even field decomposition	15
3.3.2	Rotation alignment and merge	15
3.4	Lens Distortion Correction	18
3.4.1	Mathematical model	18
3.4.2	Distortion magnitude determination	19

3.5	Performance Benchmark	20
3.5.1	Implementation	20
3.5.2	Onboard feasibility analysis	20
3.6	Final Output	22
4	Horizon Detection	24
4.1	Problem statement & Motivation	24
4.2	Gradient-based Edge Tracking	24
4.3	Circle Fitting Method	25
4.3.1	Validity of local circle approximation	25
4.3.2	Algorithm selection and validation	26
4.4	Comparison with SPICE output	28
4.5	Conclusions	29
5	Star Mapping	30
5.1	Star Detection	30
5.1.1	Dual mask construction	31
5.1.2	Centroid extraction	32
5.2	Star Catalogue	33
5.2.1	Catalogue selection	33
5.2.2	Negligibility of Parallax effects	33
5.2.3	Inertial vector conversion	34
5.3	Attitude Estimation	34
5.3.1	Pattern ambiguity	34
5.3.2	Camera ray projection	35
5.3.3	Fast candidate search	36
5.3.4	Robust solver (RANSAC/Wahba)	37
5.3.5	Boresight calculation	39
5.3.6	Mapping results	39
6	Applications	40
6.1	Juno's position vector	40

6.1.1	Constructing the position vector	41
6.1.2	Consistency checks	44
6.1.3	Transforming to Jupiter-centric coordinates	45
6.2	Atmospheric Investigation	46
6.2.1	Spatial resolution	46
6.2.2	Validation of thermosphere detection	47
7	Discussion and Conclusions	50
7.1	Discussion	50
7.1.1	What works best for this Juno dataset	50
7.1.2	Limitations and sources of error	51
7.2	Conclusions	54
7.2.1	Main findings	54
7.2.2	Recommendations for future work	54
A	Additional Material	58
A.1	Star recognition output from whole dataset	58
A.2	Source Code	60
A.2.1	Main script	60
A.2.2	Star tracker script	90

Chapter 1

Introduction

1.1 JUNO Mission

JUNO is a NASA spacecraft dedicated to the investigation of Jupiter’s interior, atmosphere, and magnetospheric environment from a polar, highly elliptical orbit. This trajectory creates short perijove windows with rapidly changing observation geometry, separated by long arcs far from the planet. An overview of the orbit evolution and close-approach passes is shown in Fig. 1.1.

JUNO is spin-stabilized, which simplifies attitude dynamics but introduces image motion for instruments integrating during spacecraft rotation. The original mission baseline envisioned two initial 53-day orbits followed by a period reduction to 14 days. However, due to concerns in the main-engine feed system (check-valve behavior), this maneuver was not executed and the mission remained in the longer-period orbit, reducing operational risk while preserving the perijove science geometry [1].

Following completion of the prime objectives, NASA approved an extended mission in which the orbit continues to evolve and enables additional science, including targeted encounters in the Jovian system [2, 3]. At end-of-mission, the spacecraft is planned to be disposed by a controlled entry into Jupiter’s atmosphere to satisfy planetary protection requirements [2].

The spacecraft carries four DTU Advanced Stellar Compass (ASC) star trackers, primarily providing precise attitude information to the magnetometer system [4]. During

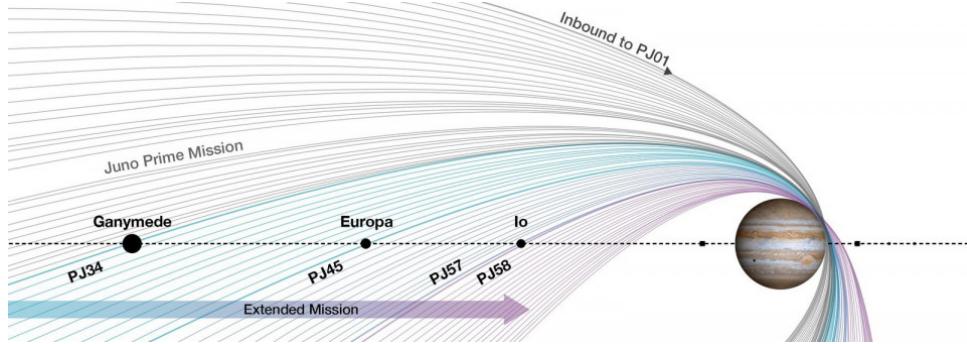


Figure 1.1: Overview of JUNO’s orbit evolution and perijove passes around Jupiter (prime and extended mission). [2]

selected perijove intervals, these cameras are commanded to acquire imagery containing Jupiter’s limb. Although the ASC is not designed as a science imager, these horizon views provide insights about Jupiter’s atmosphere.

1.2 Problem statement

The ASC does not output a fully synchronous frame. Due to the staggered readout scheme (Section 2.1), each acquisition contains two interleaved fields captured at different times, which introduces intra-frame misalignment during spacecraft spin. In addition, the orbit environment and sensor characteristics generate strong outliers and background structure that bias direct feature extraction.

The dataset therefore presents two coupled challenges. First, the raw frames must be converted into a calibrated consistent representation. Second, the signals of interest (a short limb arc and a sparse star field) must be extracted robustly. The required corrections and their implementation are treated in Chapter 3.

1.3 Research goals

The assignment specification was intentionally open-ended. The work therefore focuses on the aspects of the dataset that are most relevant for geometric reconstruction and atmospheric interpretation. The objectives are:

- Construct an image processing sequence that converts raw ASC acquisitions into

calibrated, denoised, and geometrically corrected frames (Chapter 3).

- Compensate the odd/even readout misalignment by despinning and merging both fields into a single consistent image (Section 3.3).
- Extract Jupiter’s apparent horizon from the corrected frames and validate the result against a SPICE-derived reference (Chapter 4).
- Detect and match background stars to an inertial catalogue to estimate camera pointing from the same frames (Chapter 5).
- Combine limb geometry and star-derived attitude to recover an image-constrained spacecraft position estimate, and derive basic horizon-referenced atmospheric observables where supported by the data (Chapter 6).

Chapter 2

Background and Data

2.1 Juno ASC's camera

The Advanced Stellar Compass (ASC) is an instrument providing attitude measurements to Juno's magnetometers. It consists of two pairs of Camera Head Units (CHUs) placed at the end of one of Juno's 3 solar arrays, 8 and 10 meters from the spacecraft's body to mitigate electromagnetic disturbance. Each CHU is a star tracker monochrome 8-bit camera with a $19^\circ \times 13.5^\circ$ field of view and a 752x580 pixel CCD sensor. For optimal performance in a high radiation environment, the sensor lowers its integration time by using the built-in electronic shutter. The full integration time is divided evenly between two image fields : the odd-numbered rows are read out in the first field, followed by the even rows in the second field after a fixed delay Δt . Due to Juno's spin during this delay, the result consists of two clearly misaligned row sets, which must be accounted for during image processing [4].

In addition to the odd-even row misalignment, the ASC operates under low-light conditions with short exposure times, making the images inherently photon-limited. As a result, shot noise and readout noise contribute significantly to the background intensity fluctuations. This noise floor noticeably limits the detectability of faint stars.

Furthermore, the sensor is affected by spatial non-uniformities, including pixel-to-pixel gain variations and fixed-pattern noise. These artefacts can introduce bias in local background estimation and star flux measurements. These effects should likewise be considered

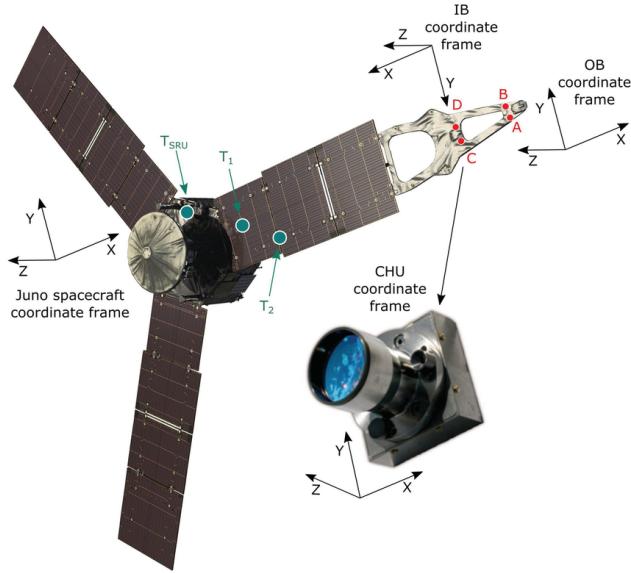


Figure 2.1: CHU disposition on Juno [5]

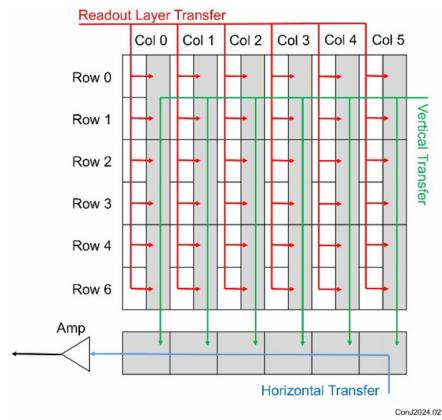


Figure 2.2: Principle of CCD operation. Photosensitive integration layer in white, shuttered layer in gray. Sketch shows integration of first image field [6]

during processing (e.g. using reference images).

Scenes containing bright objects, such as Jupiter or its illuminated limb, can induce intense stray light and scattering within the lens system. This results in extended intensity gradients and halo structures that are not associated with point sources. Such features distort the local background and can suppress or falsely trigger star detections if not explicitly masked.

Finally, residual geometric distortions introduced by the lens lead to systematic deviations from an ideal pinhole camera assumption.

2.2 Dataset description

All of this paper's work and conclusions are derived from 6 images captured by the ASC's Camera Head Unit D. They were all taken during the rare perijove passage (when the orbit has the lowest altitude) over Jupiter's North Pole. During this window, the Jovian horizon is observed at a very high slant angle, providing imaging of the upper atmosphere for a short time, before returning to deep space star field observation.

The dataset consists of 3 pairs of images taken 30 seconds apart. Within a pair of images, they are taken exactly 500ms apart. Due to the spacecraft's high speed and rotation, the observation geometry drastically changes from one pair to the next. All ac-

quisitions were made on September 12 2019 from 3:22:25.714 to 3:23:26.714 UTC. The filename of each acquisition encodes a timestamp, formatted as a floating-point scalar. This value denotes the number of seconds elapsed since the J2000 epoch (January 1, 2000, 12:00:00 UTC), allowing for sub-millisecond synchronization with the spacecraft’s trajectory data. [7]

The first pair is right before Jupiter enters the camera’s field of view. Naturally, they provide little interest in analyzing Juno’s atmosphere.

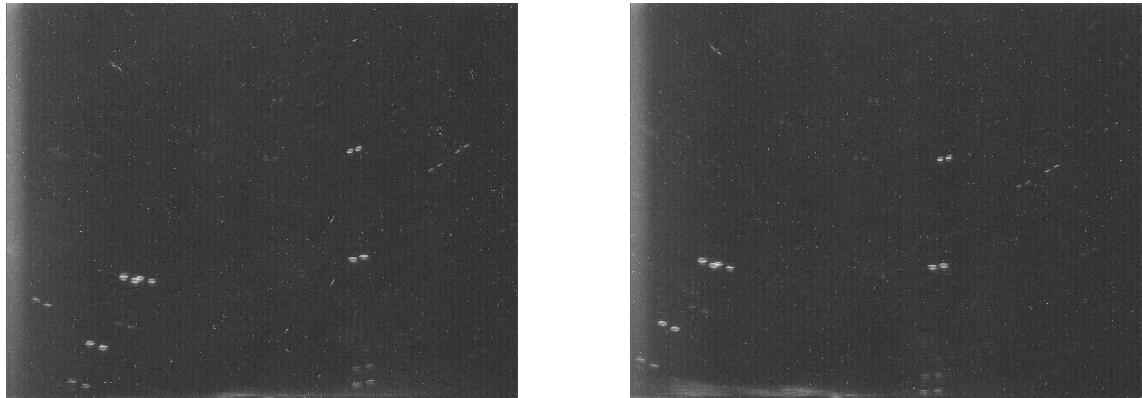


Figure 2.3: Pre-encounter observation (gained x4)

The second pair is right after the planet enters the camera’s field of view. A very short and dark limb of the night side if the planet is observed. The primary features of interest in these frames are the atmospheric intensity drop-off and an aurora visible in the bottom right corner.

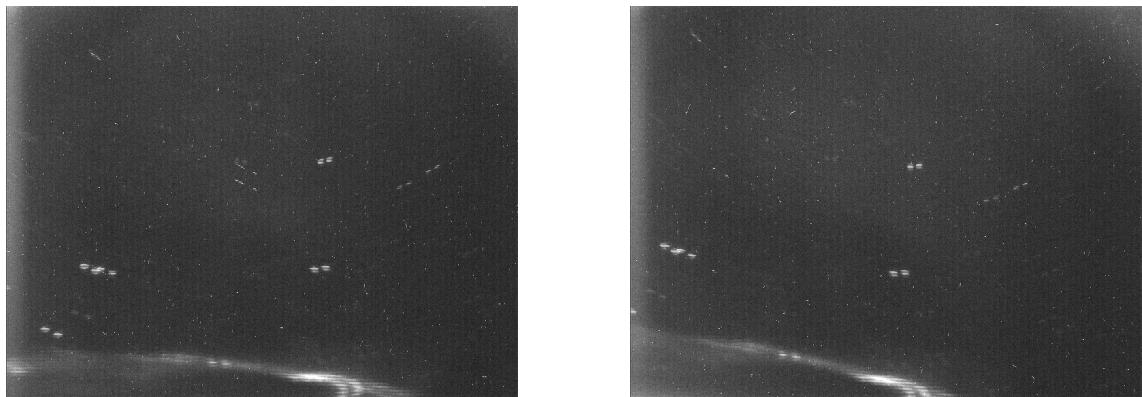


Figure 2.4: Atmospheric encounter (gained x4)

Finally, the third pair presents the greatest scientific interest. Visually, the image

appears to capture a terminator line separating a bright region (day side) from a dimmer one (night side). However, this feature is not the true day/night terminator, but rather an effect of atmospheric scattering of sunlight. The terminator line is out of bounds, positioned below the image frame.

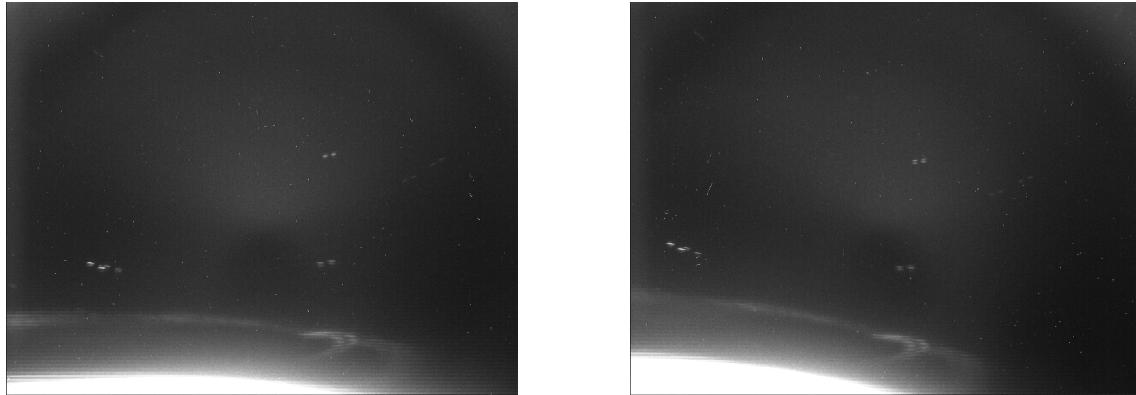


Figure 2.5: Atmospheric scattering

2.3 SPICE toolkit

To accurately interpret the 2D imagery acquired by the ASC, it is necessary to reconstruct the precise 3D observation geometry at the exact moment of capture. This was achieved using the SPICE toolkit, a library provided by NASA's Navigation and Ancillary Information Facility. SPICE stands for Spacecraft, Planet, Instrument, C-matrix and Events. It is used to handle planetary mission geometry, time conversions, and reference frame transformations. It relies on data files known as "kernels," which model the Juno mission environment:

- SPK (Spacecraft Planet Kernel): Provides the trajectory of the Juno spacecraft and the position of the Sun and Jupiter
- CK (C-matrix Kernel): Provides the attitude of the spacecraft, allowing to determine the pointing vector of the camera
- PCK (Planetary Constants Kernel): Provides the physical properties of Jupiter

In this study, CSPICE was used for these critical tasks:

- Surface intercept and horizon determination : By combining the spacecraft's position vector with the camera's pointing vector, each pixel's line-of-sight was projected onto Jupiter. This maps the 2D image plane to the 3D planetary surface, allowing to determine the "true" boundary between the planet and space on the image. This gives a comparison metric for the horizon detection algorithm.
- Distance and scale validation: The toolkit computes the exact slant range, i.e the distance between the spacecraft and the horizon line, which is about 60,000 km for the dataset. This measurement, alongside camera parameters, serves as a baseline to derive spatial resolution, required to convert pixels into physical kilometers.

As handling these tools is complex enough and not the primary focus of this study, LLMs were used for code generation in MATLAB to retrieve geometric data. These measurements were double-checked for plausibility and consistency given the mission's constraints.

Chapter 3

Image Processing

3.1 Radiometric Calibration

The first image processing step is to isolate the signal from sensor-intrinsic artifacts. This process addresses two categories of noise: stochastic thermal noise (dark current) and systematic readout structures (fixed pattern noise).

3.1.1 Dark frame subtraction

The primary calibration step involves the subtraction of a dark reference frame, generated by averaging multiple exposures taken with the shutter closed at the same operating temperature and integration time as the science data. This process serves a dual purpose:

1. Thermal correction: it removes the stochastic dark current noise. This is a thermal phenomenon where silicon atoms spontaneously release electrons, mimicking light signals. They accumulate during exposure and appear as signal in the image, making dark current noise extremely temperature and exposure dependent. Since it is purely additive, the reference dark frame can be subtracted from the raw image.
2. Fixed pattern noise removal: The ASC sensor exhibits a vertical oscillatory pattern caused by gain variations in the column readout amplifiers. Since this pattern is systematic and stable over time, it is captured in the dark frame.

By subtracting the reference frame from the raw image, both the thermal baseline and

the vertical stripes are effectively eliminated in a single operation:

$$I_{corr}(x, y) = I_{raw}(x, y) - I_{dark}(x, y)$$

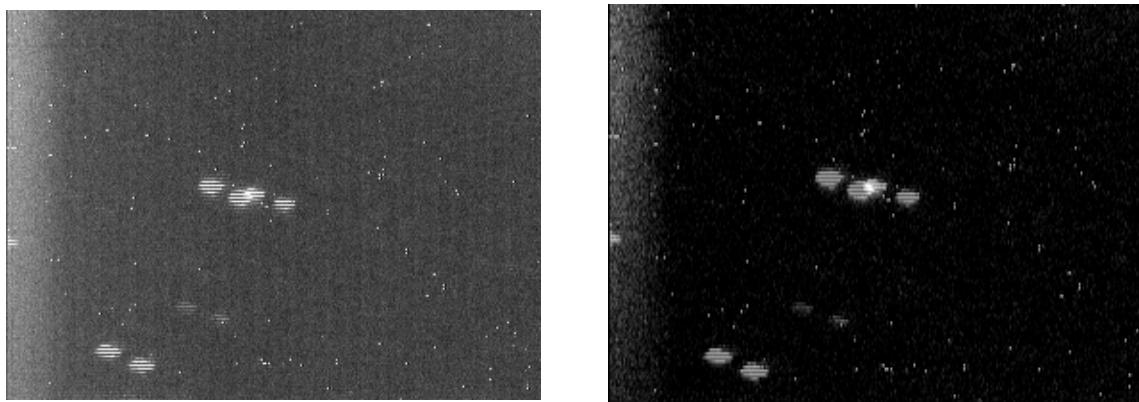


Figure 3.1: Before and after dark frame subtraction (images gained x10). The background is much darker and vertical stripes are removed.

Alternatively, the oscillatory pattern and the DC component could be filtered in the frequency domain. Since the readout noise is a vertical periodic oscillation, its spectrum is concentrated on the horizontal axis of the 2D Fourier transform (see 3.2). Notching the peak at frequency 0 would reduce overall brightness. Notching the secondary peaks would remove the periodic noise. However, this approach is only of theoretical interest and suboptimal compared to dark frame subtraction. Indeed, notching introduces ringing artifacts near high-frequency features such as stars, degrading the signal. In addition, FFT is computationally far heavier than simple integer subtraction.

3.1.2 Residual sensor artifacts

Following the dark subtraction, residual systematic artifacts persist in the calibrated image. The most prominent feature is a localized, additive brightness gradient visible on the left edge of the sensor. Physically, this phenomenon (known as amplifier glow) is a consequence of the sensor's architecture. The power supply and readout electronics are located along the left side of the silicon die. During the readout process, these components dissipate power, leading to localized heating. This power dissipation creates a thermal gradient across the silicon substrate. Since dark current generation is exponentially dependent

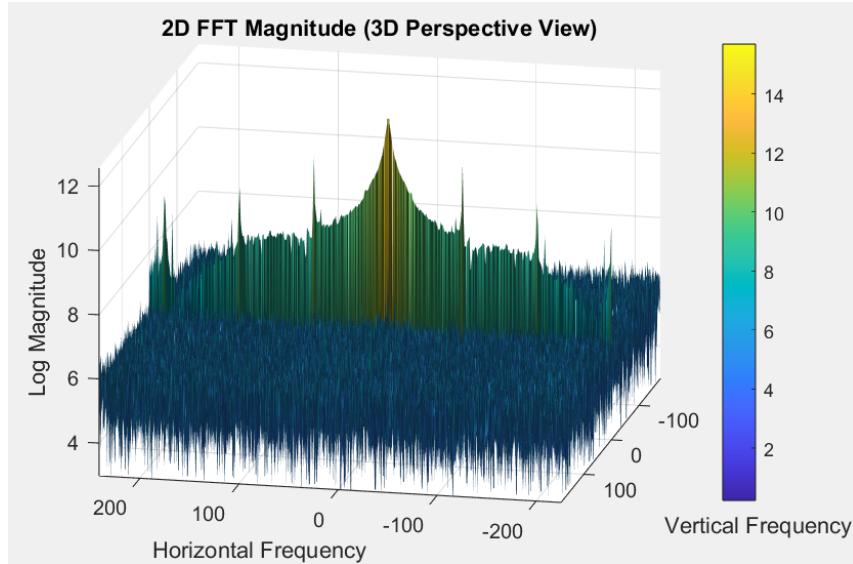


Figure 3.2: 2D Fourier transform of dark frame

dent on temperature, pixels in close proximity to the warm power supply rails accumulate significantly more thermal charge than those in the cooler center of the array.

To eliminate these column-correlated artifact, a column-median subtraction filter is employed. This method relies on the fact that the majority of field of view captures deep space. Consequently, the statistical median of any column provides an estimate of the background bias level for that column, independent of bright features such as stars or the planetary disk.

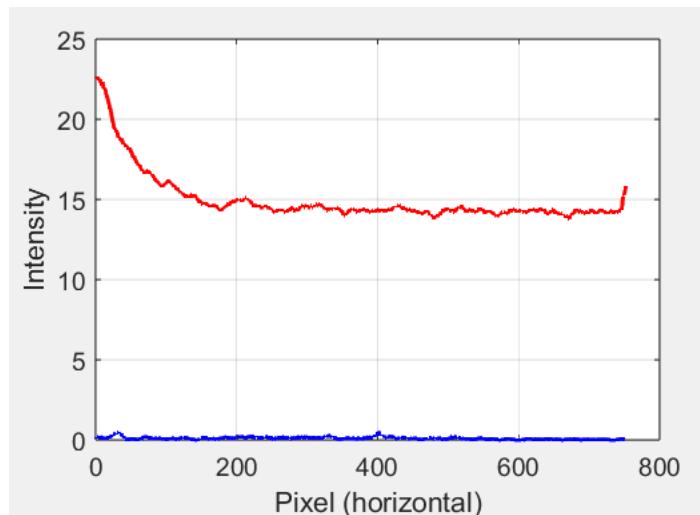


Figure 3.3: Line scan of raw (in red) and calibrated (in blue) background intensities (moving average). The calibration successfully removes amplifier glow (on the left-hand side) and background bias.

3.2 External Noise Mitigation

3.2.1 Radiation noise characterization

Jupiter's radiative environment The Juno spacecraft operates in a harsh radiation environment, primarily dominated by the Jovian radiation belts. As it approaches the poles, the flux of particles in the magnetosphere increases by several orders of magnitude. The CCD is exposed to ionizing radiation flux, specifically high-energy electrons. For a star tracker, which relies on identifying stable constellations, this creates a false depiction of the sky by adding thousands of transient bright spots, interpreted as stars. Therefore, removing these bright features is not an enhancement step : it is absolutely necessary to retrieve scientific data from the image.

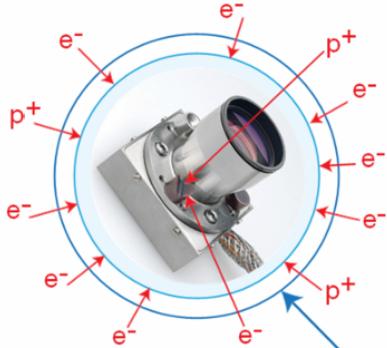


Figure 3.4: Sensor-particle interaction [6]

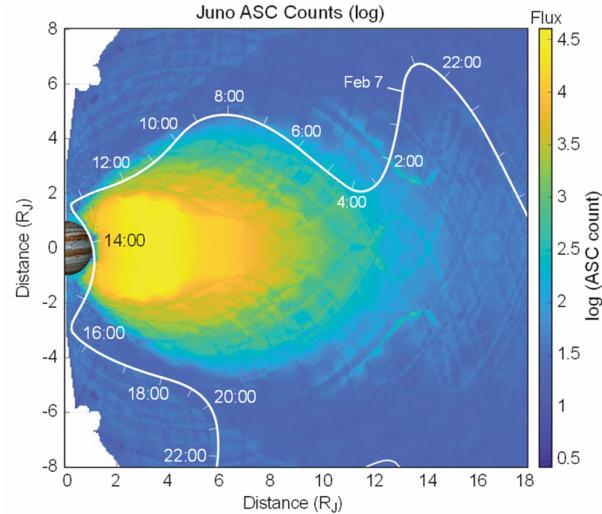


Figure 3.5: Jovian magnetosphere particle flux [6]

Particle-sensor interaction The ASC sensor is equipped with radiation shielding that significantly reduces the flux of particles reaching the CCD. While it efficiently blocks protons and electrons $<10\text{MeV}$, higher energy particles have enough momentum to penetrate the shielding enclosure. The particles that do reach the CCD liberate charge that is collected for a single frame. The noise is therefore transient : it appears in one frame and is completely uncorrelated in the next one.

In the image, an electron appears as a very high intensity isolated pixel because it deposits a large amount of kinetic energy in a single potential well of the CCD. This

"impulse" characteristic makes noise distinguishable from other light sources. The optical system is designed to spread out focal objects (i.e. stars) over a Point Spread Function with an area larger than 4 pixels [6] (see 3.7).

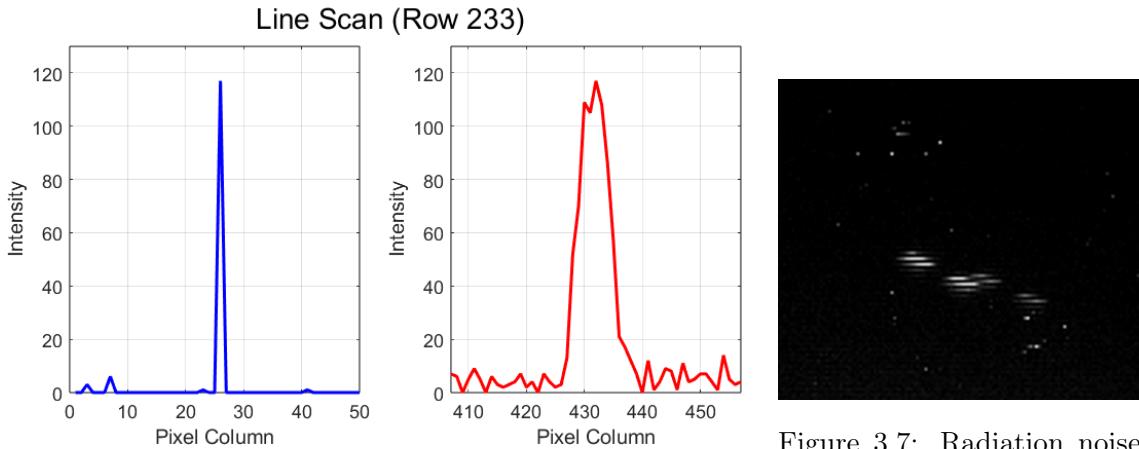


Figure 3.6: Blue line indicates radiation hit, characterized by a pixel-wide impulse (Dirac delta function). Red line represents a star, displaying a broader Gaussian distribution consistent with the optical PSF.

Figure 3.7: Radiation noise (dots) and stars from calibrated image

3.2.2 Gradient-based filtering

The algorithm determines which pixels are valid data and which are radiation speckles. To preserve the edges of the planets and the star field, the chosen topology is a gradient-based switching median filter. Unlike a standard median filter that processes every pixel (potentially blurring fine details), this method is conditional, meaning it only modifies pixels identified as defective based on a gradient threshold. The detection logic relies on the morphological difference between the sharp profile of a radiation speckle and the smoother gradient of a Point Spread Function (cf. 3.6).

Step A : Gradient detection The main challenge is to wisely choose the threshold. If it is set too low, false detections will occur. Stars will be treated as detected particles and removed. If set too high, radiation speckles will be interpreted as stellar candidates. There are 2 cases to separate :

- Light sources : The optics are calibrated to have a finite Point Spread Function. The

signal from a star is normally distributed.

$$I_{star}(r) \propto e^{-r^2/2\sigma^2}$$

- Radiation noise : a particle enters and exit the same pixel sensor (in most cases). This results in a high-intensity impulse signal. The gradient between the speckle and the neighbor is nearly equal to the amplitude of the hit itself/

$$I_{rad}(x, y) = A \cdot \delta(x - x_0, y - y_0)$$

Therefore, by setting a gradient threshold T_{grad} such that:

$$\max(\nabla I_{star}) < T_{grad} < \min(\nabla I_{rad})$$

the filter can effectively separate the two sets.

Step B : Median Replacement Pixels flagged by the gradient check are replaced using a kernel restricted to the nearest orthogonal neighbors (top bottom left right). This method is preferred to a box or Gaussian filter to preserve the edge of Jupiter's limb.

$$P_{new}(x, y) = \text{median}\{P(x, y - 1), P(x, y + 1), P(x - 1, y), P(x + 1, y)\}$$

This approach is supported by literature [8], where Van Dokkum demonstrated that cosmic rays can be robustly distinguished from stars by analyzing the sharpness of their edges (using Laplacian or gradient derivatives) rather than just their intensity.

3.3 Despinning

Due to the ASC staggered readout (cf. 2.1), a single “image” actually contains two interleaved fields captured at different times. During the fixed delay Δt between the odd-and even-row readouts, Juno rotates, causing a systematic duplication and misalignment of all scene features. Despinning compensates this intra-frame motion by mapping the

even field back to the odd-field acquisition time, after which both fields are merged into a single geometrically consistent frame.

3.3.1 Odd/even field decomposition

Let $I(u, v)$ denote the raw 752×580 image, with u the column index and v the row index. The two half-height fields are extracted by row decimation:

$$I_{\text{even}}(u, y) = I(u, 2y + 1) \quad (3.1)$$

$$I_{\text{odd}}(u, y) = I(u, 2y) \quad y \in [0, 289] \quad (3.2)$$

I_{even} and I_{odd} seem to be inverted because the index starts with 0. Line number 1 (odd) corresponds to $I(u, 0)$.

3.3.2 Rotation alignment and merge

For accurate alignment, the motion between both fields is modeled as a 3D rotation of the camera's line-of-sight vectors. The process begins by transforming the spacecraft's angular velocity vector ω_{SC} , which is defined in the spacecraft body frame, into the camera's reference frame (see figure 3.8). This is achieved using the fixed mounting rotation matrix $R_{SC \rightarrow cam}$:

$$\omega_{cam} = R_{SC \rightarrow cam} \omega_{SC} \quad (3.3)$$

From this angular velocity vector, the unit rotation axis \mathbf{k} and the scalar rotation angle θ are defined as :

$$\mathbf{k} = \frac{\omega_{cam}}{\|\omega_{cam}\|}, \quad \theta = \|\omega\|\Delta t \quad (3.4)$$

Using $\omega = 11.967^\circ/\text{s}$ and $\Delta t = 125 \text{ ms}$ yields $\Delta\theta = 1.496^\circ$

Each pixel (u, v) is assigned a 3D vector which represents a pinhole ray using the

intrinsic camera parameters :

$$\mathbf{r} = \begin{bmatrix} (u - c_x) p_x / f \\ (v - c_y) p_y / f \\ 1 \end{bmatrix} \quad (3.5)$$

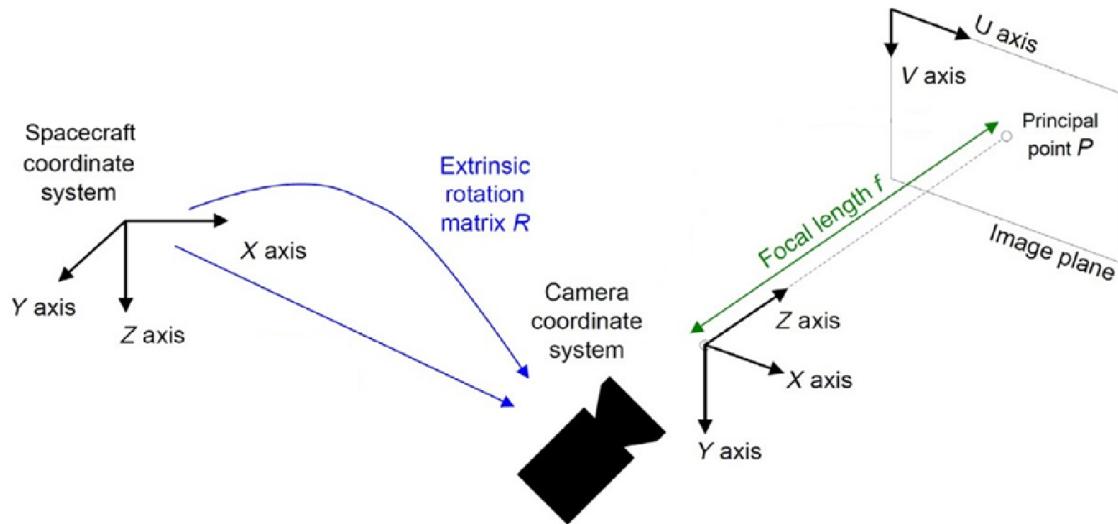


Figure 3.8: Rotation between SC and camera coordinate systems [9].

The rotation from a frame to another is applied using Rodrigues' rotation formula 3.6, allowing to spin a vector around an arbitrary axis \mathbf{k} [10]. For a target pixel in the even field, its corresponding 3D vector \mathbf{r} is rotated around the axis \mathbf{k} by angle θ to align it with the odd-field timestamp (see figure 3.9). The rotated vector \mathbf{r}' is given by:

$$\mathbf{r}' = \mathbf{r} \cos \theta + (\mathbf{k} \times \mathbf{r}) \sin \theta + \mathbf{k}(\mathbf{k} \cdot \mathbf{r})(1 - \cos \theta) \quad (3.6)$$

Finally, this rotated vector \mathbf{r}' is reprojected onto the 2D sensor plane to determine the motion-compensated pixel coordinates (u', v') , allowing the even field to be merged seamlessly with the odd field. Since the even field has half the vertical sampling, the corresponding even-field row index is

$$y' = \frac{v' - 1}{2}$$

The intensity is obtained by bilinear interpolation of $I_{\text{even}}(u', y')$. Pixels whose back-

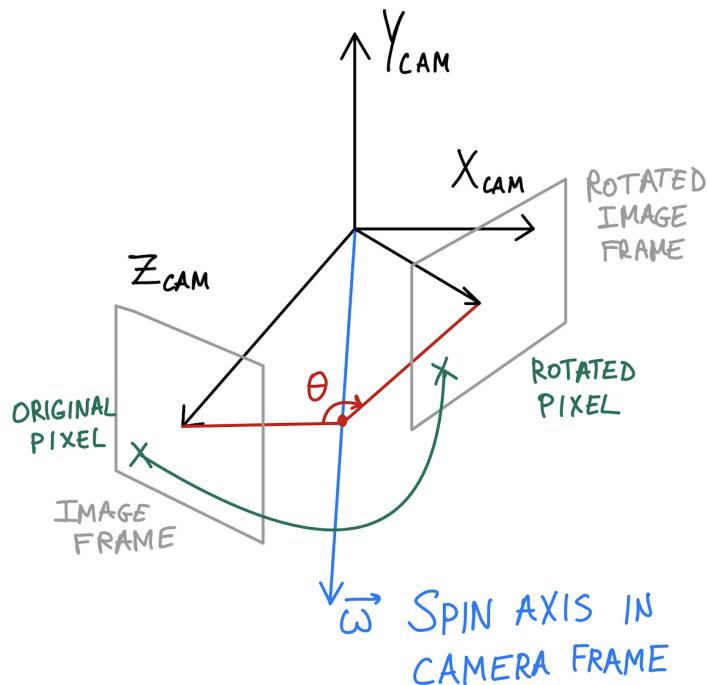


Figure 3.9: Rotation of the camera's pointing vector around the SC's spin axis (expressed in camera frame).

projected coordinates fall outside the even-field bounds are assigned zero, which explains the residual border artifacts observed after despinning.

The two fields are finally merged into a single full-resolution image by interlacing. The odd rows are taken directly from I_{odd} , while the even rows are filled with the motion-compensated samples from $I_{even \rightarrow odd}$ evaluated on the even-row grid. This preserves the original 752×580 sampling. The total despinning result is shown in Fig. 3.10.

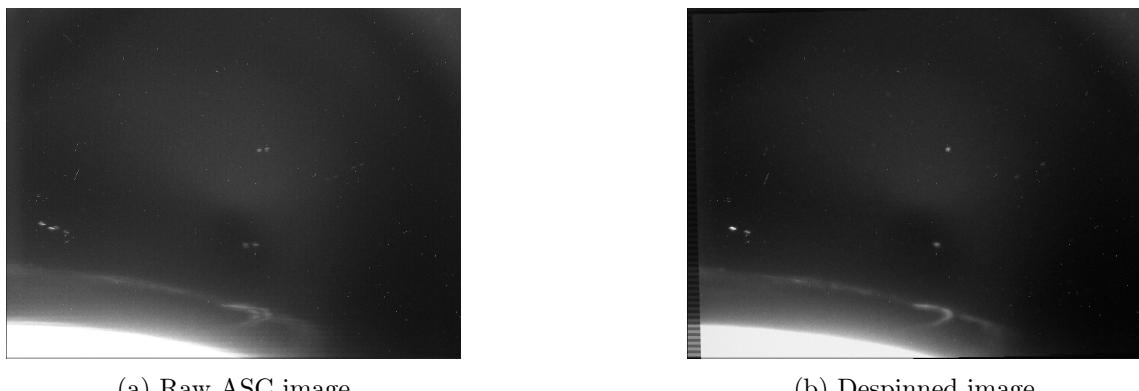


Figure 3.10: Effect of the despinning procedure.

3.4 Lens Distortion Correction

The preceding analysis implicitly relied on the pinhole camera model, which assumes the lens is a single point without geometry. Lens distortion is a deviation from the ideal projection in the pinhole model. It is a form of optical aberration where straight lines appear bent in the CCD due to lens curvature (cf 3.11).

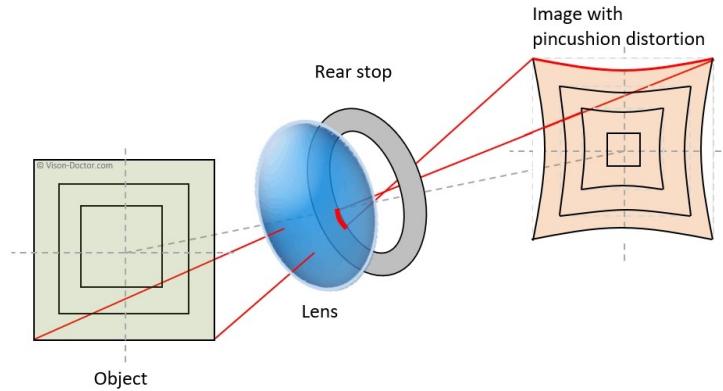


Figure 3.11: Lens distortion origin [11]

3.4.1 Mathematical model

The image can be corrected using a simplified Brown-Conrady transformation, considering only the first-order radial distortion term. To handle the sensor's non-square pixel geometry ($dx = 8.6\mu m$, $dy = 8.3\mu m$), a pixel aspect ratio $\alpha = dy/dx$ is introduced. The transformation relates a point in the undistorted image space to the distorted image space:

$$\begin{cases} x = u_n - u_c \\ y = \alpha(v_n - v_c) \\ u_d = u_c + x(1 + \kappa r^2) \\ v_d = v_c + y(1 + \kappa r^2)/\alpha \end{cases} \quad (3.7)$$

Figure 3.12: Transformation in image spaces [9]

where (u_n, v_n) are the coordinates in the undistorted image, (u_d, v_d) are the coordinates

in the distorted image, $r = \sqrt{x^2 + y^2}$ is the distance from optical center, (u_c, v_c) the coordinates of the optical center, and κ a distortion coefficient.. A positive coefficient indicates a pincushion distortion, i.e the lines bend inwards [9].

It is important to note that these equations 3.7 solve the inverse problem : they solve for the input (distorted image) based on the output (corrected image). An inverse mapping is therefore implemented : instead of iterating over the input image and projecting pixels forward, the algorithm iterates over every coordinate (u_n, v_n) in the target image, computes the corresponding (u_d, v_d) and samples the correct intensity value.

3.4.2 Distortion magnitude determination

To assess the scale of the distortion in the Juno data, a simulation was run with the optical parameters of the ASC's camera:

- $\kappa = 3.3 \cdot 10^{-8}$
- $(u_c, v_c) = (383, 257)$
- $(W, H) = (752, 580)$
- $(DX, DY) = (8.6, 8.3) \mu\text{m}$
- $\alpha = DY/DX = 0.965$

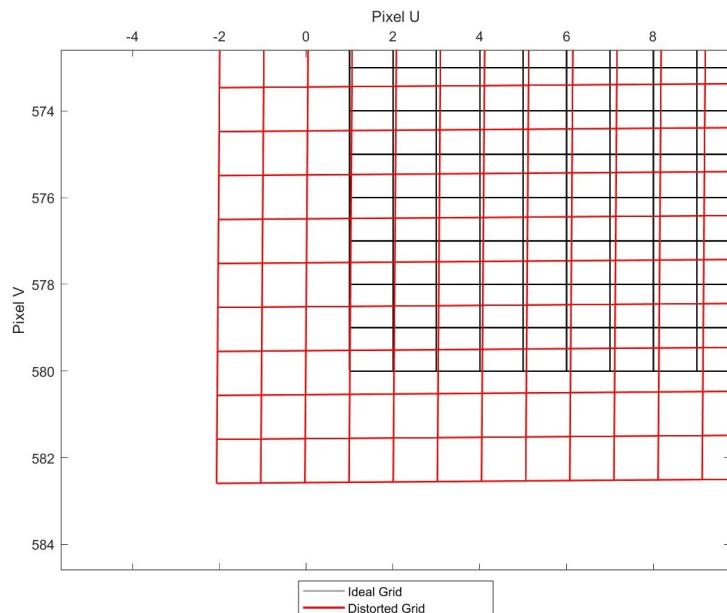


Figure 3.13: Lens distortion simulation (zoomed on corner)

In the simulation 3.13, the distorted grid (red) is projected over the ideal grid of pixels (gray). The distortion is radially dependent, reaching its maximum at the image corners, . In these peripheral regions, the displacement is 3-4 pixels. While visually unnoticeable, it is absolutely non-negligible in the context of atmospheric investigation. Neglecting those few pixels can translate in a geolocation error of hundreds of kilometers.

3.5 Performance Benchmark

This section evaluates the computational resources needed to execute the image processing sequence. The feasibility of running the process on the onboard μ ASC is assessed given the hardware specifications. One crucial detail is that only the star tracking needs to be performed in flight. If a processing step is numerically too intensive to determine attitude, it can be left out if it does not affect the result too much. The atmospheric investigation is ground-based, therefore computational resources are not a limitation.

3.5.1 Implementation

The most computationally demanding stages of the pipeline are the geometric transformations: despinning and lens correction. To assess their viability on onboard hardware, they were implemented in C++. The logic relies on a standard optimization technique that decouples time-invariant map generation from real-time pixel processing.

Pre-computation Since the lens distortion and rotation parameters are constant for the ASC instrument, the coordinate transformation map is time-invariant. Therefore, the computationally heavy fields are computed only once during initialization. Two floating-point matrices (Look-Up Tables) are generated and stored in memory.

Remapping In flight, the correction is applied to incoming raw frames by referencing the pre-computed LUTs. Because the calculated source coordinates rarely align with integer pixel centers, the algorithm employs bilinear interpolation to resample the pixel intensity. This separation allows the heavy trigonometric and polynomial calculations to be front-loaded, leaving only memory lookups and interpolation for the real-time loop.

3.5.2 Onboard feasibility analysis

The analysis in table 3.1 demonstrates that the full image processing sequence is infeasible for the onboard μ ASC microcomputer due to two primary constraints :

- Processing Latency: The ASC operates on a 250 ms cycle (4 Hz) [4]. While the geometric correction takes only 22 ms on a 3.6 GHz laptop, onboard processors (op-

Processing Step	Laptop Benchmark	Onboard Feasibility
<i>Memory (RAM)</i>		
Input Image (8-bit raw)	0.42 MB	Feasible
Output Image (8-bit processed)	0.42 MB	Feasible
Despin Maps ($2 \times$ 32-bit float)	3.33 MB	Inefficient
Lens Correction Maps ($2 \times$ 32-bit float)	3.33 MB	Inefficient
Total RAM Required	~7.5 MB	Borderline
<i>Processing Time</i>		
Initialization & Map Computing	102.4 ms	N/A (One-time)
Dark Frame Calibration	4.6 ms	Feasible
Column-Median Subtraction	72.1 ms	Critical
Speckle Removal	49.3 ms	High Load
Geometric Correction (Despin + Lens)	22.4 ms	Critical (FPU Load)
Total Latency per Frame	~148.4 ms	Infeasible

Table 3.1: Feasibility analysis: processing time and memory requirements for onboard execution (single frame)

erating at <100 MHz with limited Floating-Point Unit capabilities) would execute these floating-point interpolations orders of magnitude slower. The estimated onboard execution time would exceed 1 second per frame, causing a processing timeout that would disrupt the star tracker’s critical attitude determination function.

- RAM Efficiency: Storing the high-precision LUTs for geometric correction requires approximately 7.5 MB of RAM. While technically possible within the hardware limits, allocating so much of the available memory to static maps is highly inefficient and competes with critical buffers required for star catalogs.

In conclusion, the image processing sequence is NOT feasible in its current form on the μ ASC due to strict 250ms time budget. This was expected, as the star-tracking microcontroller was not optimized for image cleaning, but sparse feature extraction (analyzing star shapes) rather than transforming the full 436,000 pixel array. The hardware specifications (RAM, CPU clock speed) are constrained to the bare minimum required for this primary function. The μ ASC should remain focused on lightweight tasks such as simple particle enumeration (thresholding), while the high-load image analysis must be reserved for the ground segment. Further work could analyze the feasibility of skipping certain numerically

heavy steps, such as lens correction and median subtraction, and still keeping acceptable tolerance for attitude determination. That would leave only crucial steps such as bias subtraction and de-rotation in the lighter processing sequence.

sta

3.6 Final Output

The image processing chain has successfully transformed raw, noisy sensor data into calibrated scientific imagery. As summarized in figure 3.14, the process systematically addressed the environmental and hardware-specific challenges of the JUNO mission. The

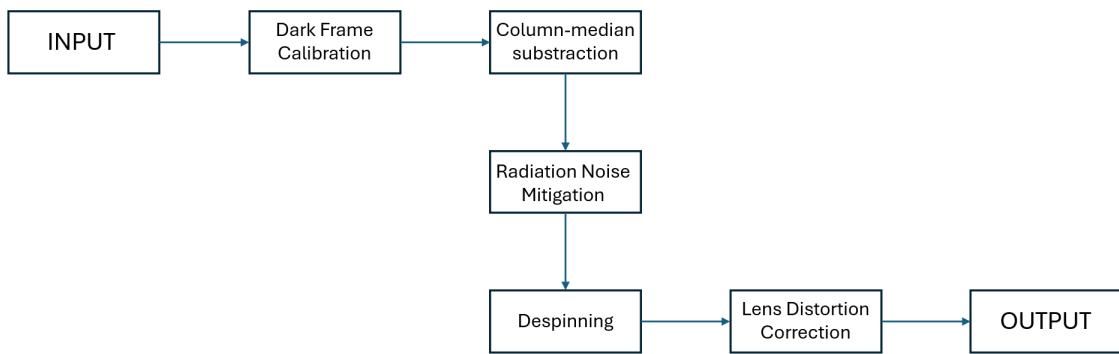


Figure 3.14: Image processing flowchart

final output exhibits an improved SNR and important upgrades :

- Dark calibration has normalized background intensity to nearly zero.
- Gradient-based filtering proved effective in mitigating the high-energy particle speckles due to harsh radiation environment. By selectively targeting high-frequency impulse noise without applying a global smoothing kernel, atmospheric and stellar features are preserved.
- Lens correction and despinning steps have aligned the image features, merging the rotation-induced duplicates into one single object.

While the core features of the atmosphere are now distinct, residual artifacts remain, specifically at the periphery of the frame. The left-hand side of the image exhibits edge-effects from the despinning. These artifacts are attributed to the lack of overlapping data

for interpolation at the image border. Some vignetting is still present in the top right corner.

Despite these minor localized issues, the central region of interest is calibrated and geometrically corrected, rendering the dataset suitable for the intended scientific analysis.

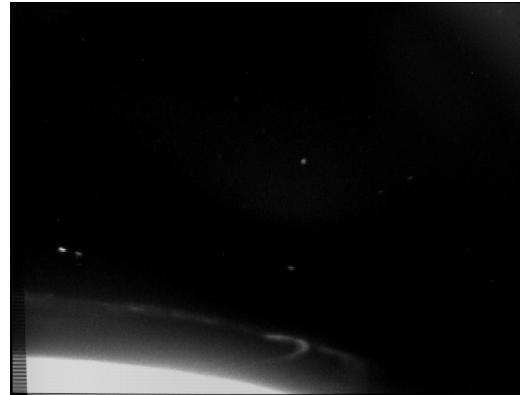


Figure 3.15: Pre & Post-Processing

Chapter 4

Horizon Detection

4.1 Problem statement & Motivation

This chapter investigates the feasibility of extracting Jupiter's horizon purely based on the dataset. Combined with star mapping, the spacecraft's position may be recovered. Accurate detection of the planetary horizon is required for these 2 purposes :

- Optical navigation
- Atmospheric investigation

Fitting a curve to such partial data presents a mathematical challenge. Traditional algebraic methods, such as the standard least-squares fit, are known to be statistically inconsistent when applied to short arcs. They underestimate the curvature radius, leading to substantial errors in estimating the planet's center and the spacecraft's relative position.

The motivation for this study is to implement and validate a robust fitting algorithm capable of overcoming these limitations.

4.2 Gradient-based Edge Tracking

The goal of this section is to extract Jupiter's horizon purely from the dataset, independently of trajectory kernels. In an ideal case, the horizon is defined as the set of points where the intensity gradient is maximized. Standard edge detection methods, such as the Canny Detector, prove unsuitable for the ASC images. Due to the faintness of the limb

(thus low SNR), these methods misidentify high contrast features such as stars, auroras and atmospheric haze as part of the horizon.

To overcome this limitation, a custom algorithm was made just for this purpose. Unlike global edge detectors, this algorithm enforces spatial continuity, tracking the horizon across adjacent columns to reject isolated noise. While generally robust, the method remains sensitive to high-intensity transient features, such as auroras, which can intersect the limb and bias the method. The detection logic is implemented as follows:

- Scanning: Iterate through columns j .
- Gradient Computation: Calculate the vertical gradient vector ∇I_j for the column.
- Thresholding: Calculate an adaptive threshold $T_j = \mu_j + k \cdot \sigma_j$, where μ_j and σ_j are the mean and standard deviation of the gradients in column j .
- Peak Finding: If a gradient $\nabla I_{i,j} > T_j$ is found, verify it is the local maximum within a window $[i - 3, i + 3]$. This step ensures the data is continuous. Update the horizon row $R_j = i$.
- Zero-Order Hold: If no gradient in the window exceeds T_j , the algorithm maintains the previous horizon row: $R_j = R_{j-1}$, assuming the limb has not moved.

This method does not work as is on the last pair of images 2.5. The detected "horizon" is of course the separation between the bright and dark regions, even though in reality it's an effect of sunlight scattering. These regions need to be explicitly masked out for the algorithm to work.

4.3 Circle Fitting Method

The edge tracking yielded a set of vaguely connected points. The goal is now to fit a smooth curve that approximates best this discontinuous line.

4.3.1 Validity of local circle approximation

Jupiter is an oblate spheroid, with a semi-major axis (equatorial radius) and semi-minor axis (polar radius) measuring respectively 71492km and 66854km. These values

correspond to the point of 1 bar of pressure. This gives a flattening ratio of :

$$f = \frac{a - b}{a} = 6.49\%$$

This makes Jupiter the second most flattened planet in the Solar System after Saturn, approximating it as a sphere would yield a non-negligible error.

However, the spacecraft is looking only at an arc of $\sim 16^\circ$ of Jupiter (data derived from SPICE). A simulation was performed to evaluate whether a local spherical approximation is valid (see 4.1). The outcome shows that the curvature of the limb within the narrow FOV is indistinguishable from a spherical arc. The simulation indicated that fitting a circle to the true ellipsoid segment yielded a radial deviation of 12.68 km, which is less than the spatial resolution of a single pixel. An important distinction is that the radius of the fitted circle is not the radius of Jupiter, but rather a local radius of curvature.

In conclusion, fitting a circle rather than an ellipse is valid given the narrow view angle.

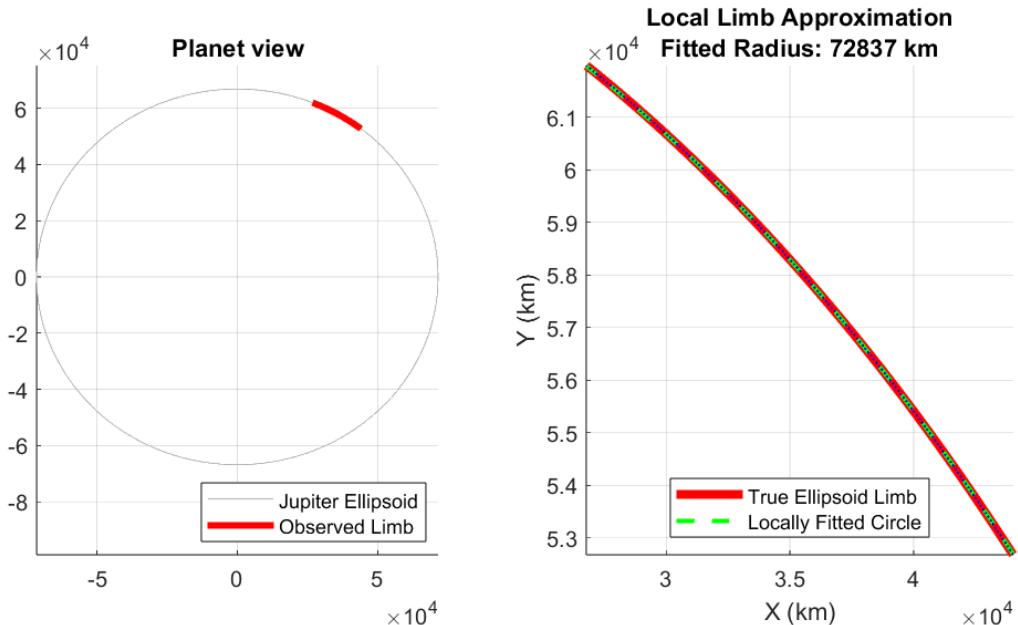


Figure 4.1: Simulation of Jupiter's ellipsoid and locally fit circle

4.3.2 Algorithm selection and validation

Fitting a circle to a set of 2D points is a common problem in image analysis. The most classic approach is performing Kåsa's fit. It is a simple least-square problem where

the squared distance between the estimated circle and the set of points is minimized. It performs well if the set of points is at least half a circle. It tends to be heavily biased towards smaller circles when the set is a short arc, which is precisely the case for Juno's dataset.

For this study's purpose, the most accurate method is a Hyperaccurate (Hyper) algebraic circle fit proposed by Al-Sharadqah and Chernov [12]. This method minimizes the algebraic error for the general circle equation defined by the parameter vector $\mathbf{A} = (A, B, C, D)^T$:

$$A(x^2 + y^2) + Bx + Cy + D = 0$$

The method seeks to minimize the mean square algebraic distances subject to a constraint that prevents the trivial solution $\mathbf{A} = 0$. This is formulated as a generalized eigenvalue problem:

$$M\mathbf{A} = \eta N\mathbf{A}$$

where M is the matrix of data moments, η is the eigenvalue, and N is the constraint matrix, which for centered data is defined as:

$$N = \begin{pmatrix} 8\bar{z} & 0 & 0 & 2 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 0 \end{pmatrix}$$

where \bar{z} is the mean of $z = x^2 + y^2$. By selecting the eigenvector corresponding to the smallest positive eigenvalue, this constraint neutralizes the geometric bias to the order of $O(n^{-2})$, ensuring a stable estimation of the radius even when the visible horizon represents a short limb. Algorithm implemented in C++ with use of AI tools.

Unsurprisingly, the Hyper method outperforms the basic least-square algorithm (see 4.4 for more details). Due to the short visible limb and therefore large circles, the uncertainty on the radius is significant. The radii yielded by Kåsa's and Hyper methods are respectively 1141 and 2248 pixels, which is about 100% difference. This was expected, as the first one is biased towards shorter circles. Still, the efficiency of any circle fitting method is limited

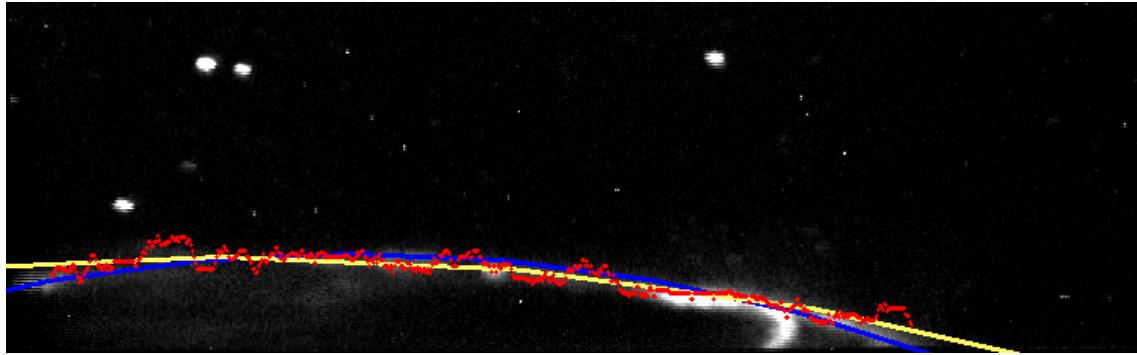


Figure 4.2: Red discontinuous data points show the output of edge-tracking. Blue line is Kåsa’s fit. Yellow line is Hyper fit. Image gained x6.

by the quality of the edge tracking algorithm, and ultimately the data itself.

4.4 Comparison with SPICE output

Only by specifying the exact time of the image capture, the SPICE-based program computes the exact location and field of view of the camera. It can then represent where the 1-bar planet edge is on the image without even analyzing it. This gives a baseline for comparing horizon detection algorithms.

The evaluation metrics for both methods is the error area (total number of mismatched pixels) and the deviation from the radius estimate.

Algorithm	Error Area (px)	Radius (px) (% error)
SPICE (kernel data)	0	2458.16 (+0%)
Kåsa	7937	1141.39 (-54%)
Hyper	7261	2248.52 (-8%)

The Hyper fit outperforms the least-square in both metrics. The Kåsa fit exhibits a massive bias towards smaller circles, underestimating the radius by 54%. Conversely, the Hyper fit maintains the curve shape, recovering the radius with an 8% deviation despite processing only a 16° arc.

However, a persistent error area of approximately 7,000 pixels remains for both methods. This discrepancy is not algorithmic but physical. The edge-tracking detects the limb based on the strongest gradient. This corresponds to high-altitude stratosphere, scattered sunlight, auroral emissions, rather than the 1-bar pressure level defined by the SPICE

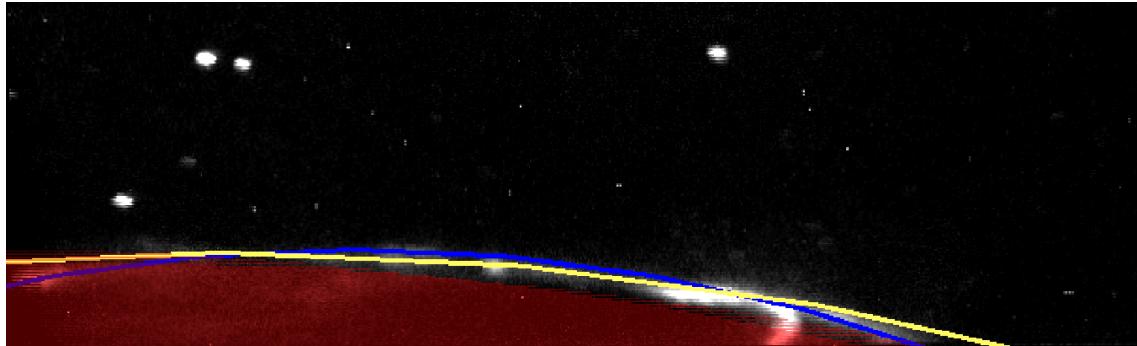


Figure 4.3: Red area represents the 1-bar planet edge extracted from SPICE. Blue line is Kåsa’s fit. Yellow line is Hyper fit. Image gained x6.

model. As a consequence, the visible limb detected by the camera is offset from the physical limb provided by the kernels.

4.5 Conclusions

The investigation of horizon detection on Juno images reveals the discrepancies between physical reality and optical observations. A horizon in a planetary image is not a binary boundary but a continuous gradient of atmospheric scattering, especially when observed from close. While SPICE kernels provide an exact mathematical model of the 1-bar edge, the camera observes the cloud tops and stratospheric hazes. Consequently, any image-based edge detection will inherently deviate from the kernel data. This atmospheric offset is a physical reality of the dataset, not a failure of the fitting algorithm.

For scientific analysis where the spacecraft’s trajectory is well-defined, SPICE kernels should be the primary source for geometry. Image-based horizon detection should be reserved for scenarios where the spacecraft position is the unknown variable (e.g., autonomous navigation). In such cases, the algorithm must combine the Hyper fit with an atmospheric model to account for the offset between the visible atmosphere and the 1-bar edge.

Chapter 5

Star Mapping

The JUNO images contain a limited number of background stars that appear as compact intensity 'blobs', distributed over a small number of pixels. By computing the center-of-mass for each blob, every star can be reduced to a single point vector in space.

Using these vectors, the relative spatial configuration of these stars can be determined. This set can then be compared to the relative angular separations of stars listed in an existing celestial catalogue. Once the correct mapping between the image - catalogue is found, the orientation of the star set with respect to the image plane can be computed.

From this information, the sensor attitude can be determined accurately. Since the sensor is fixed on the spacecraft, this directly yields the spacecraft orientation in inertial space.

5.1 Star Detection

In the ASC images, stars appear with strongly varying intensities. On the one hand, a small number of stars are recorded as highly concentrated, near-saturated intensity blobs. These starts are clearly distinguishable from the background, and can be detected reliably by applying a high absolute intensity threshold. On the other hand, the majority of stars are significantly fainter and closer to the noise floor, which makes them way harder to detect. To attempt to detect both groups of stars, a hybrid approach is implemented in `star_tracker.py`:

5.1.1 Dual mask construction

Let $I(x, y)$ be the grayscale image (after denoising, see Chapter 3). A binary detection mask is formed as

$$M_{\text{final}} = M_{\text{sat}} \vee M_{\text{faint}}, \quad (5.1)$$

where M_{sat} flags saturated/near-saturated pixels and M_{faint} targets small-amplitude point sources.

(1) Saturated blobs. Saturated or near-saturated stars are detected directly from the raw intensity image using a fixed saturation threshold, ensuring robust localisation of the brightest point sources. A saturation threshold T_{sat} is applied directly on I :

$$M_{\text{sat}}(x, y) = 1\{I(x, y) \geq T_{\text{sat}}\}. \quad (5.2)$$

This preserves large bright stars even when adaptive statistics (mean/std) are dominated by Jupiter glow or background gradients.

(2) Faint stars via background subtraction. In parallel, faint stars are detected through background subtraction followed by adaptive thresholding, which enhances small-amplitude point sources while remaining locally robust to background variations. A smooth background estimate $B(x, y)$ is computed with a large median filter kernel (e.g. 61×61 in our default preset):

$$B = \text{MedianFilter}(I, k_{\text{bg}}), \quad I'(x, y) = \max(I(x, y) - B(x, y), 0). \quad (5.3)$$

Optionally, I' is normalized to $[0, 1]$ by its maximum. An adaptive threshold is then computed using the global mean and standard deviation:

$$T = \mu(I') + \alpha \sigma(I'), \quad (5.4)$$

where α is the `threshold_multiplier` (e.g. $\alpha \approx 1.4$ in the default settings). The faint-star mask is

$$M_{\text{faint}}(x, y) = 1\{I'(x, y) \geq T\}. \quad (5.5)$$

Optionally, a morphological opening removes isolated pixels and small artifacts.

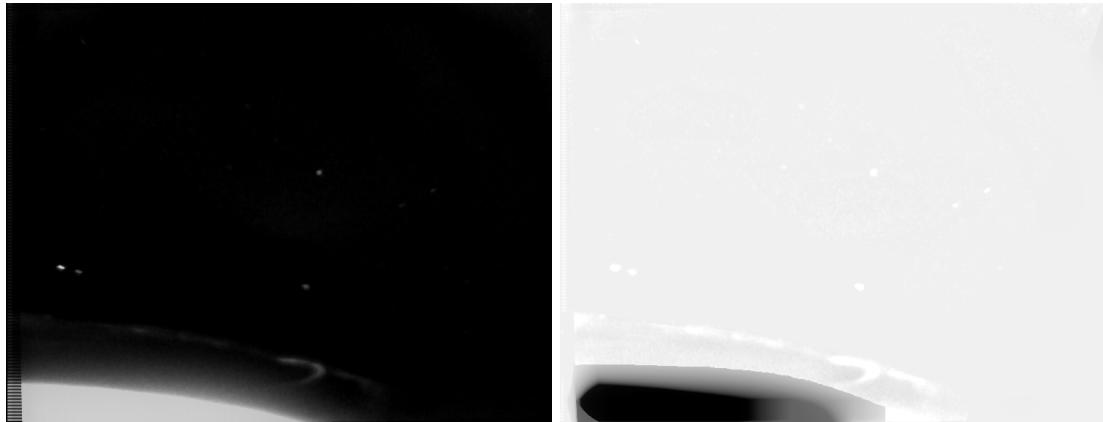


Figure 5.1: (Left) Input frame used for star detection. (Right) Saturated mask

5.1.2 Centroid extraction

Connected components are extracted from M_{final} . For each connected component Ω , we compute:

- its area $A = |\Omega|$ (used to reject oversized artifacts), and
- an intensity-weighted centroid (\hat{x}, \hat{y}) using the original (non-binary) image intensities.

Let $w(x, y) = I(x, y)$ for $(x, y) \in \Omega$. The weighted centroid is:

$$\hat{x} = \frac{\sum_{(x,y) \in \Omega} x w(x, y)}{\sum_{(x,y) \in \Omega} w(x, y)}, \quad \hat{y} = \frac{\sum_{(x,y) \in \Omega} y w(x, y)}{\sum_{(x,y) \in \Omega} w(x, y)}. \quad (5.6)$$

The sum $\sum w(x, y)$ is used as a flux proxy to rank detections; the pipeline keeps only the top- N brightest detections (e.g. $N = 400$), because the subsequent matching is designed around bright catalogue stars.

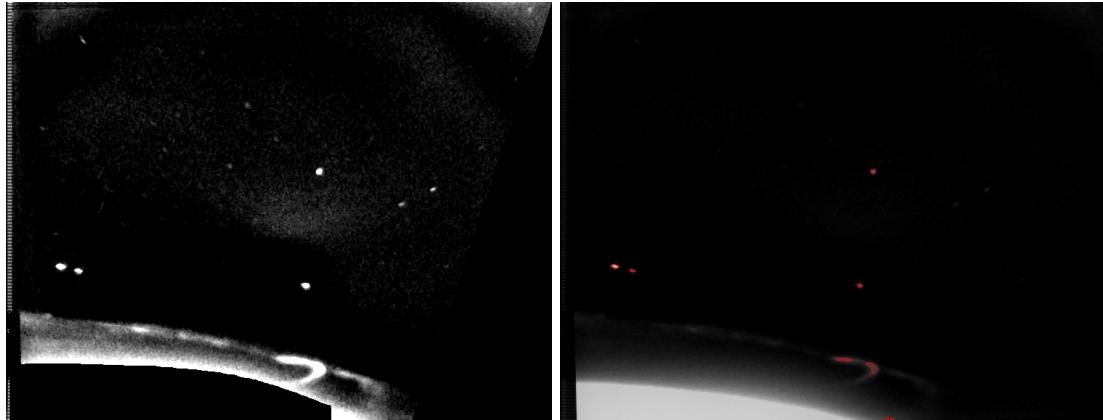


Figure 5.2: (Left) Mask-subtracted image M_{sat} . (Right) Detected star candidates

Figure 5.3: Detected star candidates overlaid on the image. Each blob is reduced to a subpixel centroid via Eq. (5.6).

5.2 Star Catalogue

5.2.1 Catalogue selection

A practical limitation is that not every catalogue is ideal across the full magnitude range of interest:

- **Bright stars:** a HIP/Bright-Star-Catalogue [13] style list is reliable for very bright objects (roughly mag $\in [-2, 4]$).
- **Fainter stars:** The Gaia DR3 catalogue [14] provides dense coverage and accurate astrometry, but a given curated export may miss some of the extremely bright stars or require special handling of identifiers.

Therefore we use a mixed approach: HIP-based entries for the brightest stars, and Gaia entries beyond that range, with a crossmatch/name-lookup layer for consistent labeling (HIP numbers, Gaia IDs, and friendly names/IAU where available).

5.2.2 Negligibility of Parallax effects

Catalogue coordinates are provided as inertial directions from the Solar System barycentric perspective (J2000/ICRF). For star-tracking, we treat stars as fixed points on the celestial sphere. The error from observing from Jupiter instead of Earth is dominated by

stellar parallax:

$$\theta_{\text{parallax}} \approx \frac{B}{D}, \quad (5.7)$$

where B is the observer baseline (Earth–Jupiter distance, order of a few AU) and D is the star distance (typically many parsecs, i.e., light-years). Even for very nearby stars, the resulting angular difference is on the order of arcseconds, which is negligible compared to the matching tolerance used in our robust solver (order of 10^{-1} degrees). Hence, using the same inertial star directions is sufficiently accurate for this study.

5.2.3 Inertial vector conversion

Each catalogue star is represented by its equatorial coordinates (α, δ) in radians (RA, Dec). The corresponding inertial unit vector \mathbf{k} is:

$$\mathbf{k}(\alpha, \delta) = \begin{bmatrix} \cos \delta \cos \alpha \\ \cos \delta \sin \alpha \\ \sin \delta \end{bmatrix}, \quad \|\mathbf{k}\| = 1. \quad (5.8)$$

In the code, these are stored as `vec_inertial` per catalogue entry and are the fundamental primitives used for angular comparisons and attitude estimation.

5.3 Attitude Estimation

5.3.1 Pattern ambiguity

If we only consider relative geometry between stars, the key invariant is the angle between two lines of sight:

$$\theta_{ij} = \arccos(\mathbf{u}_i^\top \mathbf{u}_j), \quad (5.9)$$

where \mathbf{u}_i and \mathbf{u}_j are unit direction vectors.

With three stars, there are three pairwise angles; with four stars, there are six. In a large catalogue (thousands to tens of thousands of stars), a triple of angles may still occur for multiple different triples, especially under measurement noise. Adding a fourth star provides additional constraints (six angles total), and the probability of accidental

collisions becomes extremely small.

However, even if the identity of the stars is known, the attitude is still not obtained from angles alone: the same 4-star shape can appear at different locations on the sensor, which corresponds to different boresight directions. In other words, inter-star angles constrain the pattern but not the absolute pointing. The absolute pixel locations must be mapped to absolute rays in the camera frame, and these rays must be aligned with inertial catalogue vectors.

In our implementation, this disambiguation is achieved by (i) generating hypotheses from angle-consistent pairs, and (ii) validating them against many additional stars (RANSAC inliers). This plays the same role as “adding the 4th star and beyond”: it collapses the remaining ambiguity to a unique attitude.

5.3.2 Camera ray projection

For each detected centroid (\hat{x}, \hat{y}) (in image pixel coordinates), we compute a unit ray \mathbf{v}_{cam} using a pinhole camera model with the ASC intrinsics. Let (f_x, f_y) be focal lengths in pixels and (c_x, c_y) the principal point. Then:

$$x_n = \frac{\hat{x} - c_x}{f_x}, \quad y_n = -\frac{\hat{y} - c_y}{f_y}, \quad (5.10)$$

where the minus sign accounts for the image y -axis pointing downward while the camera frame convention uses $+y$ upward.

A small radial distortion term is optionally applied using a coefficient κ :

$$r^2 = x_n^2 + y_n^2, \quad \begin{bmatrix} x_d \\ y_d \end{bmatrix} = (1 - \kappa r^2) \begin{bmatrix} x_n \\ y_n \end{bmatrix}. \quad (5.11)$$

Finally, the (unnormalized) ray is $\tilde{\mathbf{v}} = [x_d, y_d, 1]^\top$ and the unit ray is

$$\mathbf{v}_{\text{cam}} = \frac{\tilde{\mathbf{v}}}{\|\tilde{\mathbf{v}}\|}. \quad (5.12)$$

Note on resizing. Since figures in this report may be exported/rescaled, the code rescales

(f_x, f_y, c_x, c_y) by the ratio of the current image dimensions to the hardware sensor resolution before applying Eq. (5.10).

5.3.3 Fast candidate search

To efficiently propose catalogue correspondences, angles between pairs of the N brightest catalogue stars are precomputed:

$$\theta_{ab}^{\text{cat}} = \arccos(\mathbf{k}_a^\top \mathbf{k}_b). \quad (5.13)$$

In classical star trackers, Mortari's Pyramid star pattern recognition algorithm [15] introduces the so-called K-vector technique to accelerate this search. All catalogue pair-angles are stored in a single sorted array. A small auxiliary index (the K-vector) then maps a range query

$$\theta_{\text{obs}} \pm \varepsilon$$

to (approximately) the corresponding lower and upper indices in that sorted list. This converts an exhaustive scan over all pairs into a fast range lookup over a short contiguous interval.

A discretized equivalent is used here: the catalogue angles are grouped into bins of width $\Delta\theta$ (e.g. 0.01°), yielding a dictionary

$$\text{bin} \rightarrow \{(a, b, \theta_{ab}^{\text{cat}})\}.$$

Given an observed detection pair-angle, only the corresponding bin (and a small neighborhood of adjacent bins to account for measurement noise) is queried. This reduces the candidate set by several orders of magnitude, while preserving the same core objective as a K-vector range search: limiting hypothesis generation to catalogue pairs with compatible inter-star angles.

5.3.4 Robust solver (RANSAC/Wahba)

Let $\{\mathbf{v}_i\}_{i=1}^M$ be the detected camera rays (Eq. (5.12)) and $\{\mathbf{k}_j\}_{j=1}^K$ the catalogue inertial vectors (Eq. (5.8)). We seek the rotation matrix $\mathbf{R} \in SO(3)$ such that

$$\mathbf{k}_{\pi(i)} \approx \mathbf{R} \mathbf{v}_i, \quad (5.14)$$

where $\pi(i)$ is the catalogue index matched to detection i .

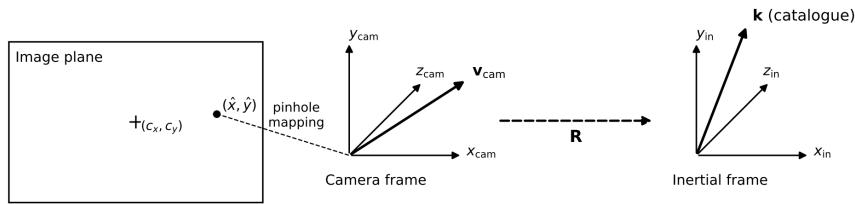


Figure 5.4: From Image Plane Coordinates to Inertial Frame Vectors

(1) RANSAC hypothesis generation from one angle-consistent pair. We precompute all detection-pair angles θ_{ij}^{det} via Eq. (5.9). Each RANSAC iteration:

1. randomly selects a detection pair (i, j) with measured θ_{ij}^{det} ,
2. retrieves candidate catalogue pairs (a, b) whose θ_{ab}^{cat} lies within the corresponding bin neighborhood,
3. builds a candidate attitude \mathbf{R} using the TRIAD method from the two vector correspondences.

TRIAD [16] construction. Given two non-collinear camera rays $(\mathbf{v}_1, \mathbf{v}_2)$ and their inertial matches $(\mathbf{k}_1, \mathbf{k}_2)$, TRIAD forms orthonormal bases:

$$\mathbf{t}_1 = \mathbf{v}_1, \quad \mathbf{t}_2 = \frac{\mathbf{v}_1 \times \mathbf{v}_2}{\|\mathbf{v}_1 \times \mathbf{v}_2\|}, \quad \mathbf{t}_3 = \mathbf{t}_1 \times \mathbf{t}_2, \quad (5.15)$$

and similarly $(\mathbf{s}_1, \mathbf{s}_2, \mathbf{s}_3)$ from $(\mathbf{k}_1, \mathbf{k}_2)$. Writing

$$\mathbf{C}_{\text{cam}} = [\mathbf{t}_1 \ \mathbf{t}_2 \ \mathbf{t}_3], \quad \mathbf{C}_{\text{in}} = [\mathbf{s}_1 \ \mathbf{s}_2 \ \mathbf{s}_3],$$

the attitude mapping camera \rightarrow inertial is

$$\mathbf{R} = \mathbf{C}_{\text{in}} \mathbf{C}_{\text{cam}}^{\top}. \quad (5.16)$$

(2) Inlier finding. For a candidate \mathbf{R} , every detected ray is projected into the inertial frame:

$$\hat{\mathbf{k}}_i = \mathbf{R} \mathbf{v}_i. \quad (5.17)$$

We assign it to the catalogue star with maximum dot product:

$$j^*(i) = \arg \max_j \mathbf{k}_j^{\top} \hat{\mathbf{k}}_i, \quad (5.18)$$

and declare an inlier if the angular error is below a tolerance ε :

$$\arccos(\mathbf{k}_{j^*(i)}^{\top} \hat{\mathbf{k}}_i) \leq \varepsilon. \quad (5.19)$$

The hypothesis with the highest inlier count is kept.

(3) Wahba [17] refinement using Davenport's Q-method [18]. Once a consensus set of inliers is found, the attitude is refined by solving Wahba's problem:

$$\mathbf{R}^* = \arg \min_{\mathbf{R} \in SO(3)} \sum_{i \in \mathcal{I}} w_i \|\mathbf{k}_{\pi(i)} - \mathbf{R} \mathbf{v}_i\|^2, \quad (5.20)$$

where \mathcal{I} is the inlier set and w_i are weights (unity in our default configuration). The implementation uses Davenport's Q-method to obtain the optimal rotation via a quaternion eigenproblem, then converts the quaternion to \mathbf{R}^* .

5.3.5 Boresight calculation

The final output of the attitude solver is \mathbf{R} mapping camera-frame rays to inertial directions. The camera boresight is the camera z -axis:

$$\mathbf{b}_{\text{cam}} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}, \quad \mathbf{b}_{\text{in}} = \mathbf{R}\mathbf{b}_{\text{cam}} = \begin{bmatrix} x \\ y \\ z \end{bmatrix}. \quad (5.21)$$

The corresponding pointing angles are:

$$\delta = \arcsin(z), \quad \alpha = \text{atan2}(y, x), \quad (5.22)$$

with α wrapped to $[0, 2\pi)$ and both angles reported in degrees. This yields the final **RA/Dec of the optical axis** for the considered frame.

5.3.6 Mapping results

After the refined attitude is obtained, every detection is matched again using the nearest-by-dot criterion (Eq. (5.18)) within a slightly looser angular threshold (e.g. 0.3°). The image (Fig. 5.5) is then annotated with detected centroids, matches star names/identifiers (HIP / Gaia / friendly name when available) and the estimated boresight RA/Dec.

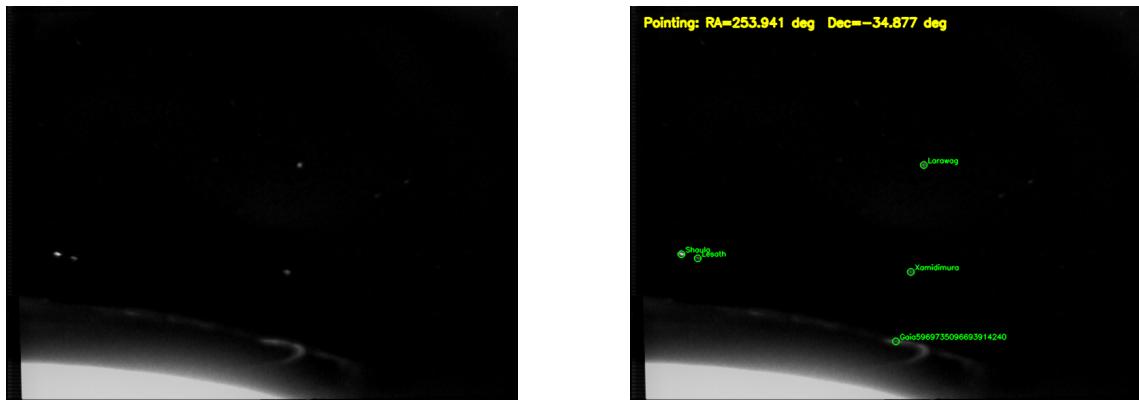


Figure 5.5: Final star mapping result. Matched stars are labeled in the image plane, and the estimated camera pointing (α, δ) is printed.

Chapter 6

Applications

6.1 Juno’s position vector

As defined in the research objectives (Chapter 1), the final goal of this study was to combine the limb geometry and the star-derived attitude to recover an image-constrained spacecraft position estimate. Ideally, this would close the loop on the geometric reconstruction.

However, the investigation in Chapter 4 demonstrated that the optical horizon detected in the ASC images does not correspond to the physical 1-bar reference surface provided by the SPICE kernels. Instead, it represents a diffuse boundary of stratospheric haze and scattered sunlight, introducing a systematic offset that cannot be fully corrected without a complex atmospheric radiative transfer model.

Consequently, applying the navigation algorithm to this specific dataset yields a biased position fix. Therefore, the following subsection focuses on the **theoretical derivation** of the horizon navigation method. This explains the mathematical logic intended to be applied, demonstrating how the position vector is constructed in a scenario where the visual limb is detected reliably.

6.1.1 Constructing the position vector

Limb pixels to camera-frame rays

Let (x_i, y_i) be pixel coordinates sampled from the detected horizon curve, where $i = 1, \dots, N$ indexes valid limb points. Because the images are already lens-corrected and despun (Chapter 3), a pinhole model is sufficient for ray construction.

Using the camera intrinsics (f_x, f_y, c_x, c_y) in pixels, each pixel is mapped to a normalized direction in the camera frame:

$$x_{n,i} = \frac{x_i - c_x}{f_x}, \quad y_{n,i} = -\frac{y_i - c_y}{f_y}, \quad (6.1)$$

where the minus sign accounts for the downward image y -axis convention. The corresponding (unnormalized) ray is

$$\tilde{\mathbf{v}}_i = \begin{bmatrix} x_{n,i} \\ y_{n,i} \\ 1 \end{bmatrix}, \quad \mathbf{v}_i = \frac{\tilde{\mathbf{v}}_i}{\|\tilde{\mathbf{v}}_i\|}. \quad (6.2)$$

Camera intrinsics. For the JUNO ASC camera, the focal length and principal point are known from calibration and are used directly. The focal lengths in pixel units are given by

$$f_x = \frac{f}{D_X}, \quad f_y = \frac{f}{D_Y}, \quad (6.3)$$

with $f = 20006 \mu\text{m}$ and $(D_X, D_Y) = (8.6, 8.3) \mu\text{m}$. The principal point is fixed at $(c_x, c_y) = (383, 257)$ pixels.

Camera rays to inertial rays

Let $\mathbf{R}_{I \leftarrow C}$ be the rotation matrix from camera frame to inertial frame estimated by the star mapping solver (Chapter 5). Each limb ray becomes an inertial unit vector:

$$\mathbf{u}_i = \mathbf{R}_{I \leftarrow C} \mathbf{v}_i, \quad \|\mathbf{u}_i\| = 1. \quad (6.4)$$

Spherical horizon constraint (tangent-ray cone)

Assume Jupiter is locally approximated as a sphere of radius R (as motivated in Chapter 4). Let \mathbf{r} be the unknown spacecraft position vector in a Jupiter-centered inertial frame (origin at Jupiter's center, axes aligned with J2000/ICRF). A limb ray from the spacecraft has the parametric form:

$$\mathbf{p}_i(s) = \mathbf{r} + s \mathbf{u}_i, \quad s \geq 0. \quad (6.5)$$

Tangency to the sphere $\|\mathbf{p}\| = R$ occurs when the quadratic in s has zero discriminant. Expanding $\|\mathbf{r} + s\mathbf{u}_i\|^2 = R^2$ yields:

$$s^2 + 2s(\mathbf{r}^\top \mathbf{u}_i) + (\|\mathbf{r}\|^2 - R^2) = 0. \quad (6.6)$$

The tangency condition is:

$$(\mathbf{r}^\top \mathbf{u}_i)^2 = \|\mathbf{r}\|^2 - R^2, \quad i = 1, \dots, N. \quad (6.7)$$

This implies all limb rays lie on a cone with axis \mathbf{r} , i.e. the dot product with the axis direction is (ideally) constant over all i (see Fig. 6.1).

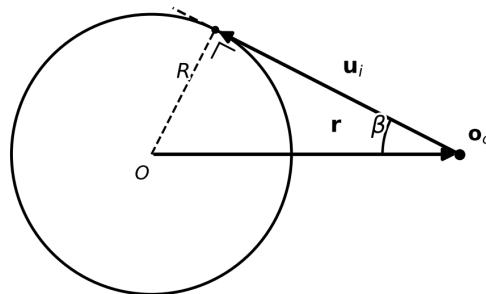


Figure 6.1: Spherical horizon tangent geometry

Cone-axis fit and range recovery

Define the unit position direction:

$$\mathbf{n} = \frac{\mathbf{r}}{\|\mathbf{r}\|}, \quad \|\mathbf{n}\| = 1. \quad (6.8)$$

Using Eq. (6.7) and $\mathbf{r} = \rho\mathbf{n}$ with $\rho = \|\mathbf{r}\|$, one obtains:

$$(\rho \mathbf{n}^\top \mathbf{u}_i)^2 = \rho^2 - R^2 \Rightarrow (\mathbf{n}^\top \mathbf{u}_i)^2 = 1 - \left(\frac{R}{\rho}\right)^2. \quad (6.9)$$

Hence $\mathbf{n}^\top \mathbf{u}_i$ should be (up to sign) constant across limb rays. Because the observed limb is only a short arc, we estimate \mathbf{n} and the constant c by least squares:

$$\min_{\|\mathbf{n}\|=1, c} \sum_{i=1}^N (\mathbf{n}^\top \mathbf{u}_i - c)^2. \quad (6.10)$$

Remark. Because the detected horizon typically covers only a short arc, the cone-axis estimation in Eq. (6.10) can become weakly constrained. In direction space, the limb rays $\{\mathbf{u}_i\}$ occupy only a small segment of the cone–sphere intersection, which limits the information available to determine the cone axis. As a result, the smallest-eigenvalue eigenvector of the covariance-like matrix \mathbf{C} may become ill-conditioned [19] and sensitive to outliers in the limb points and small biases in the attitude solution.

For a given \mathbf{n} , the optimal c is $c = \frac{1}{N} \sum_i \mathbf{n}^\top \mathbf{u}_i$. Substituting this back shows Eq. (6.10) minimizes the variance of the projections $\mathbf{n}^\top \mathbf{u}_i$. Writing $\boldsymbol{\mu} = \frac{1}{N} \sum_i \mathbf{u}_i$ and the covariance-like matrix

$$\mathbf{C} = \sum_{i=1}^N (\mathbf{u}_i - \boldsymbol{\mu})(\mathbf{u}_i - \boldsymbol{\mu})^\top, \quad (6.11)$$

the minimizing \mathbf{n} is the eigenvector of \mathbf{C} associated with its smallest eigenvalue. This is equivalent to performing Principal Component Analysis (PCA) on the set of ray vectors and selecting the component of minimum variance [20, Ch. 3]. We then set

$$c = \frac{1}{N} \sum_{i=1}^N \mathbf{n}^\top \mathbf{u}_i. \quad (6.12)$$

For a valid configuration, \mathbf{u}_i points approximately toward Jupiter, while \mathbf{n} points from

Jupiter's center to the spacecraft; therefore c is expected to be negative. If the eigenvector sign yields $c > 0$, we flip $\mathbf{n} \leftarrow -\mathbf{n}$.

Finally, combining $c = \mathbf{n}^\top \mathbf{u}_i \approx -\cos \beta$ with the geometric relation $\sin \beta = R/\rho$ gives:

$$\rho = \frac{R}{\sqrt{1 - c^2}}, \quad \mathbf{r} = \rho \mathbf{n}. \quad (6.13)$$

Equation (6.13) yields the full Jupiter-centered position vector in inertial coordinates and follows directly from the cone geometry introduced in Fig. 6.1.

Optional ellipsoid refinement

If the spherical approximation is insufficient, Jupiter can be represented as an oblate spheroid (semi-major axis a , semi-minor axis b). In a Jupiter-centered frame aligned with the ellipsoid axes, the surface is described by

$$\mathbf{x}^\top \mathbf{Q} \mathbf{x} = 1, \quad \mathbf{Q} = \text{diag}\left(\frac{1}{a^2}, \frac{1}{a^2}, \frac{1}{b^2}\right). \quad (6.14)$$

Tangency of the line $\mathbf{x}(s) = \mathbf{r} + s\mathbf{u}$ to the quadric is again obtained by a zero-discriminant condition, yielding

$$(\mathbf{u}^\top \mathbf{Q} \mathbf{r})^2 = (\mathbf{u}^\top \mathbf{Q} \mathbf{u}) (\mathbf{r}^\top \mathbf{Q} \mathbf{r} - 1). \quad (6.15)$$

A least-squares solve over Eq. (6.15) can be used to refine \mathbf{r} , using the spherical solution from Eq. (6.13) as initialization. In this work, the ellipsoidal model is employed only as a refinement step and not as a primary estimator.

6.1.2 Consistency checks

Tangency residuals

After estimating \mathbf{r} , the spherical tangency constraint is evaluated per limb ray using

$$e_i = (\mathbf{r}^\top \mathbf{u}_i)^2 - (\|\mathbf{r}\|^2 - R^2). \quad (6.16)$$

In the ideal noiseless model, $e_i = 0$. In practice, the distribution of $\{e_i\}$ provides an immediate diagnostic of:

- horizon detection quality (outliers from aurora, haze, or intensity gradient artifacts),
- residual camera modeling errors (e.g. imperfect distortion correction or intrinsics),
- attitude error from star mapping (rotation bias in $\mathbf{R}_{I \leftarrow C}$),
- physical limb offsets relative to the chosen reference radius R .

The spherical form of Eq. (6.16) is intentionally retained, as it yields a uniform and radius-independent consistency metric.

Cone projection dispersion

Because the cone model predicts $\mathbf{n}^\top \mathbf{u}_i \approx c$, the standard deviation

$$\sigma_c = \sqrt{\frac{1}{N} \sum_{i=1}^N (\mathbf{n}^\top \mathbf{u}_i - c)^2} \quad (6.17)$$

should remain small for a consistent solution. This metric is independent of the assumed planetary radius and isolates purely angular consistency.

Inter-frame stability

The dataset contains image pairs separated by 500 ms. Applying the full pipeline to both frames of a pair should yield nearly identical position vectors \mathbf{r} , with deviations consistent with the spacecraft displacement over 0.5 s. Large discrepancies indicate that the detected horizon points are not sampling the same physical limb (e.g., contamination by the bright scattering boundary in Fig. 2.5), rather than variations in camera intrinsics, which are fixed by calibration.

6.1.3 Transforming to Jupiter-centric coordinates

The position vector \mathbf{r} from Eq. (6.13) is expressed in a Jupiter-centered inertial frame aligned with J2000. For interpretation on the planet (latitude and longitude), or for com-

parison with body-fixed products, it is often desirable to express this vector in a Jupiter body-fixed frame.

Let $\mathbf{R}_{J \leftarrow I}(t)$ denote the rotation from inertial J2000 to the Jupiter-fixed frame at epoch t . The transformation is given by

$$\mathbf{r}_J(t) = \mathbf{R}_{J \leftarrow I}(t) \mathbf{r}_I. \quad (6.18)$$

This transformation is applied solely for interpretation and comparison; all estimation steps are performed in the inertial frame.

From $\mathbf{r}_J = [x \ y \ z]^\top$, the planetocentric longitude and latitude are obtained as

$$\lambda = \text{atan}2(y, x), \quad \varphi = \arctan 2\left(z, \sqrt{x^2 + y^2}\right), \quad (6.19)$$

with λ wrapped to $[0, 2\pi)$.

6.2 Atmospheric Investigation

6.2.1 Spatial resolution

The first step before any atmospheric analysis is to convert pixels into physical kilometers. This scaling factor, or spatial resolution S , is determined by the specific viewing geometry and intrinsic camera parameters. Using the principle of similar triangles (see sketch 6.2),

$$\frac{D}{f} = \frac{S}{p} \quad (6.20)$$

Where:

- D is the slant range to the limb provided by SPICE kernels (60,621 km for this observation).
- p is the width of a pixel ($8.6 \mu\text{m}$).
- f is the focal length ($20,006 \mu\text{m}$).

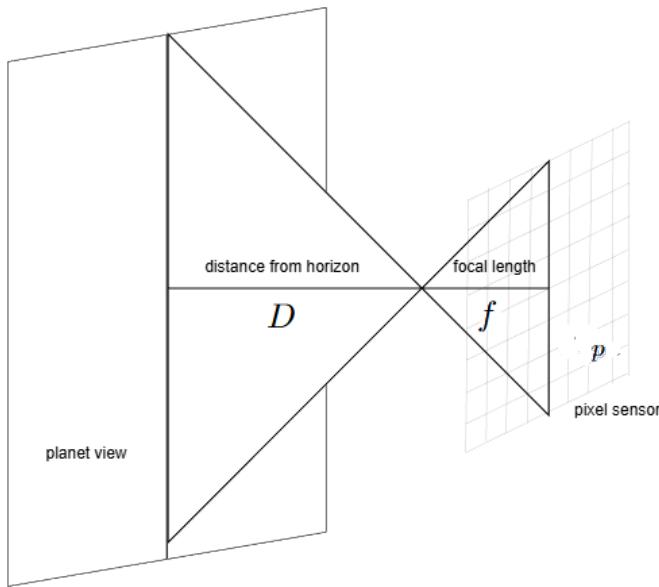


Figure 6.2: Schematic illustrating the projection of sensor pixel onto the target plane.

Substituting the value gives :

$$S = \frac{D \cdot p}{f} = \frac{60621\text{km} \cdot 8.6\mu\text{m}}{20006\mu\text{m}} = 26.06\text{km/px}$$

It is important to note that this value is valid strictly at the limb tangent point. For any features projected below the horizon (on the planetary disk), the distance D decreases slightly, resulting in a finer spatial resolution (smaller km/px value).

6.2.2 Validation of thermosphere detection

The atmospheric region extending above the limb is identified as the thermosphere, the upper boundary layer that separates Jupiter's envelope from the vacuum. This layer begins approximately 350 km above the 1-bar level [21]. Unlike the underlying troposphere, which is governed by solar heating and convection, the thermosphere is defined by a dramatic rise in temperature and intense interaction with the Jovian magnetosphere.

The defining characteristic of this layer is its immense extent, driven by temperatures that rise dramatically with altitude, ranging from 160 K to 900 K. This extreme thermal expansion is fueled by high-energy inputs: the layer absorbs solar ultraviolet radiation and is bombarded by precipitating charged particles from the magnetosphere. This particle

precipitation is the primary driver of Jupiter's powerful auroral emissions. The detection of these auroral signatures in the data confirms that the observed limb profile corresponds to this active upper-atmospheric layer.

Scale Height Correlation

A critical parameter for characterizing this atmospheric layer is the scale height (H), which represents the vertical distance over which the atmospheric density decreases by a factor of e . In a hydrostatic atmosphere, it is defined as

$$H = \frac{kT}{mg} \quad (6.21)$$

where k is the Boltzmann constant, T is the kinetic temperature, m is the mean molecular mass, and g is the local gravitational acceleration. Due to very high temperatures, the atmosphere expands significantly, resulting in scale heights that are orders of magnitude larger than those found in the cloud decks below ($H \approx 27$ km at 1-bar level).

To analyze the intensity profile, radial line scans were performed on the cleaned data (cf. 6.3). As discussed in Chapter 4, the horizon is derived from SPICE kernels, eliminating approximation errors. These line scans are averaged and fit to an exponential curve (cf. 6.4) of the following form :

$$I(r) = I_0 e^{-r/H} \quad (6.22)$$

The fitting provides an estimated scale height H of 4.8 pixels, translating to ~ 125 km using the spatial resolution derived earlier.

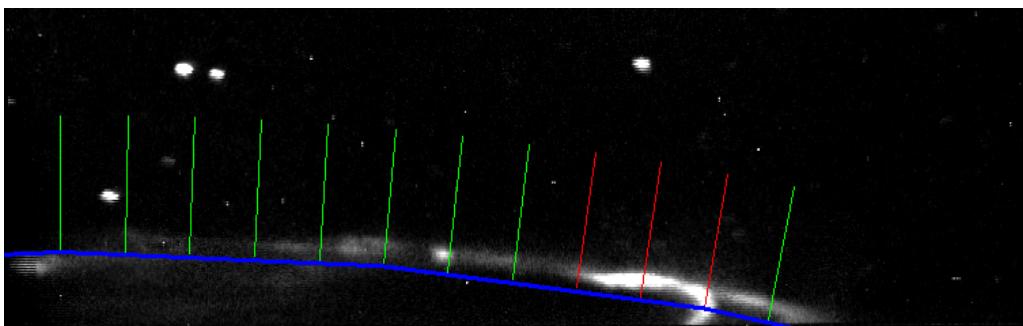


Figure 6.3: Radial line scans of Jovian atmosphere. Red lines are omitted because aurora interferes with the signal. Image gained x6.

This value correlates strongly with theoretical predictions, driven primarily by the thermal structure of the layer. As defined in equation 6.21, the scale height is directly proportional to the temperature ($H \propto T$). While the troposphere is cold (130 K), Yelle & Miller [21] describe the thermosphere as a region of extreme heating, where temperatures rise to between 800 K and 1000 K. Consequently, the scale height is effectively multiplied by a factor of roughly 5 to 7 relative to the lower atmosphere, where the scale height is 27 km. The derived value of 125 km is consistent with these high temperatures, confirming that the exponential decay detected in the imagery corresponds to the hot, expanded thermosphere.

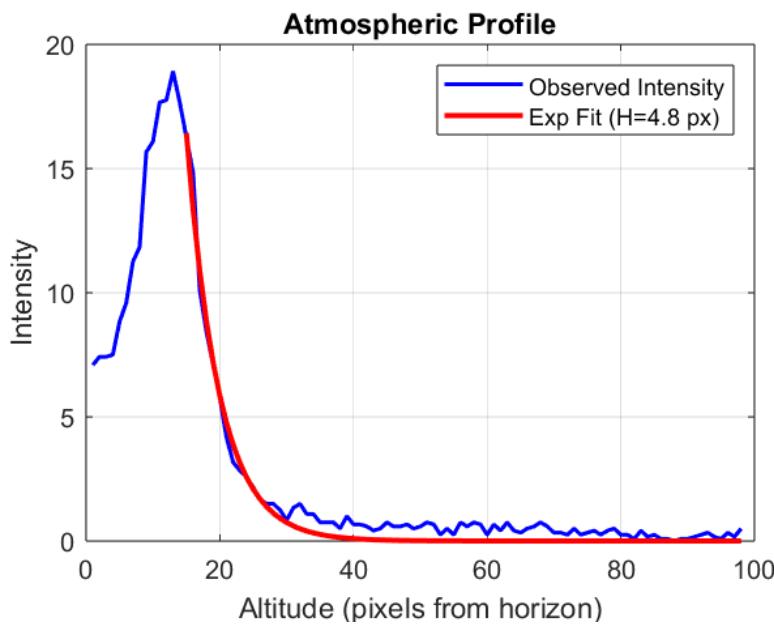


Figure 6.4: Averaged radial line scan starting from 1 bar horizon

Chapter 7

Discussion and Conclusions

7.1 Discussion

This study investigated whether the Juno ASC, originally designed as a navigation star tracker, can be repurposed into a scientifically useful imager for atmospheric inspection and as a supporting sensor for optical navigation. The results show that meaningful information can be extracted, but only when a dedicated processing chain (briefly summarized by Figure 7.1) designed for sensor bias, radiation contamination, and rotation-induced smear is applied.

7.1.1 What works best for this Juno dataset

Radiometric calibration and structured-bias suppression Dark-frame subtraction combined with optional column-bias removal removes sensor-induced background bias. In this dataset, the large fraction of deep-space pixels makes a column-median estimator effective for suppressing column-correlated residuals.

Radiation mitigation Radiation hits appear as compact, high-gradient impulses that must be removed without attenuating the star point-spread function or the faint limb gradient. Global smoothing reduces speckles but degrades both centroiding and edge localization. The gradient-based switching median filter provides the most suitable trade-off by selectively suppressing impulse artifacts while preserving broader structures. The effect

of this step is illustrated by the reduction of radiation speckles between the intermediate frames in Fig. 7.1.

Robust star mapping for inertial attitude recovery Despite the limited number of detectable background stars and strong intensity variation across the frame, the star mapping pipeline successfully recovers the camera-to-inertial rotation. A hybrid detection strategy increases the usable star set, while the robust matching stage resolves catalogue ambiguities by validating hypotheses against many additional detections (RANSAC inliers). After a consistent correspondence set is found, Wahba refinement directly provides the spacecraft orientation in inertial space.

Reconstruction through despinning and lens correction The 125 ms delay between odd and even readout fields corresponds to a measurable angular rotation, making raw frames inconsistent with a single pinhole camera model. Rotation-compensated despinning restores a coherent capture. While lens distortion is visually subtle, it becomes relevant when pixel-level errors are converted to physical distance at the limb.

Horizon estimation on short arcs Because the limb occupies only a short arc in a 16° field of view, standard algebraic least-squares circle fitting (Kåsa) becomes strongly biased. The Hyperaccurate fit is more stable for this regime and yields a radius estimate that is consistent with SPICE to within the reported residual level. This improves geometric extraction from the image, but does not remove the systematic mismatch that originates from the horizon definition.

7.1.2 Limitations and sources of error

Figure 7.1 provides an overview of the full processing pipeline: from the raw image (1) through despinning (2) and denoising (3) to an image suitable for feature extraction. The final outputs illustrate two parallel branches, star detection and catalogue-based attitude annotation (4) and horizon extraction via limb detection and geometric overlay (5).

Horizon detection compared to SPICE reference The dominant discrepancy in the horizon comparison is physical rather than numerical. The ASC detects a radiometric transition shaped by scattering and haze at altitudes above the 1-bar level used by SPICE. Consequently, horizon-based optical navigation inherits a systematic range bias unless an atmospheric offset model (or equivalent correction) is introduced.

Dataset scope Conclusions are based on a limited observation set, restricting validation across illumination conditions radiation variability.

Star detection and matching limitations Star mapping performance is constrained by the photon-limited regime and the presence of strong stray light gradients near the Jovian limb. Faint stars near the noise floor can be missed or fragmented into spurious detections, while radiation residuals can introduce false candidates. Although the robust solver mitigates these effects through inlier screening, frames with very low star counts or heavy contamination reduce attitude confidence and can increase sensitivity to centroiding and catalog matching errors.

Onboard infeasibility Full-frame calibration, filtering, and dense remapping exceed the μ ASC processing constraints at 4 Hz, confining the complete reconstruction to ground-based processing.

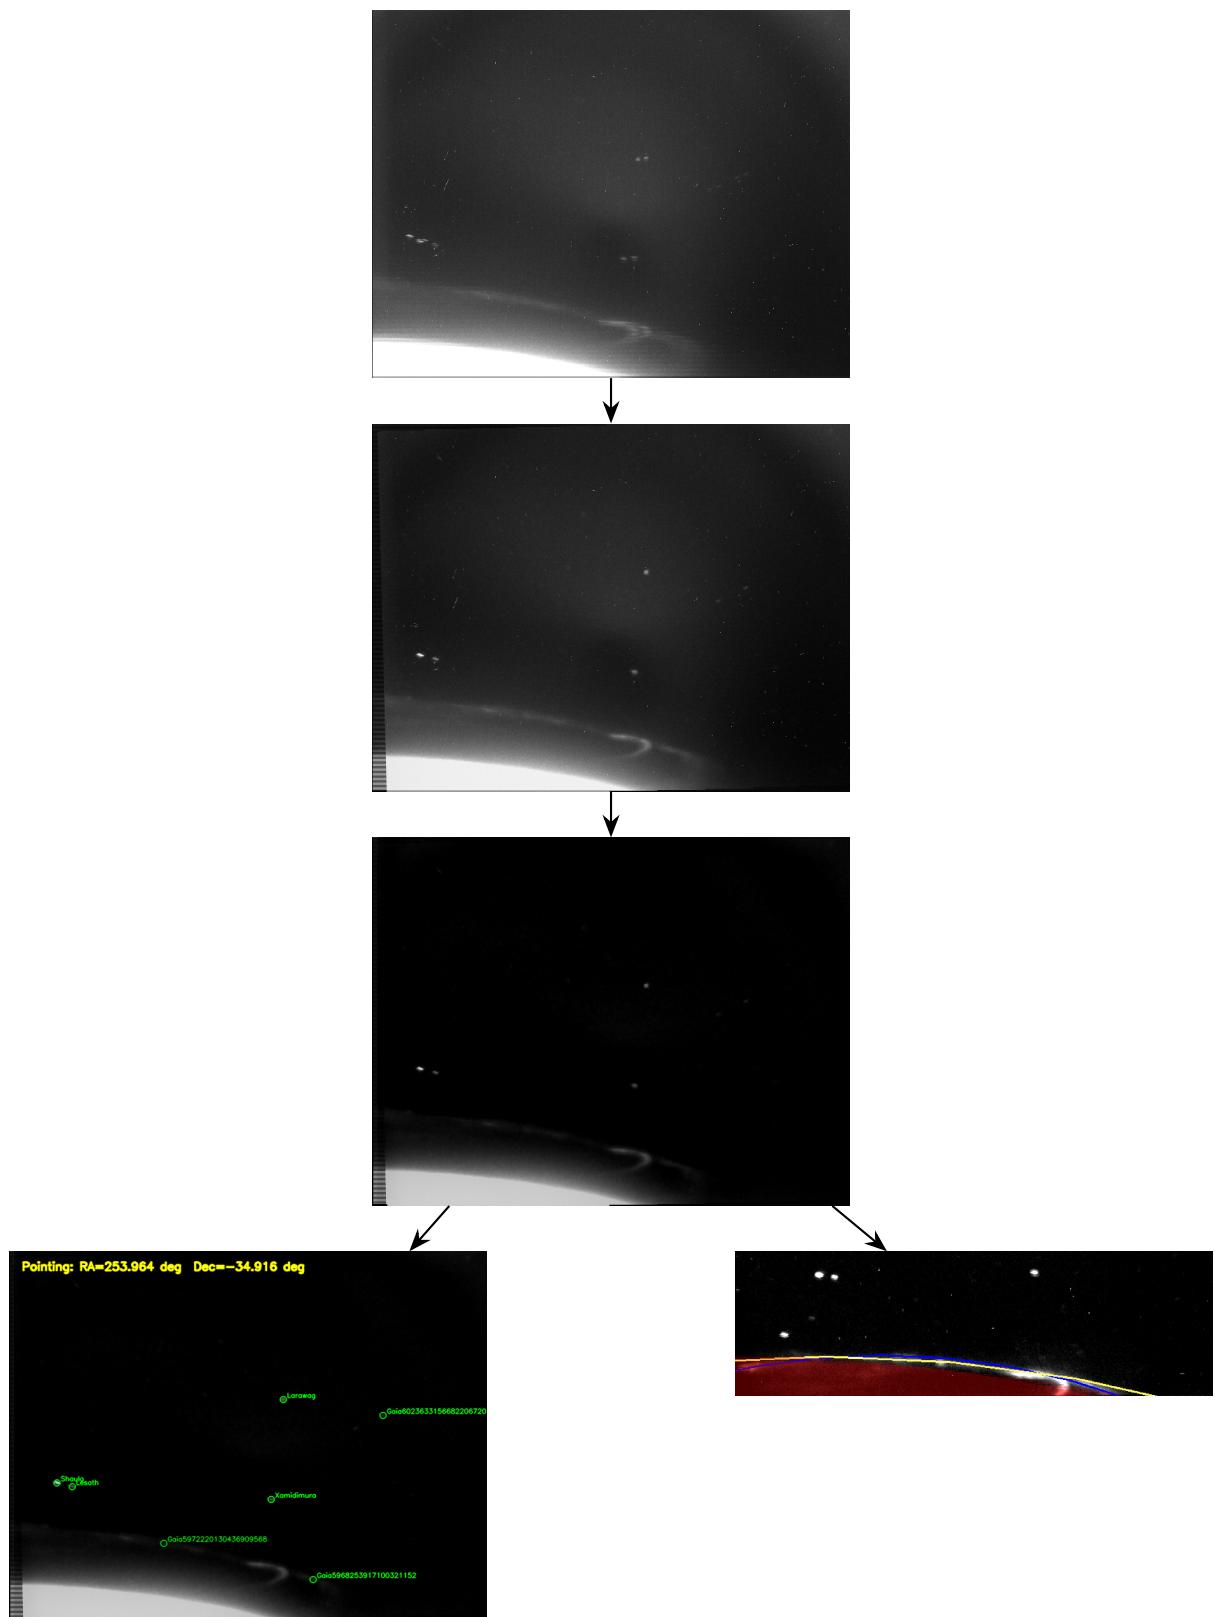


Figure 7.1: From raw image to star detection and horizon extraction.

7.2 Conclusions

7.2.1 Main findings

- A tailored image processing sequence can convert raw ASC frames into scientifically interpretable imagery.
- The gradient-based switching median filter preserves both images features more effectively than global smoothing in this dataset.
- Hyperaccurate circle fitting provides a good estimate of the horizon line, but optical navigation accuracy is inherently limited by the data itself without additional modeling.
- Robust star mapping (hybrid detection + RANSAC/Wahba refinement) yields a stable camera-to-inertial attitude estimate from sparse star fields.

7.2.2 Recommendations for future work

1. Parameterize the atmospheric offset. Introduce $R_{eff} = R + \Delta h$ or an intensity-based limb model so that the radiometric-to-geometric discrepancy becomes an estimable correction rather than a persistent bias.
2. Estimate the feasibility of removing heavy processing steps and still get a god accuracy on attitude determination for onboard attitude and, potentially, position determination.

Overall, the ASC images can be repurposed into a useful scientific and geometric dataset, but the limiting factors are dominated by the little amount of available data, sensor architecture, compute constraints (for the onboard attitude determination), and challenges in interpreting the observed limb.

Bibliography

- [1] National Aeronautics and Space Administration. NASA's Juno Mission to Remain in Current Orbit at Jupiter, February 2017. News release (Release 17-018), Feb 17, 2017.
- [2] NASA Science. Juno, 2025. Mission overview page (page last updated Sep 03, 2025).
- [3] NASA Jet Propulsion Laboratory. Juno, 2025. JPL mission overview page.
- [4] Space Instrumentation Group. Advanced Stellar Compass JUNO Software Interface Specification. Software Interface Specification JN-DTU-SP-3001, National Space Institute, Technical University of Denmark, Kongens Lyngby, Denmark, 2024. Issue 1.6.
- [5] Matija Herceg, P. Jørgensen, J. Jørgensen, and J. Connerney. Thermoelastic response of the juno spacecraft's solar array/magnetometer boom and its applicability to improved magnetic field investigation. *Earth and Space Science*, 7, 11 2020.
- [6] Troelz Denver, Julia Sushkova, John L. Jørgensen, Leonardo Ghizoni, Matija Herceg, Christina Toldbo, Mathias Benn, Peter S. Jørgensen, René Fléron, J. E. P. Connerney, Heidi N. Becker, and Scott J. Bolton. The Juno ASC as an energetic particle counter. *Space Science Reviews*, 220(8):86, 2024.
- [7] Navigation and Ancillary Information Facility. Time Required Reading. https://naif.jpl.nasa.gov/pub/naif/toolkit_docs/C/req/time.html. Accessed: 2025-12-10.

- [8] Pieter G. van Dokkum. Cosmic-ray rejection by laplacian edge detection. *Publications of the Astronomical Society of the Pacific*, 113(789):1420, nov 2001.
- [9] Olgierd Stankiewicz, Gauthier Lafruit, and Marek Domański. Chapter 1 - multiview video: Acquisition, processing, compression, and virtual view rendering. In Rama Chellappa and Sergios Theodoridis, editors, *Academic Press Library in Signal Processing, Volume 6*, pages 32–33. Academic Press, 2018.
- [10] Carlo Tomasi. Vector representation of rotations. Duke University, Computer Science 527 Lecture Notes, 2013. Accessed: 2025-11-20.
- [11] Vision Doctor. Optical errors - distortion, n.d.
- [12] A. Al-Sharadqah and N. Chernov. Error analysis for circle fitting algorithms. *Electronic Journal of Statistics*, 3:886–911, 2009.
- [13] European Space Agency. Hipparcos catalogues. <https://www.cosmos.esa.int/web/hipparcos/catalogues>, 1997. ESA SP-1200.
- [14] European Space Agency. Gaia data release 3 (dr3) catalogue. <https://www.cosmos.esa.int/web/gaia/dr3>, 2022. Gaia DR3.
- [15] Daniele Mortari, Malak A. Samaan, Christian Bruccoleri, and John L. Junkins. The “pyramid” star pattern recognition algorithm. *Navigation*, 51(3):171–183, 2004.
- [16] Malcolm D Shuster and SD Oh. Three-axis attitude determination from vector observations. *Journal of Guidance and Control*, 4(1):70–77, 1981.
- [17] Grace Wahba. A least squares estimate of satellite attitude. *SIAM review*, 7(3):409–409, 1965.
- [18] James R Wertz. *Spacecraft attitude determination and control*, volume 73. Springer Science & Business Media, 1978.
- [19] Aurelie Bellemans, Tim De Troyer, Mark Runacles, and Chris Lacor. *Basic Techniques for Computer Simulations*. Vrije Universiteit Brussel, 2023. Lecture notes, Faculty of Engineering.

- [20] Tue Herlau, Mikkel N. Schmidt, and Morten Mørup. *Introduction to Machine Learning and Data Mining*. Technical University of Denmark, 2023. Lecture notes, Fall 2023, version 1.0.
- [21] Roger V. Yelle and Steve Miller. Jupiter's thermosphere and ionosphere. In Fran Bagenal, Timothy E. Dowling, and William B. McKinnon, editors, *Jupiter: The Planet, Satellites and Magnetosphere*, chapter 9, pages 185–218. Cambridge University Press, Cambridge, UK, 2004.

Appendix A

Additional Material

A.1 Star recognition output from whole dataset

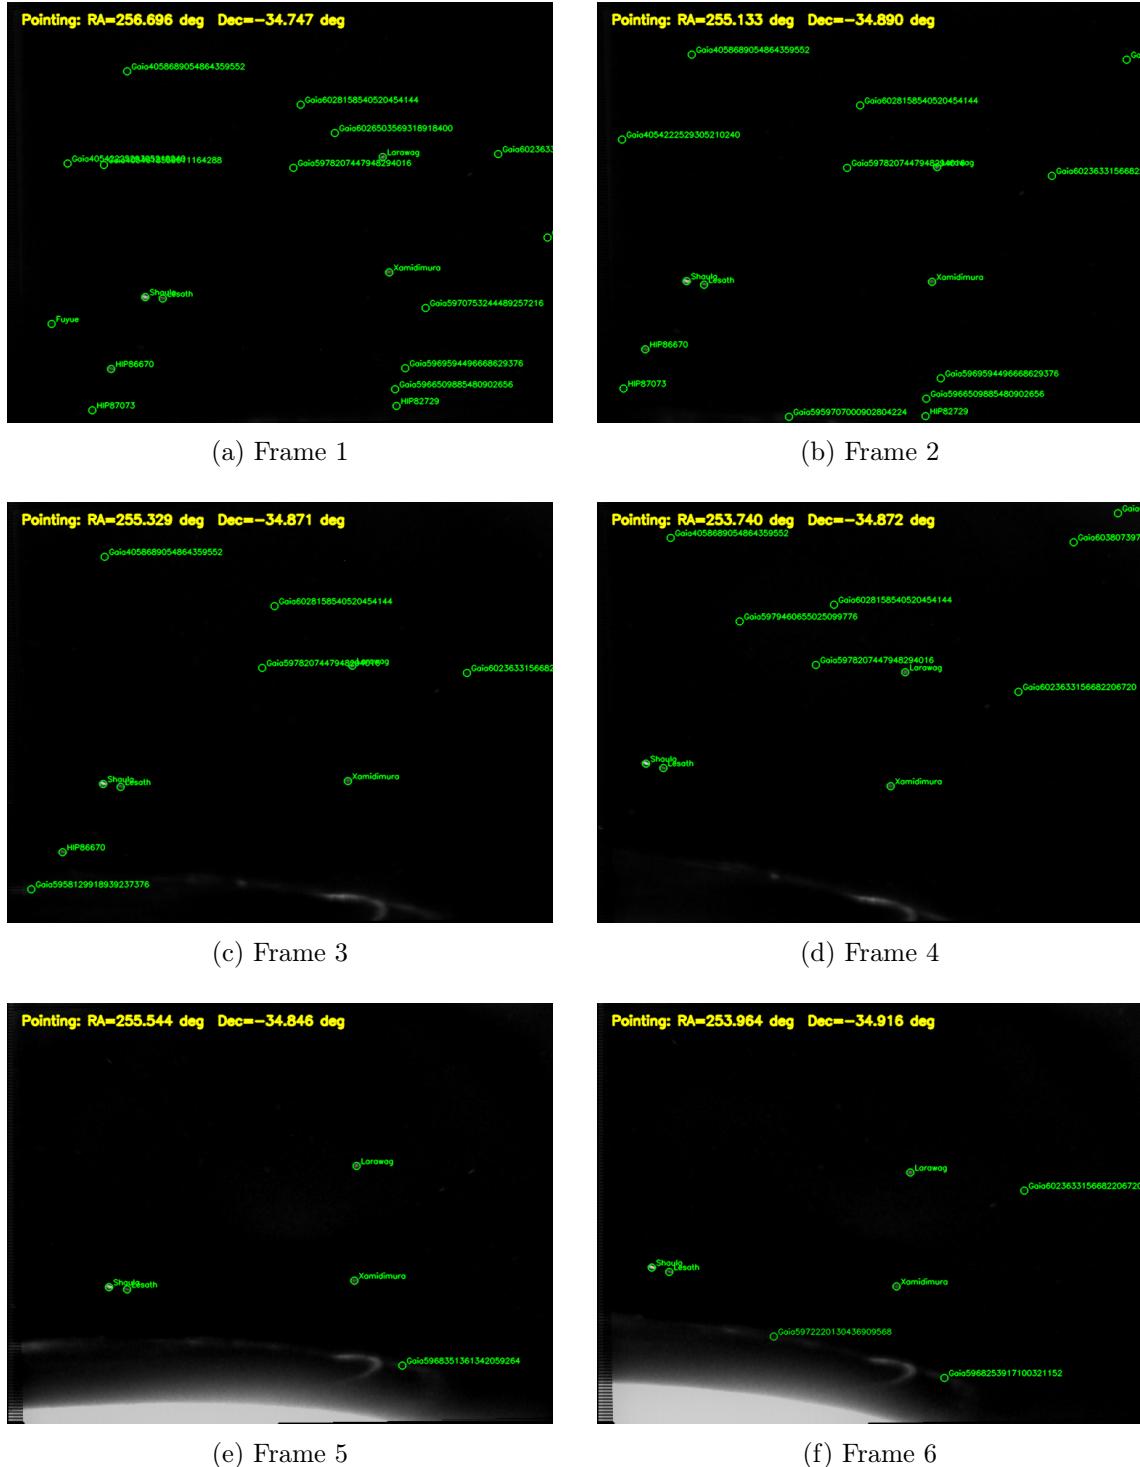


Figure A.1: Star-tracking results for six representative ASC frames.

A.2 Source Code

Note: The source code presented in this appendix was assembled, refactored, and cleaned with the assistance of AI-based tools to improve code structure, readability, and consistency across different scripts, while preserving the original algorithms and logic designed by the authors.

A.2.1 Main script

```
1 #!/usr/bin/env python3
2 """
3 TotalCode: unified JUNO pipeline with all toggles + star tracker in
4     one file.
5
6 Includes:
7 - Despin (per-pixel rotation)
8 - Lens distortion correction
9 - Denoising (dark frame, flatten, spike, median/Gaussian)
10 - Column-median glow removal
11 - Fourier notch filtering
12 - Edge tracking, circle fits, limb profiles
13 - Star tracker (catalog loading, detection, attitude, annotation)
14
15 Use the feature toggles below to enable/disable each stage.
16 """
17
18 from __future__ import annotations
19
20 from dataclasses import dataclass, field
21 from pathlib import Path
22 from typing import Dict, Iterable, List, Optional, Sequence, Tuple
23
24 import cv2
25 import numpy as np
```

```
25 from PIL import Image
26 import star_tracker
27
28 # --- Feature toggles
29
30 APPLY_DESPIN = True
31 APPLY_LENS = False
32 APPLY_DENOISE = True
33 APPLY_FFT_NOTCH = False
34 APPLY_EDGE_TRACKING = True
35 APPLY_CIRCLE_FITS = True
36 APPLY_LIMB_PROFILES = True
37 APPLY_RESIDUAL_GLOW = False # column median subtract
38 APPLY_STAR_TRACKER = True
39 STAR_USE_DESPIN_MASK = False # if True, pass the despin overlap
40 mask into star tracker detection
41
42 # --- Paths
43
44 SCIENCE_DIR = Path("/Users/vojtadeconinck/Downloads/
45     ImageAnalysis30300/JUNO_Despin/Images") / "ScienceCalibrated"
46 OUTPUT_DIR = Path("/Users/vojtadeconinck/Downloads/
47     ImageAnalysis30300/JUNO_Despin/Images") / "Combined"
48 MASK_OUTPUT_DIR = OUTPUT_DIR / "Masks"
49 STAR_CATALOG_PATH = Path("/Users/vojtadeconinck/Downloads/
50     ImageAnalysis30300/combined_stars.csv")
51
52 # --- Camera / geometry constants
53
54
```

```
50 OMEGA = np.array([-8.54302424e-05, -2.26443686e-04, -2.08965026e
51     -01], dtype=np.float64)
52 OMEGA_FRAME = "spacecraft" # "spacecraft" or "camera"
53 DELTA_T = 0.13 # seconds between odd and even fields
54
55 SC_TO_CHUD_MATRIX = np.array([
56     [0.99995302, 0.00569167, 0.00784704],
57     [0.00383956, -0.97582589, 0.21851563],
58     [0.00890106, -0.21847523, -0.9758019],
59 ])
60
61 EFL_MICRONS = 20_006.0
62 PIXEL_SIZE_X_MICRONS = 8.6
63 PIXEL_SIZE_Y_MICRONS = 8.3
64 PRINCIPAL_POINT = (383.0, 257.0) # (cx, cy)
65 SENSOR_SHAPE = (580, 752) # (H, W)
66
67 # Lens distortion
68 LENS_KAPPA = 3.3e-8
69
70 # --- Denoise settings
71
72 @dataclass
73 class FlattenConfig:
74     mask_percentile: float = 95.0
75     dilate_kernel: int = 11
76     sigma: float = 40.0
77     column_correction: bool = True
78     column_mask_percentile: float = 99.0
79     column_dilate: int = 5
80
81
```

```
82 @dataclass
83 class DenoiseConfig:
84     apply_dark_frame: bool = False
85     dark_frame_path: str | None = None
86     dark_frame: np.ndarray | None = field(default=None, repr=False)
87
88     apply_column_median: bool = False
89
90     apply_flatten: bool = True
91     flatten: FlattenConfig = field(default_factory=FlattenConfig)
92
93     spike_mode: str = "gradient" # "mean" or "gradient"
94     spike_neighborhood: int = 1
95     spike_threshold: float = 45.0
96     gradient_threshold: float = 45.0
97
98     median_kernel: int = 3
99     gaussian_kernel: int = 0
100    gaussian_sigma: float = 0.0
101
102
103 DENOISE_SETTINGS = DenoiseConfig(
104     apply_dark_frame=False,
105     dark_frame_path=str(Path("/Users/vojtadeconinck/Downloads/
106         ImageAnalysis30300/BlackRefImages/master_black.png")),
107     apply_column_median=False,
108     apply_flatten=True,
109     flatten=FlattenConfig(),
110     spike_mode="gradient",
111     spike_neighborhood=1,
112     spike_threshold=45.0,
113     gradient_threshold=45.0,
114     median_kernel=3,
115     gaussian_kernel=0,
```

```
115     gaussian_sigma=0.0,
116 )
117
118 # --- Fourier notch default (legacy manual list)
119
120 MANUAL_NOTCHES = [
121     (75, 0),
122     (115, 0),
123     (155, 0),
124     (0, 40),
125     (0, 80),
126 ]
127 NOTCH_RADIUS = 25.0
128
129 # --- Basic helpers
130
131
132 def _load_grayscale(path: Path) -> np.ndarray:
133     arr = np.array(Image.open(path))
134     if arr.ndim != 2:
135         raise ValueError(f"{path} is not a single-channel image.")
136     return arr
137
138
139 def _save_grayscale(data: np.ndarray, path: Path, dtype: np.dtype)
140     -> None:
141     info = np.iinfo(dtype)
142     clipped = np.clip(np.rint(data), info.min, info.max).astype(
143         dtype)
144     path.parent.mkdir(parents=True, exist_ok=True)
145     Image.fromarray(clipped).save(path)
```

```
145
146 def _to_8bit(arr: np.ndarray, dtype: np.dtype) -> np.ndarray:
147     info = np.iinfo(dtype)
148     scaled = np.clip(arr, info.min, info.max) / float(info.max)
149     scaled *= 255.0
150
151     return np.clip(np.rint(scaled), 0, 255).astype(np.uint8)
152
153 # --- Lens correction
154
155 def _build_lens_maps(kappa: float) -> tuple[np.ndarray, np.ndarray]:
156     H, W = SENSOR_SHAPE
157     cx, cy = PRINCIPAL_POINT
158     pix_ratio = PIXEL_SIZE_Y_MICRONS / PIXEL_SIZE_X_MICRONS
159
160     x = np.arange(W, dtype=np.float32) - cx
161     y = (np.arange(H, dtype=np.float32)[:, None] - cy) * pix_ratio
162     r2 = x[None, :] * x[None, :] + y * y
163     S = 1.0 + kappa * r2
164
165     map_x = (x[None, :] * S + cx).astype(np.float32)
166     map_y = ((y * S) / pix_ratio + cy).astype(np.float32)
167
168     return map_x, map_y
169
170 def _apply_lens_correction(
171     img: np.ndarray,
172     maps: tuple[np.ndarray, np.ndarray],
173     *,
174     interpolation: int = cv2.INTER_LINEAR,
175     border_value: float = 0.0,
176 ) -> np.ndarray:
```

```
177     map_x, map_y = maps
178
179     corrected = cv2.remap(
180         img.astype(np.float32),
181         map_x,
182         map_y,
183         interpolation=interpolation,
184         borderMode=cv2.BORDER_CONSTANT,
185         borderValue=border_value,
186     )
187
188     return corrected
189
190
191 # --- Denoising
192
193
194
195 def _as_float(gray: np.ndarray) -> np.ndarray:
196
197     return gray.astype(np.float32, copy=False) if gray.dtype != np.
198         float32 else gray
199
200
201
202
203 def _restore_dtype(clean: np.ndarray, dtype) -> np.ndarray:
204
205     if np.issubdtype(dtype, np.floating):
206
207         return clean.astype(dtype)
208
209     info = np.iinfo(dtype)
210
211     clipped = np.clip(np.rint(clean), info.min, info.max)
212
213     return clipped.astype(dtype)
214
215
216
217
218 def _load_dark_frame(path: str) -> np.ndarray:
219
220     arr = cv2.imread(path, cv2.IMREAD_UNCHANGED)
221
222     if arr is None:
223
224         raise FileNotFoundError(f"Dark frame not found: {path}")
225
226     if arr.ndim != 2:
```

```
208     raise ValueError(f"Dark frame must be single-channel: {path}
209         }")
210
211
212 def subtract_dark_frame(gray: np.ndarray, dark: np.ndarray) -> np.
213     ndarray:
214     if gray.shape != dark.shape:
215         raise ValueError(f"Dark frame shape {dark.shape} does not
216             match image {gray.shape}")
217
218     return _as_float(gray) - _as_float(dark)
219
220
221
222 def _build_mask(arr: np.ndarray, percentile: float, dilate_kernel:
223     int) -> np.ndarray:
224     thr = np.percentile(arr, percentile)
225     mask = arr > thr
226     k = max(1, int(dilate_kernel))
227     if k > 1:
228         kernel = np.ones((k, k), np.uint8)
229         mask = cv2.dilate(mask.astype(np.uint8), kernel, iterations
230             =1).astype(bool)
231
232     return mask
233
234
235
236 def _column_correction(flat: np.ndarray, config: FlattenConfig) ->
237     np.ndarray:
238     work = flat.copy()
239     mask = _build_mask(work, config.column_mask_percentile, config.
240         column_dilate)
241     H, W = work.shape
242     for x in range(W):
243         col = work[:, x]
244         good = ~mask[:, x]
```

```
235     if not np.any(good):
236         continue
237     offset = np.median(col[good])
238     work[:, x] -= offset
239
240
241
242 def column_median_subtract(gray: np.ndarray) -> np.ndarray:
243     f = _as_float(gray)
244     med = np.median(f, axis=0, keepdims=True)
245
246
247
248 def flatten_background(gray: np.ndarray, config: FlattenConfig | None = None) -> np.ndarray:
249     if gray.ndim != 2:
250         raise ValueError("flatten_background expects a single-
251                         channel image.")
252     cfg = config or FlattenConfig()
253     f = _as_float(gray)
254
255     mask = _build_mask(f, cfg.mask_percentile, cfg.dilate_kernel)
256     bg_level = np.median(f[~mask]) if np.any(~mask) else float(np.
257         median(f))
258
259     f_bgfit = f.copy()
260     f_bgfit[mask] = bg_level
261
262     bg = cv2.GaussianBlur(f_bgfit, (0, 0), sigmaX=cfg.sigma, sigmaY
263                           =cfg.sigma)
264     f_flat = f - bg
265
266
267     if cfg.column_correction:
268         f_flat = _column_correction(f_flat, cfg)
```

```
265     return f_flat
266
267
268 def remove_radiation_spikes(gray: np.ndarray, neighborhood: int,
269     threshold: float) -> np.ndarray:
270     if gray.ndim != 2:
271         raise ValueError("remove_radiation_spikes expects a single-
272             channel image.")
273     n = max(1, int(neighborhood))
274     k = 2 * n + 1
275     f = _as_float(gray)
276     kernel = np.ones((k, k), dtype=np.float32)
277     summed = cv2.filter2D(f, -1, kernel, borderType=cv2.
278             BORDER_REFLECT)
279     neighbor_sum = summed - f
280     neighbor_count = k * k - 1
281     neighbor_mean = neighbor_sum / float(neighbor_count)
282     mask = (f - neighbor_mean) > threshold
283     cleaned = f.copy()
284     cleaned[mask] = neighbor_mean[mask]
285     return cleaned
286
287
288 def remove_gradient_spikes(gray: np.ndarray, threshold: float) ->
289     np.ndarray:
290     if gray.ndim != 2:
291         raise ValueError("remove_gradient_spikes expects a single-
292             channel image.")
293     f = _as_float(gray)
294     out = f.copy()
295
296     center = f[1:-1, 1:-1]
297     left = f[1:-1, :-2]
298     right = f[1:-1, 2:]
```

```
294     up = f[:-2, 1:-1]
295     down = f[2:, 1:-1]
296
297     horiz_peak = (center - left > threshold) & (center - right >
298                 threshold)
299     vert_peak = (center - up > threshold) & (center - down >
300                 threshold)
301     mask = horiz_peak & vert_peak
302
303
304     if np.any(mask):
305         neighbors = np.stack([left, right, up, down], axis=0)
306         med = np.median(neighbors, axis=0)
307         out_view = out[1:-1, 1:-1]
308         out_view[mask] = med[mask]
309
310     return out
311
312
313
314
315
316
317
318
319
320
321
322
323
```

```
def clean_image(gray: np.ndarray, config: DenoiseConfig | None =
None) -> np.ndarray:
    if gray.ndim != 2:
        raise ValueError("clean_image expects a single-channel
image.")
    cfg = config or DenoiseConfig()
    work = _as_float(gray)

    if cfg.apply_dark_frame:
        dark = cfg.dark_frame
        if dark is None and cfg.dark_frame_path:
            dark = _load_dark_frame(cfg.dark_frame_path)
            cfg.dark_frame = dark
        if dark is not None:
            work = subtract_dark_frame(work, dark)

    if cfg.apply_column_median:
```

```
324     work = column_median_subtract(work)
325
326     if cfg.apply_flatten:
327         work = flatten_background(work, cfg.flatten)
328
329     if cfg.spike_mode == "gradient":
330         work = remove_gradient_spikes(work, cfg.gradient_threshold)
331     else:
332         work = remove_radiation_spikes(work, cfg.spike_neighborhood,
333                                         cfg.spike_threshold)
334
335     if cfg.median_kernel and cfg.median_kernel >= 3:
336         k = cfg.median_kernel if cfg.median_kernel % 2 == 1 else
337             cfg.median_kernel + 1
338         work = cv2.medianBlur(work.astype(np.float32), k)
339
340     if cfg.gaussian_kernel and cfg.gaussian_kernel > 0:
341         k = cfg.gaussian_kernel if cfg.gaussian_kernel % 2 == 1
342             else cfg.gaussian_kernel + 1
343         work = cv2.GaussianBlur(work, (k, k), cfg.gaussian_sigma)
344
345     return _restore_dtype(work, gray.dtype)
346
347 # --- Despin (per-pixel rotation)
348
349 def rodrigues(omega: np.ndarray, delta_t: float) -> np.ndarray:
350     theta = float(np.linalg.norm(omega) * delta_t)
351
352     if theta == 0.0:
353         return np.eye(3, dtype=np.float64)
354
355     k = omega / np.linalg.norm(omega)
356
357     kx, ky, kz = k
358
359     K = np.array(
```

```
354     [
355         [0.0, -kz, ky],
356         [kz, 0.0, -kx],
357         [-ky, kx, 0.0],
358     ],
359     dtype=np.float64,
360 )
361 sin_t = np.sin(theta)
362 cos_t = np.cos(theta)
363 return np.eye(3) + sin_t * K + (1.0 - cos_t) * (K @ K)
364
365
366 def _pixel_rays(x_full: np.ndarray, y_full: np.ndarray) -> np.
367 ndarray:
368     cx, cy = PRINCIPAL_POINT
369     x_cam = (x_full - cx) * (PIXEL_SIZE_X_MICRONS / EFL_MICRONS)
370     y_cam = (y_full - cy) * (PIXEL_SIZE_Y_MICRONS / EFL_MICRONS)
371     rays = np.stack([x_cam, y_cam, np.ones_like(x_cam)], axis=-1)
372     return rays
373
374
375 def _project_to_full_pixels(rays: np.ndarray) -> Tuple[np.ndarray,
376 np.ndarray]:
377     cx, cy = PRINCIPAL_POINT
378     x = rays[..., 0] / rays[..., 2]
379     y = rays[..., 1] / rays[..., 2]
380     u = x * (EFL_MICRONS / PIXEL_SIZE_X_MICRONS) + cx
381     v = y * (EFL_MICRONS / PIXEL_SIZE_Y_MICRONS) + cy
382     return u, v
383
384
385 def _warp_even_field_to_reference(
386     even_img: np.ndarray,
387     output_rows_full: np.ndarray,
```

```
386     R_odd_to_even: np.ndarray,
387 ) -> tuple[np.ndarray, np.ndarray]:
388     H_full, W = SENSOR_SHAPE
389     assert even_img.shape[0] == H_full // 2
390
391     x_full = np.broadcast_to(np.arange(W, dtype=np.float64), (
392         output_rows_full.size, W))
393     y_full = np.broadcast_to(output_rows_full[:, None].astype(np.
394         float64), (output_rows_full.size, W))
395
396     rays_odd = _pixel_rays(x_full, y_full).reshape(-1, 3)
397     rays_even = (R_odd_to_even @ rays_odd.T).T
398
399     u_full, v_full = _project_to_full_pixels(rays_even)
400     v_even_idx = (v_full - 1.0) * 0.5
401
402     x = u_full.ravel()
403     y = v_even_idx.ravel()
404     H_even, W_even = even_img.shape
405
406     valid = (x >= 0) & (x <= W_even - 1) & (y >= 0) & (y <= H_even
407         - 1)
408
409     x0 = np.clip(np.floor(x).astype(np.int64), 0, W_even - 1)
410     y0 = np.clip(np.floor(y).astype(np.int64), 0, H_even - 1)
411     x1 = np.clip(x0 + 1, 0, W_even - 1)
412     y1 = np.clip(y0 + 1, 0, H_even - 1)
413
414     wx = x - x0
415     wy = y - y0
416
417     Ia = even_img[y0, x0]
418     Ib = even_img[y0, x1]
419     Ic = even_img[y1, x0]
```

```
417     Id = even_img[y1, x1]
418
419     sampled = (
420         (1 - wx) * (1 - wy) * Ia
421         + wx * (1 - wy) * Ib
422         + (1 - wx) * wy * Ic
423         + wx * wy * Id
424     )
425     sampled[~valid] = 0.0
426     valid_mask = valid.reshape(output_rows_full.size, W).astype(np.
427         uint8)
428
429
430     def recombine_with_rotation(img: np.ndarray, omega: np.ndarray,
431         delta_t: float) -> tuple[np.ndarray, np.ndarray, np.ndarray]:
432         H, W = img.shape
433         if (H, W) != SENSOR_SHAPE:
434             raise ValueError(f"Unexpected image shape {img.shape}; "
435                 f"expected {SENSOR_SHAPE}")
436
437         odd = img[0::2, :]
438         even = img[1::2, :]
439
440         R_odd_to_even = rodrigues(omega, delta_t)
441
442         odd_rows_full = np.arange(0, H, 2)
443         even_rows_full = np.arange(1, H, 2)
444
445         aligned_even_on_odd, _ = _warp_even_field_to_reference(even,
446             odd_rows_full, R_odd_to_even)
447         aligned_even_on_even, mask_on_even =
448             _warp_even_field_to_reference(even, even_rows_full,
```

```
        R_odd_to_even)

445
446     recombined = np.empty_like(img, dtype=np.float32)
447     recombined[0::2, :] = odd
448     recombined[1::2, :] = aligned_even_on_even
449
450     mask_full = np.empty_like(img, dtype=np.uint8)
451     mask_full[0::2, :] = 1
452     mask_full[1::2, :] = mask_on_even
453
454     return recombined, mask_full, R_odd_to_even
455
456 # --- FFT / notch filtering
457
458 def fft_shift(mag: np.ndarray) -> np.ndarray:
459     return np.fft.fftshift(mag, axes=(0, 1))
460
461
462 def make_symmetric_notches(manual: Sequence[Tuple[int, int]], width
463 : int, height: int) -> list[Tuple[int, int]]:
464     out: list[Tuple[int, int]] = []
465     for x, y in manual:
466         out.append((x, y))
467         out.append(((width - x) % width, (height - y) % height))
468
469     return out
470
471 def apply_gaussian_notch_filter(complex_fft: np.ndarray,
472     notch_centers: Sequence[Tuple[int, int]], radius: float) -> None
473     :
474
475     rows, cols = complex_fft.shape[:2]
476     sigma = radius / 2.0
477     two_sigma2 = 2.0 * sigma * sigma
```

```
474     yy, xx = np.mgrid[0:rows, 0:cols]
475
476     H = np.ones((rows, cols), dtype=np.float32)
477
478     for cx, cy in notch_centers:
479
480         d2 = (xx - cx) ** 2 + (yy - cy) ** 2
481
482         notch = 1.0 - np.exp(-d2 / two_sigma2)
483
484         H *= notch.astype(np.float32)
485
486         complex_fft[:, 0] *= H
487
488         complex_fft[:, 1] *= H
489
490
491
492
493     def fourier_notch_reconstruct(img: np.ndarray, manual_notches:
494         Sequence[Tuple[int, int]], radius: float) -> np.ndarray:
495
496         padded = cv2.copyMakeBorder(
497
498             img,
499             top=0,
500             bottom=cv2.getOptimalDFTSize(img.shape[0]) - img.shape[0],
501             left=0,
502             right=cv2.getOptimalDFTSize(img.shape[1]) - img.shape[1],
503             borderType=cv2.BORDER_CONSTANT,
504             value=0,
505         )
506
507         planes = [padded.astype(np.float32), np.zeros_like(padded,
508             dtype=np.float32)]
509
510         complex_i = cv2.merge(planes)
511
512         cv2.dft(complex_i, complex_i)
513
514         complex_i = fft_shift(complex_i)
515
516
517         h, w = complex_i.shape[:2]
518
519         notches = make_symmetric_notches(manual_notches, w, h)
520
521         apply_gaussian_notch_filter(complex_i, notches, radius)
522
523
524         complex_i = fft_shift(complex_i)
525
526         reconstructed = cv2.idft(complex_i, flags=cv2.DFT_REAL_OUTPUT |
527             cv2.DFT_SCALE)
```

```
505     return reconstructed.astype(np.float32)
506
507
508 # --- Edge tracking + circles
509
510 def _column_gradients(gray: np.ndarray, mask: np.ndarray | None =
511     None) -> np.ndarray:
512     """Return |vertical gradient| per column; masked values set to
513     0."""
514     col_diff = np.abs(np.diff(gray.astype(np.int32), axis=0))
515     if mask is not None:
516         m = mask[1:, :] & mask[:-1, :]
517         col_diff = np.where(m, col_diff, 0)
518     return col_diff
519
520
521 def _find_strongest_edge_in_column(gray: np.ndarray, col: int) ->
522     int:
523     grad = np.abs(np.diff(gray[:, col].astype(np.int32)))
524     if grad.size == 0:
525         return -1
526     idx = int(np.argmax(grad))
527     return idx if grad[idx] > 0 else -1
528
529
530 def _compute_column_threshold(gray: np.ndarray, col: int, k: float =
531     1.0) -> int:
532     grads = np.abs(np.diff(gray[:, col].astype(np.int32)))
533     mean = grads.mean() if grads.size else 0.0
534     std = grads.std() if grads.size else 0.0
535     return int(round(mean + k * std))
```

```
534 def track_edge_andrey(gray: np.ndarray, start_col: int = 0, k:
535     float = 3.0, searchrange: int = 3) -> list[Tuple[int, int]]:
536     """
537         Original C++-style tracker: per column strongest local gradient
538             vs adaptive threshold,
539             zero-order hold if below threshold.
540     """
541     edges: list[Tuple[int, int]] = []
542     current_col = start_col
543     current_row = _find_strongest_edge_in_column(gray, current_col)
544     if current_row == -1:
545         return edges
546     edges.append((current_col, current_row))
547
548     for current_col in range(start_col + 1, gray.shape[1]):
549         best_row = current_row
550         best_val = -1
551         for offset in range(-searchrange, searchrange + 1):
552             r = current_row + offset
553             if r < 0 or r >= gray.shape[0] - 1:
554                 continue
555             grad = abs(int(gray[r + 1, current_col]) - int(gray[r,
556                                                 current_col]))
557             if grad > best_val:
558                 best_val = grad
559                 best_row = r
560             thresh = _compute_column_threshold(gray, current_col, k)
561             if best_val >= thresh:
562                 current_row = best_row
563                 edges.append((current_col, current_row))
564             if current_row > gray.shape[0] - 20:
565                 return edges
566     return edges
```

```
565
566 def fit_circle_ls(points: Sequence[Tuple[float, float]]) -> tuple[
567     float, float, float] | None:
568
569     if len(points) < 3:
570
571         return None
572
573     A = []
574
575     Bv = []
576
577     for x, y in points:
578
579         A.append([x, y, 1.0])
580
581         Bv.append(x * x + y * y)
582
583     A = np.asarray(A, dtype=np.float64)
584
585     Bv = np.asarray(Bv, dtype=np.float64)
586
587     sol, _, _, _ = np.linalg.lstsq(A, Bv, rcond=None)
588
589     a, b, c = sol
590
591     cx = a / 2.0
592
593     cy = b / 2.0
594
595     r = np.sqrt(c + cx * cx + cy * cy)
596
597     return cx, cy, r
598
599
600
601
602
603
604 def fit_circle_hyper(points: Sequence[Tuple[float, float]]) ->
605     tuple[float, float, float] | None:
606
607     n = len(points)
608
609     if n < 3:
610
611         return None
612
613     pts = np.asarray(points, dtype=np.float64)
614
615     mean = pts.mean(axis=0)
616
617     x = pts[:, 0] - mean[0]
618
619     y = pts[:, 1] - mean[1]
620
621     z = x * x + y * y
622
623
624     Mxx = np.mean(x * x)
625
626     Myy = np.mean(y * y)
627
628     Mxy = np.mean(x * y)
```

```
597     Mxz = np.mean(x * z)
598     Myz = np.mean(y * z)
599     Mzz = np.mean(z * z)
600     mean_z = np.mean(z)
601
602     M = np.array(
603         [
604             [Mzz, Mxz, Myz, mean_z],
605             [Mxz, Mxx, Mxy, 0.0],
606             [Myz, Mxy, Myy, 0.0],
607             [mean_z, 0.0, 0.0, 1.0],
608         ]
609     )
610
611     N_inv = np.array(
612         [
613             [0.0, 0.0, 0.0, 0.5],
614             [0.0, 1.0, 0.0, 0.0],
615             [0.0, 0.0, 1.0, 0.0],
616             [0.5, 0.0, 0.0, -2.0 * mean_z],
617         ]
618     )
619
620     P = N_inv @ M
621     eigvals, eigvecs = np.linalg.eig(P)
622     mask_real = np.isclose(eigvals.imag, 0.0, atol=1e-10)
623     eigvals = eigvals.real[mask_real]
624     eigvecs = eigvecs[:, mask_real]
625
626     pos = eigvals[eigvals > 1e-12]
627     if pos.size == 0:
628         return None
629     idx = int(np.argmin(pos))
630     v = eigvecs[:, np.where(eigvals > 1e-12)[0][idx]].real
```

```
631     A, B, C, D = v
632
633     if abs(A) < 1e-7:
634
635         return None
636
637     det = B * B + C * C - 4 * A * D
638
639     if det < 0:
640
641         return None
642
643     cx = -B / (2 * A)
644
645     cy = -C / (2 * A)
646
647     r = np.sqrt(det) / (2 * abs(A))
648
649     return cx + mean[0], cy + mean[1], r
650
651
652
653
654
655
656
657
658
659
660
```

```
641
642
643 # --- Limb profiles
644
645 def get_line_profile(img: np.ndarray, start: Tuple[float, float],
646                      end: Tuple[float, float]) -> list[int]:
647
648     """
649
650     Sample pixel values along the line from start to end (inclusive).
651
652     Uses a simple integer-step interpolation (Bresenham-like) to
653     avoid cv2.LineIterator.
654
655
656
657
658
659
660
```

```
661
662     xs = np.linspace(x0, x1, steps + 1)
663     ys = np.linspace(y0, y1, steps + 1)
664     vals: list[int] = []
665     for x, y in zip(xs, ys):
666         xi, yi = int(round(x)), int(round(y))
667         if 0 <= xi < W and 0 <= yi < H:
668             vals.append(int(img[yi, xi]))
669     return vals
670
671
672 def analyze_limb_profiles(
673     img: np.ndarray,
674     center: Tuple[float, float],
675     radius: float,
676     scan_length: int = 160,
677     margin: int = 20,
678     step_deg: int = 1,
679 ) -> list[dict]:
680     H, W = img.shape
681     cx, cy = center
682     results: list[dict] = []
683     for angle in range(0, 360, step_deg):
684         rad = np.deg2rad(angle)
685         r_inner = radius - margin
686         r_outer = radius + (scan_length - margin)
687         start = (cx + r_inner * np.cos(rad), cy + r_inner * np.sin(
688             rad))
689         end = (cx + r_outer * np.cos(rad), cy + r_outer * np.sin(
690             rad))
691         if not (0 <= start[0] < W and 0 <= start[1] < H and 0 <=
692             end[0] < W and 0 <= end[1] < H):
693             continue
694         profile = get_line_profile(img, start, end)
```

```
692     grads = [abs(profile[i + 1] - profile[i]) for i in range(
693         len(profile) - 1)]
694
695     if grads:
696         max_grad = max(grads)
697         max_pos = grads.index(max_grad)
698     else:
699         max_grad = 0
700         max_pos = 0
701
702     results.append(
703         {
704             "angle_deg": angle,
705             "max_grad": max_grad,
706             "max_grad_pos": max_pos,
707             "profile_len": len(profile),
708         }
709     )
710
711     return results
712
713
714
715 # --- Star tracker data structures
716 -----
717 # --- Star tracker integration (delegated to star_tracker.py)
718 -----
719
720 STAR_MAX_CATALOG_MAG = 8.0
721
722
723 def run_star_tracker_wrapper(
724     image_path: str,
725     catalog_path: str,
726     max_catalog_mag: float = 8.0,
727     output_path: str = "annotated.png",
728     detection_mask: Optional[np.ndarray] = None,
729 ):
```

```
722     """Delegate to star_tracker.run_star_tracker for full
723         functionality."""
724
725     star_tracker.run_star_tracker(
726         image_path=image_path,
727         catalog_path=catalog_path,
728         max_catalog_mag=max_catalog_mag,
729         output_path=output_path,
730         detection_mask=detection_mask,
731     )
732
733
734
735     # --- main
736
737     -----
738
739     def process_folder(
740         science_dir: Path,
741         output_dir: Path,
742         omega: np.ndarray,
743         delta_t_s: float,
744         denoise_config: DenoiseConfig | None = None,
745         lens_maps: tuple[np.ndarray, np.ndarray] | None = None,
746     ) -> List[Path]:
747
748         frames = sorted(science_dir.glob("IMD_*.png"))
749
750         if not frames:
751             print(f"No science frames found in {science_dir}")
752
753             return []
754
755
756         written: List[Path] = []
757
758         for frame in frames:
759             raw = _load_grayscale(frame)
760
761             dtype = raw.dtype
762
763             work = raw.astype(np.float32)
```

```
754     if APPLY_DESPIN:
755
756         recombined, mask_full, R = recombine_with_rotation(work
757             , omega, delta_t_s)
758
759         work = recombined
760
761         work[mask_full == 0] = 0
762
763     else:
764
765         mask_full = np.ones_like(work, dtype=np.uint8)
766
767         R = np.eye(3, dtype=np.float64)
768
769
770     if APPLY_LENS and lens_maps is not None:
771
772         work = _apply_lens_correction(work, lens_maps)
773
774         mask_full = _apply_lens_correction(mask_full, lens_maps
775             , interpolation=cv2.INTER_NEAREST)
776
777         mask_full = (mask_full > 0.5).astype(np.uint8)
778
779         work[mask_full == 0] = 0
780
781
782     if APPLY_RESIDUAL_GLOW:
783
784         work = column_median_subtract(work)
785
786
787     if APPLY_FFT_NOTCH:
788
789         work = fourier_notch_reconstruct(work, MANUAL_NOTCHES,
790             NOTCH_RADIUS)
791
792
793     if APPLY_DENOISE and denoise_config is not None:
794
795         work = clean_image(work, denoise_config)
796
797
798     # Analysis-only steps
799
800     edges: list[Tuple[int, int]] = []
801
802     circle_ls: Optional[Tuple[float, float, float]] = None
803
804     circle_hyper: Optional[Tuple[float, float, float]] = None
805
806     limb_stats: list[dict] = []
807
808
809     if APPLY_EDGE_TRACKING:
810
811         edge_img = work.astype(np.uint8)
```

```
785     H_img = edge_img.shape[0]
786
787     # Base ROI: lower portion + overlap mask
788     roi_start = int(0.5 * H_img)
789
790     edge_mask = np.ones_like(edge_img, dtype=bool)
791     edge_mask[:roi_start, :] = False
792
793     edge_mask[mask_full == 0] = False
794
795     # Brightness gate to avoid tracking pure background
796     valid_pixels = edge_img[edge_mask]
797
798     if valid_pixels.size > 0:
799
800         bright_thresh = np.percentile(valid_pixels, 50.0)
801
802         bright_mask = edge_img >= bright_thresh
803
804         edge_mask &= bright_mask
805
806         edge_img_roi = edge_img.copy()
807
808         edge_img_roi[~edge_mask] = 0
809
810         # Estimate dominant gradient row to focus search band
811
812         row_energy = np.sum(np.abs(np.diff(edge_img_roi.astype(
813             np.int32), axis=0)), axis=1)
814
815         if np.any(row_energy):
816
817             peak_row = int(np.argmax(row_energy))
818
819         else:
820
821             peak_row = int(0.75 * H_img)
822
823             band_half = 60
824
825             band = (max(1, peak_row - band_half), min(H_img - 1,
826                 peak_row + band_half))
827
828             edges = track_edge_andrey(edge_img_roi, start_col=10, k
829             =3.0, searchrange=3)
830
831
832             if APPLY_CIRCLE_FITS and len(edges) >= 8:
833
834                 circle_ls = fit_circle_ls(edges)
835
836                 circle_hyper = fit_circle_hyper(edges)
837
838
839             if APPLY_LIMB_PROFILES and circle_ls:
840
841                 cx, cy, radius = circle_ls
```

```
815     limb_stats = analyze_limb_profiles(work.astype(np.uint8
816                                         ), (cx, cy), radius)
817
818     # Optional visualization of edge + fits
819     if APPLY_EDGE_TRACKING:
820         # Use the same dynamic range as the saved output (no
821         # extra normalization to avoid boosting noise)
822         overlay_base = _to_8bit(work, dtype)
823         overlay = cv2.cvtColor(overlay_base, cv2.COLOR_GRAY2BGR
824                               )
825
826         # Red points: edge tracking samples
827         for x, y in edges:
828             cv2.circle(overlay, (int(x), int(y)), 1, (0, 0,
829                         255), -1, lineType=cv2.LINE_AA)
830
831         # Blue line: LS fit (Kasa)
832         if circle_ls:
833             cx_l, cy_l, r_l = circle_ls
834             cv2.circle(overlay, (int(round(cx_l)), int(round(
835                         cy_l))), int(round(r_l)), (255, 0, 0), 2,
836                         lineType=cv2.LINE_AA)
837
838         # Yellow line: Hyper fit
839         if circle_hyper:
840             cx_h, cy_h, r_h = circle_hyper
841             cv2.circle(overlay, (int(round(cx_h)), int(round(
842                         cy_h))), int(round(r_h)), (0, 255, 255), 2,
843                         lineType=cv2.LINE_AA)
844
845         # Write overlay alongside the main output
846         overlay_path = output_dir / f"{frame.stem}_limb_overlay
847                                     .png"
848
849         cv2.imwrite(str(overlay_path), overlay)
850
851
852         out_path = output_dir / frame.name
853         _save_grayscale(work, out_path, dtype)
854
855         if APPLY_DESPIN:
```

```
840     mask_path = MASK_OUTPUT_DIR / frame.name
841     _save_grayscale(mask_full, mask_path, np.uint8)
842
843     if APPLY_STAR_TRACKER:
844         try:
845             annotated_path = out_path.with_name(f"{out_path.
846                                         stem}_annotated.png")
847             run_star_tracker_wrapper(
848                 image_path=str(out_path),
849                 catalog_path=str(STAR_CATALOG_PATH),
850                 max_catalog_mag=STAR_MAX_CATALOG_MAG,
851                 output_path=str(annotated_path),
852                 detection_mask=mask_full if (APPLY_DESPIN and
853                                               STAR_USE_DESPIN_MASK) else None,
854             )
855             print(f"star tracker annotated -> {annotated_path
856                  }")
857         except Exception as exc:
858             print(f"star tracker failed: {exc}")
859
860         # Report
861         tag = []
862         if APPLY_DESPIN:
863             theta_deg = np.linalg.norm(omega) * delta_t_s * 180.0 /
864             np.pi
865             tag.append(f"dtheta={theta_deg:.6f}deg")
866         if APPLY_LENS:
867             tag.append("lens")
868         if APPLY_DENOISE:
869             tag.append("denoise")
870         if APPLY_FFT_NOTCH:
871             tag.append("fft")
872         if APPLY_RESIDUAL_GLOW:
873             tag.append("colmed")
```

```

870     summary = ",".join(tag) if tag else "raw"
871     print(f"{frame.name} -> {out_path}[{summary}]")
872
873     if circle_ls:
874         cx, cy, r = circle_ls
875         print(f"  circle_ls: cx={cx:.2f}, cy={cy:.2f}, r={r:.2f}")
876
877     if circle_hyper:
878         cx, cy, r = circle_hyper
879         print(f"  circle_Hyper: cx={cx:.2f}, cy={cy:.2f}, r={r:.2f}")
880
881     if limb_stats:
882         best = max(limb_stats, key=lambda d: d["max_grad"])
883         print(f"  limb_max_grad[{best['max_grad']}] at angle {best['angle_deg']} deg")
884
885     written.append(out_path)
886
887     return written
888
889
890
891
892     def main() -> None:
893
894         output_dir = OUTPUT_DIR
895         output_dir.mkdir(parents=True, exist_ok=True)
896         denoise_cfg = DENOISE_SETTINGS if APPLY_DENOISE else None
897         if denoise_cfg and denoise_cfg.apply_dark_frame and denoise_cfg
898             .dark_frame is None and denoise_cfg.dark_frame_path:
899                 denoise_cfg.dark_frame = _load_grayscale(Path(denoise_cfg.
900
901                     dark_frame_path))
902
903
904         omega = np.array(OMEGA, dtype=np.float64)
905
906         if OMEGA_FRAME == "spacecraft":
907
908             omega = np.array(SC_TO_CHUD_MATRIX @ omega, dtype=np.
909
910                 float64)
911
912
913         lens_maps = _build_lens_maps(LENS_KAPPA) if APPLY_LENS else
914             None

```

```
897
898     process_folder(
899         science_dir=SCIENCE_DIR,
900         output_dir=output_dir,
901         omega=omega,
902         delta_t_s=DELTA_T,
903         denoise_config=denoise_cfg,
904         lens_maps=lens_maps,
905     )
906
907
908 if __name__ == "__main__":
909     main()
```

Listing A.1: Main script

A.2.2 Star tracker script

```
1 import numpy as np
2 import cv2
3 from dataclasses import dataclass
4 from typing import List, Tuple, Dict, Optional
5 import math
6 import re
7 import csv
8 from pathlib import Path
9 from denoising import DenoiseConfig, FlattenConfig, clean_image
10
11
12 # -----
13 # Data structures
14 # -----
15
16 @dataclass
17 class DetectedStar:
```

```
18     x: float              # pixel x (column)
19     y: float              # pixel y (row)
20     flux: float           # integrated brightness
21     ray_cam: Optional[np.ndarray] = None    # 3D unit vector in
22                               camera frame
23
24
25 @dataclass
26 class ImageInputConfig:
27     """Configuration for loading the input image with optional
28         rotation/resizing."""
29     path: str
30     rotate_deg: float = 0.0                      # CCW rotation for
31                               processing
32     resize_to: Optional[Tuple[int, int]] = None   # (width, height)
33                               in px
34
35
36 @dataclass
37 class CatalogStar:
38     name: str
39     ra_deg: float
40     dec_deg: float
41     mag: float
42     vec_inertial: np.ndarray
43
44
45 @dataclass
46 class ProcessingOptions:
47     """Processing toggles for star detection."""
48     background_subtraction: bool = False
49     background_kernel: int = 51
```

```
47     normalize: bool = False
48
49     threshold_multiplier: float = 2.5
50
51     morph_open: bool = False
52
53     morph_kernel: int = 3
54
55     morph_iterations: int = 1
56
57     min_area: Optional[int] = 2
58
59     max_area: Optional[int] = 20000
60
61     max_stars: int = 200
62
63     saturated_threshold: int = 225
64
65     allow_oversized_saturated: bool = True
66
67
68
69
70
71 JUNO_HARDWARE_PARAMS = {
72     "efl_um": 20006.0,           # Effective Focal Length
73     "pixel_x_um": 8.6,          # Pixel size X (um)
74     "pixel_y_um": 8.3,          # Pixel size Y (um)
75     "res_x": 752.0,             # Native width
76     "res_y": 580.0,             # Native height
77     "cx.hardware": 383.0,        # Native optical center X
78     "cy.hardware": 257.0,        # Native optical center Y
79     "kappa": 3.3e-8,            # Distortion coefficient (small but
                                included)
```

```

80 }
81
82 # JUNO ASC (CHU-D) lens parameters
83 LENS_JUNO_ASC = {
84     # derived focal lengths in pixels
85     "fx": 20006.0 / 8.6,      # ~2326.28 px
86     "fy": 20006.0 / 8.3,      # ~2410.36 px
87     "cx": 383.0,
88     "cy": 257.0,
89     "kappa": 3.3e-8,          # single-term radial distortion
90     "W": 752,
91     "H": 580,
92 }
93
94 # -----
95 # Star-tracker runconfig (adjust for your frame/catalog)
96 # -----
97 STAR_IMAGE_PATH = Path("JUNO_Despin_Images") / "Despinned" / "
98     Denoised"
99 STAR_IMAGE_FALLBACK = Path("JUNO_Despin_Images") / "Despinned" / "
100     IMD_621530842.422424.png"
101 # Use a Gaia/HIP CSV as the primary catalog for labeling.
102 STAR_CATALOG_PATH = Path("/Users/vojtadeconinck/Downloads/
103     ImageAnalysis30300/combined_stars.csv")
104 STAR_OUTPUT_PATH = Path("annotated.png")
105 STAR_MAX_CATALOG_MAG = 8.0
106 STAR_ROTATE_DEG = 0.0
107 STAR_RESIZE_TO: Optional[Tuple[int, int]] = None
108 STAR_FLIP_CODE: Optional[int] = None    # None (no flip), 0=vertical,
109                 1=horizontal, -1=both
110 STAR_MIRROR_X_IN_RAYS = True           # mirror camera-x in the ray
111                 mapping instead of flipping the image
112 STAR_PROCESSING_MODE = "real"        # fallback preset

```

```
108 STAR_PROCESSING_OPTIONS: Optional[ProcessingOptions] =
109     ProcessingOptions(
110         background_subtraction=True,
111         background_kernel=61,
112         normalize=True,
113         threshold_multiplier=1.4,
114         morph_open=True,
115         morph_kernel=3,
116         morph_iterations=1,
117         min_area=10,
118         max_area=600,
119         max_stars=400,
120         saturated_threshold=220,
121         allow_oversized_saturated=True,
122     )
123 STAR_CAMERA_OVERRIDE: Optional[Dict[str, float]] = None
124 STAR_MASK_BORDERS_PX = (20, 80, 20, 20) # (top, bottom, left,
125                                         right) mask away overlap-loss edges
126 # Debug output settings
127 STAR_SAVE_DEBUG_INTERMEDIATES = True
128 STAR_DEBUG_DIR = Path("JUNO_Despin/Images") / "star_tracker_debug"
129 # Precomputed valid-region masks (odd/even overlap) produced by
130 # juno_despin
131 STAR_VALID_MASK_DIR = Path("JUNO_Despin/Images") / "Despinned" / "
132     Masks"
133 # Jupiter glare mask settings (optional)
134 STAR_SAVE_JUPITER_MASK = True
135 STAR_JUPITER_PERCENTILE = 90.0 # high percentile to isolate glare
136     (lower to include halo)
137 STAR_JUPITER_MIN_AREA = 3000 # minimum glare component area;
138     stars are << this area
139 STAR_JUPITER_MIN_Y_FRAC = 0.25 # only suppress blobs in lower part
140     of image
141 STAR_JUPITER_DILATE = 35 # dilate glare mask to include halo
```

```
135 STAR_JUPITER_PAD_Y = 0           # no rectangular padding, rely on
136   dilation
137 STAR_JUPITER_PAD_X = 0
138 STAR_DISABLE_JUPITER_FALLBACK = True  # skip row-mean fallback if
139   no blob found
140 STAR_JUPITER_EXCLUDE_STAR_MAX_AREA = 400  # components <= this area
141   are treated as stars, not glare
142 # Auto-build overlap mask from the input image instead of fixed
143   borders (used if no precomputed mask)
144 STAR_AUTO_OVERLAP_MASK = True
145 STAR_AUTO_MASK_PERCENTILE = 70.0  # percentile to threshold non-
146   zero region
147 STAR_AUTO_MASK_CLOSE_KERNEL = 9    # odd kernel for closing
148 STAR_AUTO_MASK_ERODE = 2          # pixels to erode after closing (
149   set 0 to skip)
150
151 APPLY_DENOISING = False
152 STAR_DENOISE_SETTINGS = DenoiseConfig(
153     apply_flatten=False,
154     spike_neighborhood=1,
155     spike_threshold=40.0,
156     median_kernel=3,
157     gaussian_kernel=0,
158     gaussian_sigma=0.0,
159 )
160 STAR_ATTITUDE_MIN_INLIERS = 4
161 STAR_ATTITUDE_MAX_ERR_DEG = 0.2
162 STAR_ATTITUDE_MAX_ITER = 3000
163 # Name-lookup table to map HIP/Gaia IDs to readable labels.
164 STAR_NAME_LOOKUP_PATH = Path("gaia_to_names.csv")  # Gaia->HIP
165   mapping source
166 STAR_GAIA_HIP_PATH = Path("2025-12-05-15-20-27-425171.csv")  #
167   default Gaia->HIP crossmatch
```

```
160 # BSC TSV is only used as a positional fallback when the lookup is
161 # missing.
162 STAR_BSC_PATH = Path("bright_star_catalog.tsv") # default Bright Star Catalog
163 STAR_IAU_PATH = Path("IAU-CatalogueofStarNames(alwaysuptodate).csv") # optional IAU names
164 STAR_HIP_TO_HD_PATH: Optional[Path] = None # set if you have a HIP->HD mapping CSV
165 STAR_GAIA_POS_PATH = Path("gaia_bright_stars.csv") # optional Gaia positions to help name matching
166
167 def _rotate_image_keep_size(image: np.ndarray, angle_deg: float) -> np.ndarray:
168     """Rotate around image center while preserving dimensions."""
169     if abs(angle_deg) < 1e-6:
170         return image
171     h, w = image.shape[:2]
172     M = cv2.getRotationMatrix2D((w / 2.0, h / 2.0), angle_deg, 1.0)
173     return cv2.warpAffine(
174         image,
175         M,
176         (w, h),
177         flags=cv2.INTER_LINEAR,
178         borderMode=cv2.BORDER_CONSTANT,
179         borderValue=0,
180     )
181
182 def _hms_to_deg(hms: str) -> float:
183     """Convert 'HH:MM:SS.SSS' or compact 'HHMMSS.S' to RA in degrees."""
184     hms = hms.strip()
185     if ":" in hms:
186         h, m, s = hms.split(":")
```

```

187         h = float(h); m = float(m); s = float(s)
188     else:
189         # compact notation 'HHMMSS.S'
190         if len(hms) < 6:
191             raise ValueError(f"RA string too short:{hms!r}")
192         h = float(hms[0:2])
193         m = float(hms[2:4])
194         s = float(hms[4:])
195     return 15.0 * (h + m/60.0 + s/3600.0)
196
197
198 def _dms_to_deg(dms: str) -> float:
199     """Convert '+DD:MM:SS.SS' or '+DDMMSS' style strings to Dec in
200     degrees."""
201     dms = dms.strip()
202     if ":" in dms:
203         sign = -1.0 if dms.startswith("-") else 1.0
204         body = dms.lstrip("+-")
205         d, m, s = body.split(":")
206         d = float(d); m = float(m); s = float(s)
207     else:
208         sign = -1.0 if dms[0] == "-" else 1.0
209         body = dms.lstrip("+-")
210         if len(body) < 4:
211             raise ValueError(f"Dec string too short:{dms!r}")
212         d = float(body[0:2])
213         m = float(body[2:4])
214         s = float(body[4:]) if len(body) > 4 else 0.0
215     return sign * (d + m/60.0 + s/3600.0)
216
217 def _to_uint8_debug(img: np.ndarray) -> np.ndarray:
218     """Scale various arrays to uint8 for debug dumps."""
219     if img.dtype == np.uint8:

```

```

220     return img
221
222     if img.dtype == np.bool_:
223         return (img.astype(np.uint8) * 255)
224
225     if np.issubdtype(img.dtype, np.integer):
226         max_val = np.iinfo(img.dtype).max
227
228         return np.clip(img.astype(np.float32) * (255.0 / max_val),
229                     0, 255).astype(np.uint8)
230
231     lo = float(np.percentile(img, 1.0))
232
233     hi = float(np.percentile(img, 99.0))
234
235     if hi <= lo:
236
237         hi = lo + 1e-6
238
239     norm = np.clip((img - lo) / (hi - lo), 0.0, 1.0)
240
241     return np.clip(norm * 255.0, 0, 255).astype(np.uint8)

242
243
244
245
246
247
248
249

def _make_viz_binary(img_u8: np.ndarray) -> np.ndarray:
    """Create a more visible binary for sparse masks (dilate and
    invert)."""
    if img_u8.dtype != np.uint8:
        img_u8 = _to_uint8_debug(img_u8)
    if img_u8.max() == 0:
        return img_u8
    # If mostly 0/255, dilate to make sparse points visible and
    # invert for contrast.
    unique_vals = np.unique(img_u8)
    if unique_vals.size <= 3:
        k = np.ones((3, 3), np.uint8)
        dil = cv2.dilate(img_u8, k, iterations=1)
        inv = 255 - dil
        return inv
    return img_u8

```

```
250 def _make_overlay(mask_u8: np.ndarray, base_gray_u8: np.ndarray) ->
251     np.ndarray:
252     """Overlay binary mask as red contours on the grayscale base
253     image."""
254     if mask_u8.dtype != np.uint8:
255         mask_u8 = _to_uint8_debug(mask_u8)
256     if base_gray_u8.dtype != np.uint8:
257         base_gray_u8 = _to_uint8_debug(base_gray_u8)
258     color = cv2.cvtColor(base_gray_u8, cv2.COLOR_GRAY2BGR)
259     if mask_u8.max() == 0:
260         return color
261     contours, _ = cv2.findContours(mask_u8, cv2.RETR_EXTERNAL, cv2.
262                                     CHAIN_APPROX_SIMPLE)
263     cv2.drawContours(color, contours, -1, (0, 0, 255), 1)
264     return color
265
266
267
268 def _quadrilateral_mask_from_image(img_gray: np.ndarray,
269                                     thresh_percentile: float = 0.5,
270                                     close_kernel: int = 9,
271                                     erode_px: int = 2) -> np.ndarray
272                                     :
273     """Build a convex/quad mask from the truly non-zero region of
274     the image."""
275     positive = img_gray[img_gray > 0]
276     if positive.size == 0:
277         return np.zeros_like(img_gray, dtype=np.uint8)
278     thr_val = float(np.percentile(positive, thresh_percentile))
279     thr_val = max(1.0, thr_val)
280     mask = (img_gray > thr_val).astype(np.uint8)
281     if close_kernel and close_kernel > 1:
282         k = close_kernel if close_kernel % 2 == 1 else close_kernel
283         + 1
284         kernel = np.ones((k, k), np.uint8)
```

```
278     mask = cv2.morphologyEx(mask, cv2.MORPH_CLOSE, kernel,
279                             iterations=1)
280
281     if erode_px and erode_px > 0:
282         k = max(1, int(erode_px))
283         kernel = np.ones((k, k), np.uint8)
284         mask = cv2.erode(mask, kernel, iterations=1)
285
286     # Convex hull of non-zero pixels
287     ys, xs = np.nonzero(mask)
288
289     if ys.size == 0:
290
291         return np.zeros_like(mask, dtype=np.uint8)
292
293     pts = np.column_stack((xs, ys)).astype(np.int32)
294     hull = cv2.convexHull(pts)
295
296     # Try to approximate to 4 points if possible
297     approx = cv2.approxPolyDP(hull, 0.01 * cv2.arcLength(hull, True),
298                               True)
299
300     poly = approx if len(approx) >= 3 else hull
301
302     quad_mask = np.zeros_like(mask, dtype=np.uint8)
303     cv2.fillPoly(quad_mask, [poly], 1)
304
305     return quad_mask
306
307
308
309     def _write_debug_images(prefix: str, images: Dict[str, np.ndarray],
310                           out_dir: Optional[Path] = None) -> None:
311
312         """Write intermediate debug images (uint8 scaled) to
313             STAR_DEBUG_DIR or a custom subfolder."""
314
315         if not images:
316
317             return
318
319         out_dir = out_dir or STAR_DEBUG_DIR
320         out_dir.mkdir(parents=True, exist_ok=True)
321         base_gray = images.get("step_input_gray")
322         base_gray_u8 = _to_uint8_debug(base_gray) if base_gray is not
323             None else None
```

```
307     for name, arr in images.items():
308
309         try:
310
311             out = _to_uint8_debug(arr)
312
313         except Exception:
314
315             continue
316
317             cv2.imwrite(str(out_dir / f"{prefix}_{name}.png"), out)
318
319             # Also write a visibility-enhanced version for sparse
320
321             binaries
322
323             viz = _make_viz_binary(out)
324
325             if viz is not out:
326
327                 cv2.imwrite(str(out_dir / f"{prefix}_{name}_viz.png"),
328
329                             viz)
330
331             # Overlay on base gray if available and this looks like a
332
333             binary/sparse mask
334
335             if base_gray_u8 is not None and out.max() > 0 and np.unique
336
337                 (out).size <= 10:
338
339                 overlay = _make_overlay(out, base_gray_u8)
340
341                 cv2.imwrite(str(out_dir / f"{prefix}_{name}_overlay.png"
342
343                             ), overlay)
344
345
346
347
348 def _load_precomputed_mask(image_path: Path) -> Optional[np.ndarray]
349 ]:
350
351     """
352
353     Load a precomputed valid-region mask (odd/even overlap)
354
355         matching the image filename.
356
357     Returns a uint8 mask with 1s where valid, or None if not found
358
359         or unreadable.
360
361     """
362
363     mask_path = STAR_VALID_MASK_DIR / image_path.name
364
365     if not mask_path.exists():
366
367         return None
368
369     mask = cv2.imread(str(mask_path), cv2.IMREAD_GRAYSCALE)
370
371     if mask is None:
```

```
333     return None
334
335
336
337 def _compute_jupiter_mask(
338     img_gray: np.ndarray,
339     percentile: float = STAR_JUPITER_PERCENTILE,
340     min_area: int = STAR_JUPITER_MIN_AREA,
341     min_y_frac: float = STAR_JUPITER_MIN_Y_FRAC,
342     dilate_px: int = STAR_JUPITER_DILATE,
343     pad_y: int = STAR_JUPITER_PAD_Y,
344     pad_x: int = STAR_JUPITER_PAD_X,
345 ) -> Optional[np.ndarray]:
346     """
347     Detect bright Jupiter/glare blobs and return a uint8 mask (1 =
348         glare).
349
350     - Blur first to merge the limb/halo.
351     - Threshold on a high percentile.
352     - Keep large components low in the image (Jupiter region),
353         union them, pad and dilate.
354
355     Returns None if nothing is found.
356
357     """
358
359     # Light blur to merge halo structures
360     img.blur = cv2.GaussianBlur(img_gray, (9, 9), 0)
361
362     positive = img.blur[img.blur > 0]
363     if positive.size == 0:
364         return None
365     thr = float(np.percentile(positive, percentile))
366     if thr <= 0:
367         return None
368     bright = (img.blur >= thr).astype(np.uint8)
369     num, labels, stats, centroids = cv2.
370         connectedComponentsWithStats(bright, connectivity=8)
```

```
364     H, W = img_gray.shape
365
366     mask = np.zeros_like(img_gray, dtype=np.uint8)
367
368     if num > 1:
369
370         for idx in range(1, num):
371
372             area = stats[idx, cv2.CC_STAT_AREA]
373
374             if area < min_area:
375
376                 continue
377
378             if area <= STAR_JUPITER_EXCLUDE_STAR_MAX_AREA:
379
380                 continue # likely a star or small blob
381
382             y_c = centroids[idx][1]
383
384             if y_c < min_y_frac * H:
385
386                 continue
387
388             # Use the actual component shape (not just its bounding
389             # box)
390
391             mask = np.logical_or(mask, labels == idx).astype(np.
392                                         uint8)
393
394
395     # Fallback: if no blob matched, optionally try a row-mean band
396     # (disabled if flag set)
397
398     if mask.max() == 0 and not STAR_DISABLE_JUPITER_FALLBACK:
399
400         row_mean = img_blur.mean(axis=1)
401
402         m = row_mean.mean()
403
404         s = row_mean.std()
405
406         cutoff = max(m + s, np.percentile(row_mean, 90))
407
408         rows = np.where(row_mean >= cutoff)[0]
409
410         if rows.size > 0:
411
412             y0 = int(rows.min())
413
414             y1 = int(rows.max())
415
416             if y1 > H * 0.4: # ensure it's towards the bottom
417
418                 mask[y0:y1 + 1, :] = 1
419
420
421     if mask.max() == 0:
422
423         return None
424
```

```
395     # dilate to catch halo
396
397     if dilate_px and dilate_px > 1:
398
399         k = dilate_px if dilate_px % 2 == 1 else dilate_px + 1
400
401         kernel = np.ones((k, k), np.uint8)
402
403         mask = cv2.dilate(mask, kernel, iterations=1)
404
405     return mask
406
407 def load_yale_bsc_ascii(path: str, max_mag: float = 6.5):
408
409     """
410
411     Parse the ASCII 'ybsc5' (Yale Bright Star Catalog 5) file.
412
413
414     The parts[6] field looks like:
415
416     '000001.1+444022000509.9+451345114.44-16.88'
417
418     RA1950 Dec1950 RA2000 Dec2000    l      b
419
420
421     Parsed sequentially instead of using hard-coded indices.
422
423
424     catalog = []
425
426     with open(path, "r", encoding="ascii", errors="ignore") as f:
427
428         for line in f:
429
430             line = line.rstrip("\n")
431
432             if not line.strip():
433
434                 continue
435
436
437             # Skip headers/comment lines
438
439             if line.startswith("sbsc5") or "BSC_number" in line or
440
441                 "BSC_num" in line:
442
443                 continue
444
445             if not line.strip()[0].isdigit():
446
447                 continue
448
449
450             parts = line.split()
451
452             if len(parts) < 8:
453
454                 continue
```

```
428     # Catalog number
429
430     try:
431         num = int(parts[0])
432     except ValueError:
433         continue
434
435     ra_block = parts[6]
436
437     mag_str = parts[7]
438
439     # Minimum length check
440     if len(ra_block) < 30:
441         continue
442
443     # Sequential parsing
444     i = 0
445
446     ra1950_str = ra_block[i:i+7]; i += 7      #
447           '000001.1'
448
449     dec1950_str = ra_block[i:i+7]; i += 7      #
450           '+444022'
451
452     ra2000_str = ra_block[i:i+7]; i += 7      #
453           '000509.9'
454
455     dec2000_str = ra_block[i:i+7]; i += 7      #
456           '+451345'
457
458     # Ignore remaining galactic l,b fields
459
460
461     # Magnitude
462
463     try:
464         mag = float(mag_str)
465     except ValueError:
466         continue
467
468     if mag > max_mag:
469         continue
470
471     # Convert RA/Dec to degrees
```

```
458     try:
459
460         ra_deg    = _hms_to_deg(ra2000_str)
461         dec_deg  = _dms_to_deg(dec2000_str)
462
463     except ValueError:
464
465         # Skip malformed record
466
467         continue
468
469
470     vec = _radec_to_vec(ra_deg, dec_deg)
471
472
473     name = f"BSC{num:04d}"
474
475     catalog.append(
476         CatalogStar(
477             name=name,
478             ra_deg=ra_deg,
479             dec_deg=dec_deg,
480             mag=mag,
481             vec_inertial=vec,
482         )
483     )
484
485
486     if not catalog:
487
488         raise RuntimeError("No stars parsed from ybsc5; the format"
489                           "may be unexpected.")
490
491     catalog.sort(key=lambda c: c.mag)
492
493     print(f"Loaded {len(catalog)} stars from {path}")
494
495     return catalog
496
497
498 def load_gaia_csv(path: str, max_mag: float = 6.5) -> List[
499     CatalogStar]:
500
501     """
502
503     Read Gaia/Hipparcos CSV and tolerate mixed-case column names.
504
505     """
506
507     catalog: List[CatalogStar] = []
508
509     with open(path, "r", encoding="utf-8", newline="") as f:
```

```
490     reader = csv.DictReader(f)
491
492     # Force lowercase headers (ra, dec, mag) for robustness
493
494     if reader.fieldnames:
495
496         reader.fieldnames = [name.lower() for name in reader.
497                             fieldnames]
498
499
500     for row in reader:
501
502         try:
503
504             ra_deg = float(row["ra"])
505             dec_deg = float(row["dec"])
506             mag = float(row["mag"])
507
508         except (KeyError, ValueError):
509
510             continue
511
512         if mag > max_mag:
513
514             continue
515
516
517         name = row.get("source_id", "Unknown")
518         vec = _radec_to_vec(ra_deg, dec_deg)
519
520
521         catalog.append(
522             CatalogStar(
523                 name=str(name),
524                 ra_deg=ra_deg,
525                 dec_deg=dec_deg,
526                 mag=mag,
527                 vec_inertial=vec,
528             )
529         )
530
531
532     if not catalog:
533
534         raise RuntimeError("No stars parsed. Check the CSV file.")
535
```

```
523     catalog.sort(key=lambda c: c.mag)
524
525     print(f"Loaded {len(catalog)} stars from {path}")
526
527
528     def _parse_hip_from_name(name: str) -> Optional[int]:
529         """Extract HIP number from strings like 'HIP32349'."""
530         m = re.match(r"hip\s*(\d+)", name.strip(), flags=re.IGNORECASE)
531         if not m:
532             return None
533
534         try:
535             return int(m.group(1))
536         except ValueError:
537             return None
538
539     def _parse_gaia_from_name(name: str) -> Optional[int]:
540         """Extract Gaia source_id if the name is a pure digit string.
541
542         """
543         s = name.strip()
544         if not s.isdigit():
545             return None
546
547         try:
548             return int(s)
549         except ValueError:
550             return None
551
552     def load_gaia_to_hip_map(path: Path) -> Dict[int, int]:
553         """Load Gaia->HIP mapping from gaia_to_names.csv (HIP column).
554
555         """
556
557         if not path.exists():
558             print(f"INFO: Gaia->HIP lookup file not found at {path}, skipping mapping.")
```

```
554     return {}
555
556     gaia_to_hip: Dict[int, int] = {}
557
558     with open(path, "r", encoding="utf-8", newline="") as f:
559         reader = csv.DictReader(f)
560
561         for row in reader:
562
563             hip_raw = row.get("HIP") or row.get("hip")
564             gaia_raw = row.get("gaia_id") or row.get("GAIA_ID") or
565                         row.get("gaia")
566
567             if not gaia_raw or not hip_raw:
568
569                 continue
570
571             try:
572
573                 gaia_val = int(float(gaia_raw))
574                 hip_val = int(float(hip_raw))
575
576             except ValueError:
577
578                 continue
579
580             gaia_to_hip[gaia_val] = hip_val
581
582     print(f"Loaded {len(gaia_to_hip)} Gaia->HIP mappings from {path}")
583
584     return gaia_to_hip
585
586
587
588
589 def load_iau_names(path: Path) -> Dict[int, str]:
590     """
591
592     Load HIP->IAU proper name map from the IAU catalog CSV if
593         available.
594
595     Prefer the Simbad spelling column when present.
596
597     """
598
599     if not path.exists():
600
601         return {}
602
603     hip_names: Dict[int, str] = {}
604
605     with open(path, "r", encoding="utf-8", newline="") as f:
606
607         reader = csv.DictReader(filter(lambda ln: ln.strip(), f))
608
609         for row in reader:
```

```
585     # Strip simple HTML span wrappers and normalize keys
586     cleaned = {}
587     for k, v in row.items():
588         if k is None:
589             continue
590         key = re.sub(r"<[^>]+>", "", k).strip()
591         cleaned[key.lower()] = v
592
593     hip_raw = (
594         cleaned.get("hip")
595         or cleaned.get("hip_id")
596         or row.get("HIP")
597         or row.get("hip")
598         or row.get("Hip")
599     )
600     if not hip_raw:
601         continue
602
603     name = (
604         cleaned.get("simbad_spelling")
605         or cleaned.get("proper_names")
606         or row.get("Name")
607         or row.get("name")
608     )
609     if not name:
610         continue
611
612     try:
613         hip_val = int(float(hip_raw))
614     except ValueError:
615         continue
616     hip_names[hip_val] = name.strip()
617 print(f"Loaded {len(hip_names)} IAU HIP names from {path}")
618 return hip_names
```

```

619
620
621 def _parse_radec_from_bsc(row: Dict[str, str]) -> Optional[Tuple[
622     float, float]]:
623     ra = row.get("RAJ2000", "")
624     dec = row.get("DEJ2000", "")
625     try:
626         h, m, s = [float(x) for x in ra.replace(":", " ").split()]
627         ra_deg = 15.0 * (h + m / 60.0 + s / 3600.0)
628         sign = -1.0 if dec.strip().startswith("-") else 1.0
629         d, dm, ds = [float(x) for x in dec.replace(":", " ").split(
630             ()]
631         dec_deg = sign * (abs(d) + dm / 60.0 + ds / 3600.0)
632     return ra_deg, dec_deg
633 except Exception:
634     return None
635
636
637 def load_bsc_for_positions(path: Path) -> Optional[List[Tuple[float,
638     float, str]]]:
639     """
640     Light-weight BSC loader for positional fallback: returns list
641     of (ra_deg, dec_deg, name).
642     """
643     if not path.exists():
644         return None
645     rows = []
646     with open(path, "r", encoding="utf-8", newline="") as f:
647         reader = csv.DictReader(filter(lambda ln: not ln.startswith(
648             "#") and ln.strip(), f), delimiter="|")
649         for row in reader:
650             name = (row.get("Name") or "").strip()
651             if not name:
652                 continue

```

```

648         radec = _parse_radec_from_bsc(row)
649
650         if radec is None:
651             continue
652
653         ra_deg, dec_deg = radec
654
655         rows.append((ra_deg, dec_deg, name))
656
657     print(f"Loaded {len(rows)} BSC entries with positions for"
658           "fallback naming")
659
660     return rows
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677

```

```
678     assigned = False
679
680     if hip_val is not None and iau_name_map:
681
682         new_name = iau_name_map.get(hip_val)
683
684         if new_name:
685
686             if new_name != star.name:
687
688                 replaced += 1
689
690                 star.name = new_name
691
692                 assigned = True
693
694
695 # -----
696 # Star detection
697 # -----
698 def detect_stars(image_gray: np.ndarray, options: ProcessingOptions,
699 , mask: Optional[np.ndarray] = None,
700
701         debug_images: Optional[Dict[str, np.ndarray]] =
702
703             None) -> List[DetectedStar]:
704
705         """
706
707         Hybrid detector:
708
709         combines saturated-blob detection for very bright stars with
710
711         background-subtracted detection for faint stars.
712
713         """
714
715         if image_gray.ndim != 2:
716
717             raise ValueError("detect_stars expects a single-channel ("
718
719                 "grayscale) image.")
720
721
722         mask_bool = None
```

```
709     if mask is not None:
710
711         if mask.shape != image_gray.shape:
712
713             raise ValueError("Mask shape must match image .")
714
715         mask_bool = mask.astype(bool)
716
717         if debug_images is not None:
718
719             debug_images["step_mask"] = mask_bool.astype(np.uint8)
720
721             * 255
722
723         if debug_images is not None:
724
725             debug_images["step_input_gray"] = image_gray.copy()
726
727
728         # -----
729
730         # Path 1: saturated stars on the raw image (e.g., Antares)
731
732         # Anything above the saturation threshold in the original image
733
734             counts as a star;
735
736         # skip background subtraction to preserve bright cores.
737
738         # -----
739
740         sat_thresh = int(np.clip(options.saturated_threshold, 0, 255))
741
742         _, binary_saturated = cv2.threshold(image_gray, sat_thresh,
743
744             255, cv2.THRESH_BINARY)
745
746         if mask_bool is not None:
747
748             binary_saturated = cv2.bitwise_and(binary_saturated,
749
750                 mask_bool.astype(np.uint8))
751
752
753         # -----
754
755         # Path 2: faint-star detection using background subtraction
756
757         # -----
758
759         work = image_gray.astype(np.float32)
760
761         if mask_bool is not None:
762
763             work = np.where(mask_bool, work, 0.0)
764
765         if debug_images is not None:
766
767             debug_images["step_masked_gray"] = work.copy()
768
769
770         # Force background subtraction for the faint path unless
771
772             explicitly disabled
```

```
738     if options.background_subtraction:
739         # Ensure the kernel is large enough (minimum 51) and odd
740         k = max(options.background_kernel, 51)
741         if k % 2 == 0:
742             k += 1
743
744         bg = cv2.medianBlur(image_gray, k)
745         work = work - bg.astype(np.float32)
746
747         if debug_images is not None:
748             debug_images["step_background"] = bg
749             debug_images["step_after_bg_sub"] = work.copy()
750
751         if options.normalize:
752             work -= work.min()
753             max_val = work.max()
754             if max_val > 0:
755                 work /= max_val
756
757             if debug_images is not None:
758                 debug_images["step_after_norm"] = work.copy()
759
760             # Threshold selection for faint detections
761             mean, std = cv2.meanStdDev(work)
762             mean_val = float(mean[0][0])
763             std_val = float(std[0][0])
764             thresh_val = mean_val + options.threshold_multiplier * std_val
765             if debug_images is not None:
766                 debug_images["step_stats_mean_std"] = np.array([[mean_val,
767                                         std_val]], dtype=np.float32)
768
769             if options.normalize:
770                 _, binary_faint = cv2.threshold(work, thresh_val, 255, cv2.
771                                             THRESH_BINARY)
772
773         else:
```

```
769     _, binary_faint = cv2.threshold(work, thresh_val, 255, cv2.
770                                         THRESH_BINARY)
771
771     binary_faint = binary_faint.astype(np.uint8)
772
772     if mask_bool is not None:
773
773         binary_faint = cv2.bitwise_and(binary_faint, mask_bool.
774                                         astype(np.uint8))
774
774     if debug_images is not None:
775
775         debug_images["step_binary_faint"] = binary_faint.copy()
776
776
777     # -----
778
778     # Step 3: merge saturated + faint detections (bitwise OR)
779
779     # -----
780
780     final_binary = cv2.bitwise_or(binary_saturated, binary_faint)
781
781     if mask_bool is not None:
782
782         final_binary = cv2.bitwise_and(final_binary, mask_bool.
783                                         astype(np.uint8))
783
783     if debug_images is not None:
784
784         debug_images["step_binary_saturated"] = binary_saturated.
785                                         copy()
785
785     debug_images["step_binary_final"] = final_binary.copy()
786
786
787     # Clean up small noise
788
788     if options.morph_open:
789
789         kernel_size = max(1, options.morph_kernel)
790
790         kernel = np.ones((kernel_size, kernel_size), np.uint8)
791
791         iterations = max(1, options.morph_iterations)
792
792         final_binary = cv2.morphologyEx(final_binary, cv2.
793                                         MORPH_OPEN, kernel, iterations=iterations)
793
793     if debug_images is not None:
794
794         debug_images["step_binary_final_open"] = final_binary.
795                                         copy()
795
795     # -----
```

```
797     # Step 4: analyze blobs
798     #
799     num_labels, labels, stats, centroids = cv2.
800         connectedComponentsWithStats(final_binary)
800
801     H, W = image_gray.shape
802     detections: List[DetectedStar] = []
803
804     # Safety: if max_area is unset, allow very large blobs for
805     # bright stars
805     safe_max_area = options.max_area if options.max_area is not
806         None else 100000
806
807     for label in range(1, num_labels):
808         x, y, w, h, area = stats[label]
809
810         if options.min_area is not None and area < options.min_area
811             :
812                 continue
812
813         cx, cy = centroids[label]
814
815         # Measure flux on the original image (not the background-
816             subtracted one)
816         r = int(math.sqrt(area)/2) + 2
817         r = min(r, 40) # cap radius
818
819         x0 = max(int(cx) - r, 0)
820         x1 = min(int(cx) + r + 1, W)
821         y0 = max(int(cy) - r, 0)
822         y1 = min(int(cy) + r + 1, H)
823
824         patch = image_gray[y0:y1, x0:x1]
825         if patch.size == 0:
```

```
826         continue
827
828     patch_max = int(patch.max()) if patch.size else 0
829     is_saturated_blob = patch_max >= sat_thresh
830     if area > safe_max_area and not (options.
831         allow_oversized_saturated and is_saturated_blob):
832         continue
833
834     # Weighted centroid
835     weights = patch.astype(float)
836     s = weights.sum()
837
838     if s <= 0:
839         det = DetectedStar(x=cx, y=cy, flux=float(area))
840     else:
841         yy, xx = np.mgrid[y0:y1, x0:x1]
842         cx_refined = (xx * weights).sum() / s
843         cy_refined = (yy * weights).sum() / s
844         det = DetectedStar(x=cx_refined, y=cy_refined, flux=s)
845
846     detections.append(det)
847
848     detections.sort(key=lambda d: d.flux, reverse=True)
849     if options.max_stars and len(detections) > options.max_stars:
850         detections = detections[:options.max_stars]
851
852     print(f"DEBUG: Found {len(detections)} stars. Brightest flux: {detections[0].flux} if detections else 0")
853
854     return detections
855
856     # -----
857     # Processing presets / toggles
858     # -----
```

```
858 def processing_preset(mode: str) -> ProcessingOptions:
859     """Set default options for simulation (low noise) versus real
860     data."""
861     mode = (mode or "real").lower()
862     if mode == "simulation":
863         return ProcessingOptions(
864             background_subtraction=False,
865             background_kernel=31,
866             normalize=True,
867             threshold_multiplier=4.0,
868             morph_open=False,
869             min_area=1,
870             max_area=800,
871             max_stars=150,
872         )
873     # default: noisier real data
874     return ProcessingOptions()
875
876 # -----
877 # Pixels -> unit vectors
878 # -----
879 def pixels_to_unit_vectors(detections: List[DetectedStar],
880                             image_shape: Tuple[int, int],
881                             hardware_params: Optional[Dict[str,
882                               float]] = None) -> None:
883     """
884     Convert pixel detections into unit vectors using the Juno
885     camera parameters,
886     correcting for any resizing applied to the screenshot.
887     hardware_params can override res_x/res_y/cx/cy/efl/pixel pitch.
888     """
889     H_img, W_img = image_shape
```

```
889     params = dict(JUNO_HARDWARE_PARAMS)
890
891     if hardware_params:
892
893         params.update(hardware_params)
894
895         # 1. Compute scale factor between screenshot and physical
896         # sensor
897
898         # Assume the aspect ratio is roughly preserved.
899         scale_x = W_img / params["res_x"]
900         scale_y = H_img / params["res_y"]
901
902         # Warn if aspect ratio is strongly distorted (stretched
903         # screenshot)
904
905         if abs(scale_x - scale_y) > 0.05:
906
907             print(f"WARNING: Aspect ratio looks off! X-scale={scale_x
908                   :.2f}, Y-scale={scale_y:.2f}")
909
910         # 2. Compute hardware fx/fy
911
912         fx_hard = params["efl_um"] / params["pixel_x_um"] # ~2326 px
913         fy_hard = params["efl_um"] / params["pixel_y_um"] # ~2410 px
914
915         # 3. Scale to the screenshot dimensions
916
917         fx = fx_hard * scale_x
918         fy = fy_hard * scale_y
919
920         cx = params["cx.hardware"] * scale_x
921         cy = params["cy.hardware"] * scale_y
922
923         kappa = params["kappa"] # Distortion coefficient
924
925
926         print(f"DEBUG: Juno hardware scale correction: {scale_x:.3f}x")
927         print(f"Using fx={fx:.1f} (was {fx_hard:.1f}), cx={cx:.1
928               f} (was {params['cx.hardware']}"))
929
930
931         for det in detections:
932
933             # Pinhole model (with center correction and non-square
934             # pixels)
```

```
918     # Image: x right, y down. Camera frame: y often up, z
919     # forward.
920
920     # Step A: center and normalize with focal length
921     x_norm = (det.x - cx) / fx
922     y_norm = -(det.y - cy) / fy # invert Y to map image to
923     # camera frame
924     if STAR_MIRROR_X_IN_RAYS:
925         x_norm *= -1.0 # mirror camera-x so geometry matches
926         # without flipping the image
927
928     # Step B: distortion correction (inverse approximation for
929     # small kappa)
930     r2 = x_norm**2 + y_norm**2
931     if kappa != 0.0:
932         factor = 1.0 - kappa * r2
933         x_norm *= factor
934         y_norm *= factor
935
936
937     z = 1.0
938     v = np.array([x_norm, y_norm, z], dtype=np.float64)
939     v /= np.linalg.norm(v)
940
941
942     def build_catalog_pair_table(catalog: List[CatalogStar],
943                                 max_pairs: int = 5000,
944                                 angle_bin_deg: float = 0.01):
945         """
946         Compute pairwise angles between the brightest catalog stars.
947     
```

```

948     Returns:
949
950         bins: dict bin_index -> list of (idx1, idx2, angle_rad)
951         angle_bin_rad: bin size in radians
952
953         """
954
955         # choose N such that N(N-1)/2 ~ max_pairs
956
957         N = min(len(catalog), int(0.5 * (1 + math.sqrt(1 + 8 *
958             max_pairs))))
959
960         N = max(3, N)
961
962         angle_bin_rad = math.radians(angle_bin_deg)
963
964         bins: Dict[int, List[Tuple[int, int, float]]] = {}
965
966         for i in range(N - 1):
967
968             vi = catalog[i].vec_inertial
969
970             for j in range(i + 1, N):
971
972                 vj = catalog[j].vec_inertial
973
974                 cosang = float(np.clip(np.dot(vi, vj), -1.0, 1.0))
975
976                 ang = math.acos(cosang)
977
978                 b = int(ang / angle_bin_rad)
979
980                 bins.setdefault(b, []).append((i, j, ang))
981
982
983             return bins, angle_bin_rad
984
985
986
987
988     # -----
989
990     # TRIAD: snelle attitude uit 2 vectorparen
991
992     # -----
993
994
995
996
997
998     def triad_attitude(v_cam: np.ndarray, w_cam: np.ndarray,
999
1000             v_inertial: np.ndarray, w_inertial: np.ndarray)
1001
1002             -> np.ndarray:
1003
1004         """TRIAD: return R (3x3) mapping camera -> inertial frame."""
1005
1006         # Camera triad
1007
1008         v1 = v_cam / np.linalg.norm(v_cam)
1009
1010         w1 = w_cam / np.linalg.norm(w_cam)
1011
1012         t1_cam = v1
1013
1014         t2_cam = np.cross(v1, w1)

```

```
980     n2 = np.linalg.norm(t2_cam)
981
982     if n2 < 1e-6:
983
984         raise ValueError("TRIAD: camera vectors are nearly colinear
985                         .")
986
987     t2_cam /= n2
988
989     t3_cam = np.cross(t1_cam, t2_cam)
990
991     C_cam = np.column_stack((t1_cam, t2_cam, t3_cam))
992
993
994     # Inertial triad
995
996     v2 = v_inertial / np.linalg.norm(v_inertial)
997
998     w2 = w_inertial / np.linalg.norm(w_inertial)
999
1000    t1_in = v2
1001
1002    t2_in = np.cross(v2, w2)
1003
1004    n2 = np.linalg.norm(t2_in)
1005
1006    if n2 < 1e-6:
1007
1008        raise ValueError("TRIAD: inertial vectors are nearly
1009                         colinear.")
1010
1011    t2_in /= n2
1012
1013    t3_in = np.cross(t1_in, t2_in)
1014
1015    C_in = np.column_stack((t1_in, t2_in, t3_in))
1016
1017
1018    # R: camera -> inertiaal
1019
1020    R = C_in @ C_cam.T
1021
1022    return R
1023
1024
1025    # -----
1026
1027    # Wahba via Davenport Q
1028
1029    # -----
1030
1031
1032
1033    def wahba_davenport_q(camera_vecs: np.ndarray,
1034                           inertial_vecs: np.ndarray,
1035                           weights: Optional[np.ndarray] = None) -> np.
1036                           ndarray:
```

```
1011     """
1012     Solve Wahba's problem using Davenport's Q method.
1013
1014     camera_vecs: (N,3)
1015     inertial_vecs: (N,3)
1016     """
1017
1018     if weights is None:
1019         weights = np.ones(camera_vecs.shape[0], dtype=np.float64)
1020
1021     W = np.diag(weights)
1022
1023     B = camera_vecs.T @ W @ inertial_vecs
1024
1025     S = B + B.T
1026
1027     sigma = np.trace(B)
1028
1029     Z = np.array([
1030         B[1, 2] - B[2, 1],
1031         B[2, 0] - B[0, 2],
1032         B[0, 1] - B[1, 0],
1033     ], dtype=np.float64)
1034
1035
1036     K = np.zeros((4, 4), dtype=np.float64)
1037     K[:3, :3] = S - sigma * np.eye(3)
1038     K[:3, 3] = Z
1039     K[3, :3] = Z
1040     K[3, 3] = sigma
1041
1042     eigvals, eigvecs = np.linalg.eigh(K)
1043     q = eigvecs[:, np.argmax(eigvals)] # largest eigenvalue
1044     q_vec = q[:3]
1045     q0 = q[3]
1046
1047     # Normalize quaternion
1048     qx, qy, qz = q_vec
1049     norm_q = math.sqrt(q0*q0 + qx*qx + qy*qy + qz*qz)
1050     if norm_q == 0:
1051         raise RuntimeError("Zero quaternion in Wahba solution.")
```

```

1045     q0 /= norm_q
1046     qx /= norm_q
1047     qy /= norm_q
1048     qz /= norm_q
1049
1050     # Quaternion -> rotatiematrix
1051     R = np.array([
1052         [1 - 2*(qy**2 + qz**2),      2*(qx*qy - qz*q0),      2*(
1053             qx*qz + qy*q0)],
1054         [2*(qx*qy + qz*q0),      1 - 2*(qx**2 + qz**2),      2*(
1055             qy*qz - qx*q0)],
1056         [2*(qx*qz - qy*q0),      2*(qy*qz + qx*q0),      1 -
1057             2*(qx**2 + qy**2)],
1058     ], dtype=np.float64)
1059
1060     return R
1061
1062
1063
1064 def attitude_from_star_pairs_ransac(
1065     detections: List[DetectedStar],
1066     catalog: List[CatalogStar],
1067     pair_table,
1068     angle_bin_rad: float,
1069     max_iterations: int = 500,
1070     max_inlier_err_deg: float = 0.2,
1071     min_inliers: int = 8,
1072 ):
1073     """
1074     Estimate attitude with RANSAC:
1075     - compare angles between detected star pairs to catalog pairs

```

```
1076     - build hypotheses via TRIAD
1077     - score on inlier count
1078     - refine best solution with Davenport Q
1079     """
1080
1081     bins = pair_table
1082
1083     cam_vecs = np.array([d.ray_cam for d in detections])
1084
1085     if cam_vecs.shape[0] < 2:
1086
1087         raise RuntimeError("Need at least 2 detections for attitude
1088                           .")
1089
1090
1091     cat_vecs = np.array([c.vec_inertial for c in catalog])
1092     max_inlier_err_rad = math.radians(max_inlier_err_deg)
1093     cos_max_err = math.cos(max_inlier_err_rad)
1094
1095     best_inliers: List[Tuple[int, int]] = []
1096     best_R = None
1097
1098     # Precompute all detection pairs
1099     M = len(detections)
1100     det_pairs = []
1101
1102     for i in range(M - 1):
1103
1104         for j in range(i + 1, M):
1105
1106             vi = cam_vecs[i]
1107
1108             vj = cam_vecs[j]
1109
1110             cosang = float(np.clip(np.dot(vi, vj), -1.0, 1.0))
1111
1112             ang = math.acos(cosang)
1113
1114             det_pairs.append((i, j, ang))
1115
1116     if not det_pairs:
1117
1118         raise RuntimeError("No detection pairs found.")
1119
1120
1121     import random
1122
1123     for _ in range(max_iterations):
1124
1125         # Pick a random detection pair
```

```
1109     i_det, j_det, ang_det = random.choice(det_pairs)
1110     bin_idx = int(ang_det / angle_bin_rad)
1111     candidate_pairs = []
1112     for b in (bin_idx - 1, bin_idx, bin_idx + 1):
1113         if b in bins:
1114             candidate_pairs.extend(bins[b])
1115     if not candidate_pairs:
1116         continue
1117     i_cat, j_cat, _ = random.choice(candidate_pairs)
1118
1119     # Try both mappings (i->i_cat, j->j_cat) and swapped
1120     for mapping in [((i_det, i_cat), (j_det, j_cat)),
1121                     ((i_det, j_cat), (j_det, i_cat))]:
1122         (i1, c1), (i2, c2) = mapping
1123         v_cam1 = cam_vecs[i1]
1124         v_cam2 = cam_vecs[i2]
1125         v_cat1 = cat_vecs[c1]
1126         v_cat2 = cat_vecs[c2]
1127         try:
1128             R = triad_attitude(v_cam1, v_cam2, v_cat1, v_cat2)
1129         except ValueError:
1130             continue
1131
1132         # Score hypothesis
1133         inliers: List[Tuple[int, int]] = []
1134         for k, v_cam in enumerate(cam_vecs):
1135             v_in = R @ v_cam
1136             dots = cat_vecs @ v_in
1137             idx = int(np.argmax(dots))
1138             if dots[idx] >= cos_max_err:
1139                 inliers.append((k, idx))
1140
1141         if len(inliers) > len(best_inliers):
1142             best_inliers = inliers
```

```
1143         best_R = R
1144
1145     if best_R is None or len(best_inliers) < min_inliers:
1146         raise RuntimeError("RANSAC could not find a consistent
1147                             attitude.")
1148
1149     # Refine with Wahba on all inliers
1150     cam_list = np.array([detections[i].ray_cam for (i, _) in
1151                         best_inliers])
1152     in_list = np.array([catalog[j].vec_inertial for (_, j) in
1153                         best_inliers])
1154     R_refined = wahba_davenport_q(cam_list, in_list)
1155
1156     return R_refined, best_inliers
1157
1158
1159 def debug_residuals(R: np.ndarray,
1160                      detections: List[DetectedStar],
1161                      catalog: List[CatalogStar],
1162                      inliers: List[Tuple[int, int]]) -> None:
1163     """Print median/max angular residuals in degrees for matched
1164     stars."""
1165     import numpy as np    # local to keep global imports light
1166     errs = []
1167     for det_idx, cat_idx in inliers:
1168         v_cam = detections[det_idx].ray_cam
1169         v_cat = catalog[cat_idx].vec_inertial
1170         v_pred = R @ v_cam
1171         cosang = float(np.clip(np.dot(v_pred, v_cat), -1.0, 1.0))
1172         err = math.degrees(math.acos(cosang))
1173         errs.append(err)
1174
1175     if not errs:
1176         print("Residuals: none (no inliers).")
1177
1178     return
```

```
1173     print(f"Residuals: median={np.median(errs):.3f} deg, max={max(
1174         errs):.3f} deg, N={len(errs)})")
1175
1176 def match_all_detections(R: np.ndarray,
1177                           detections: List[DetectedStar],
1178                           catalog: List[CatalogStar],
1179                           max_err_deg: float = 0.3) -> List[Tuple[
1180                               int, int]]:
1181     """Match remaining detections to nearest catalog star within
1182     angular threshold."""
1183     import numpy as np
1184     if not detections:
1185         return []
1186     cat_vecs = np.array([c.vec_inertial for c in catalog])
1187     max_err_rad = math.radians(max_err_deg)
1188     cos_min = math.cos(max_err_rad)
1189
1190     extra_matches = []
1191     for i, det in enumerate(detections):
1192         if det.catalog_idx is not None:
1193             continue
1194         v_cam = det.ray_cam
1195         v_in = R @ v_cam
1196         dots = cat_vecs @ v_in
1197         idx = int(np.argmax(dots))
1198         if dots[idx] >= cos_min:
1199             det.catalog_idx = idx
1200             extra_matches.append((i, idx))
1201
1202     # -----
1203     # Boresight RA/Dec
```

```
1204 # -----
1205
1206 def rotation_to_boresight_radec(R: np.ndarray) -> Tuple[float,
1207     float]:
1208     """Return RA/Dec of the camera z-axis (boresight)."""
1209     boresight_cam = np.array([0.0, 0.0, 1.0], dtype=np.float64)
1210     v = R @ boresight_cam
1211     x, y, z = v
1212     dec = math.degrees(math.asin(z))
1213     ra = math.degrees(math.atan2(y, x))
1214     if ra < 0:
1215         ra += 360.0
1216     return ra, dec
1217
1218 # -----
1219 # Image annotation
1220 # -----
1221
1222 def annotate_image(
1223     image_bgr: np.ndarray,
1224     detections: List[DetectedStar],
1225     catalog: List[CatalogStar],
1226     matches: List[Tuple[int, int]],
1227     boresight_radec: Tuple[float, float],
1228 ) -> np.ndarray:
1229     """Draw circles and labels for matched stars."""
1230     out = image_bgr.copy()
1231     for det_idx, cat_idx in matches:
1232         det = detections[det_idx]
1233         star = catalog[cat_idx]
1234         center = (int(round(det.x)), int(round(det.y)))
1235         cv2.circle(out, center, 5, (0, 255, 0), 1, lineType=cv2.
1236 LINE_AA)
```

```
1236     label = star.name
1237     cv2.putText(out,
1238         label,
1239         (center[0] + 6, center[1] - 3),
1240         cv2.FONT_HERSHEY_SIMPLEX,
1241         0.35,
1242         (0, 255, 0),
1243         1,
1244         lineType=cv2.LINE_AA)
1245
1246     ra_deg, dec_deg = boresight_radec
1247     txt = f"Pointing: RA={ra_deg:7.3f} deg Dec={dec_deg:6.3f} deg"
1248     cv2.putText(out,
1249         txt,
1250         (20, 30),
1251         cv2.FONT_HERSHEY_SIMPLEX,
1252         0.6,
1253         (0, 255, 255),
1254         2,
1255         lineType=cv2.LINE_AA)
1256
1257
1258
1259 def annotate_detections_only(
1260     image_bgr: np.ndarray,
1261     detections: List[DetectedStar],
1262     header: str = "ATTITUDE_FAILED",
1263 ) -> np.ndarray:
1264     """Overlay raw detections (no catalog matches) so failures are
1265     inspectable."""
1266     out = image_bgr.copy()
1267     for i, det in enumerate(detections):
1268         center = (int(round(det.x)), int(round(det.y)))
```

```
1268     cv2.circle(out, center, 5, (0, 255, 0), 1, lineType=cv2.  
1269         LINE_AA)  
1270     cv2.putText(  
1271         out,  
1272         f"{i+1}",  
1273         (center[0] + 6, center[1] - 3),  
1274         cv2.FONT_HERSHEY_SIMPLEX,  
1275         0.35,  
1276         (0, 255, 0),  
1277         1,  
1278         lineType=cv2.LINE_AA,  
1279     )  
1280     if header:  
1281         cv2.putText(  
1282             out,  
1283             header,  
1284             (20, 30),  
1285             cv2.FONT_HERSHEY_SIMPLEX,  
1286             0.6,  
1287             (0, 0, 255),  
1288             2,  
1289             lineType=cv2.LINE_AA,  
1290         )  
1291     return out  
1292  
1293 # -----  
1294 # High-level pipeline  
1295 # -----  
1296  
1297 def run_star_tracker(  
1298     image_path: str,  
1299     catalog_path: str,  
1300     max_catalog_mag: float = 8.0,
```

```
1301     output_path: str = "annotated.png",
1302     processing_options: Optional[ProcessingOptions] = None,
1303     image_config: Optional[ImageInputConfig] = None,
1304     camera_params_override: Optional[Dict[str, float]] = None,
1305     denoise_config: Optional[DenoiseConfig] = None,
1306     detection_mask: Optional[np.ndarray] = None,
1307 ):
1308     """
1309     Full pipeline: image -> attitude -> annotation.
1310     """
1311     img_stem = Path(image_path).stem
1312     debug_dir = STAR_DEBUG_DIR / img_stem if
1313         STAR_SAVE_DEBUG_INTERMEDIATES else None
1314
1315     # Load image + optional resize/rotate
1316     cfg = image_config or ImageInputConfig(path=image_path)
1317     img_bgr = cv2.imread(cfg.path, cv2.IMREAD_COLOR)
1318     if img_bgr is None:
1319         raise RuntimeError(f"Could not read image {cfg.path}")
1320     if STAR_FLIP_CODE is not None:
1321         img_bgr = cv2.flip(img_bgr, STAR_FLIP_CODE)
1322         print(f"DEBUG: Input image flipped code={STAR_FLIP_CODE}")
1323     if cfg.resize_to:
1324         w, h = cfg.resize_to
1325         img_bgr = cv2.resize(img_bgr, (w, h), interpolation=cv2.
1326             INTER_LINEAR)
1327         print(f"DEBUG: Input image resized to {w}x{h}")
1328     if abs(cfg.rotate_deg) > 1e-6:
1329         img_bgr = _rotate_image_keep_size(img_bgr, cfg.rotate_deg)
1330         print(f"DEBUG: Input image rotated {cfg.rotate_deg:.2f} deg
1331             CCW")
1332     img_gray = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2GRAY)
1333     if denoise_config is not None:
1334         img_gray = clean_image(img_gray, denoise_config)
```

```
1332     if detection_mask is not None:
1333         if detection_mask.shape != img_gray.shape:
1334             raise ValueError("Detection mask shape must match image
1335 .")
1336         detection_mask = detection_mask.astype(bool)
1337     else:
1338         detection_mask = None
1339
1340     # Catalog (Yale ASCII or Gaia CSV)
1341     if catalog_path.lower().endswith(".csv"):
1342         catalog = load_gaia_csv(catalog_path, max_mag=
1343             max_catalog_mag)
1344     else:
1345         catalog = load_yale_bsc_ascii(catalog_path, max_mag=
1346             max_catalog_mag)
1347
1348     # Name replacement: prefer IAU names (via HIP), otherwise Gaia
1349     # IDs
1350     gaia_to_hip = load_gaia_to_hip_map(STAR_NAME_LOOKUP_PATH) if
1351         STAR_NAME_LOOKUP_PATH.exists() else {}
1352     iau_name_map = load_iau_names(STAR_IAU_PATH)
1353     apply_name_lookup(catalog, gaia_to_hip_map=gaia_to_hip,
1354         iau_name_map=iau_name_map)
1355
1356     if processing_options is None:
1357         processing_options = ProcessingOptions()
1358         print(f"Processing options: {processing_options}")
1359
1360     # Star detection
1361     debug_images = {} if STAR_SAVE_DEBUG_INTERMEDIATES else None
1362     detections = detect_stars(img_gray, options=processing_options,
1363         mask=detection_mask, debug_images=debug_images)
1364     if debug_images is not None:
```

```
1358     _write_debug_images(f"{img_stem}_detect_main", debug_images
1359         , out_dir=debug_dir)
1360
1361     # Retry with more sensitive preset if detections are too few
1362     if len(detections) < STAR_ATTITUDE_MIN_INLIERS:
1363         print(f"INFO: only {len(detections)} detections; retrying with a more sensitive preset")
1364     fallback_opts = ProcessingOptions(
1365         background_subtraction=True,
1366         background_kernel=max(processing_options.
1367             background_kernel, 71),
1368         normalize=True,
1369         threshold_multiplier=1.25,
1370         morph_open=False,
1371         morph_kernel=processing_options.morph_kernel,
1372         morph_iterations=processing_options.morph_iterations,
1373         min_area=3,
1374         max_area=processing_options.max_area,
1375         max_stars=400,
1376         saturated_threshold=min(processing_options.
1377             saturated_threshold, 215),
1378         allow_oversized_saturated=True,
1379     )
1380     debug_images_fb = {} if STAR_SAVE_DEBUG_INTERMEDIATES else
1381         None
1382     detections = detect_stars(img_gray, options=fallback_opts,
1383         mask=detection_mask, debug_images=debug_images_fb)
1384     if debug_images_fb is not None:
1385         _write_debug_images(f"{img_stem}_detect_fallback",
1386             debug_images_fb, out_dir=debug_dir)
1387     print(f"Detected {len(detections)} star candidates after
1388         fallback")
1389     if len(detections) < 2:
```

```
1384     print("WARN: Insufficient detections (<2) even after "
1385           "fallback; writing detections-only overlay.")
1386     annotated_fail = annotate_detections_only(img_bgr,
1387         detections, header="ATTITUDE FAILED: too few detections")
1388
1389     cv2.imwrite(output_path, annotated_fail)
1390
1391     return
1392
1393 # Pixels -> 3D rays
1394 pixels_to_unit_vectors(detections, img_gray.shape,
1395                         hardware_params=camera_params_override)
1396
1397 # Catalog pairs
1398 pair_table, angle_bin_rad = build_catalog_pair_table(catalog)
1399
1400
1401 # Attitude via RANSAC + TRIAD + Davenport
1402 min_inliers = max(2, min(STAR_ATTITUDE_MIN_INLIERS, len(
1403     detections)))
1404
1405 try:
1406     R, inliers = attitude_from_star_pairs_ransac(
1407         detections, catalog, pair_table, angle_bin_rad,
1408         max_iterations=STAR_ATTITUDE_MAX_ITER,
1409         max_inlier_err_deg=STAR_ATTITUDE_MAX_ERR_DEG,
1410         min_inliers=min_inliers,
1411     )
1412
1413 except Exception as exc:
1414     print(f"ATTITUDE FAILED: {exc}")
1415     annotated_fail = annotate_detections_only(img_bgr,
1416         detections, header="ATTITUDE FAILED")
1417
1418     cv2.imwrite(output_path, annotated_fail)
1419
1420     return
1421
1422     print(f"Found attitude with {len(inliers)} inliers")
1423     debug_residuals(R, detections, catalog, inliers)
```

```
1411     extra_matches = match_all_detections(R, detections, catalog,
1412                                         max_err_deg=max(STAR_ATTITUDE_MAX_ERR_DEG, 0.3))
1413
1414     all_matches = inliers + extra_matches
1415
1416     # Attach matches to detections
1417
1418     for det_idx, cat_idx in all_matches:
1419         detections[det_idx].catalog_idx = cat_idx
1420
1421     # Boresight RA/Dec
1422
1423     ra_deg, dec_deg = rotation_to_boresight_radec(R)
1424
1425     print(f"Boresight pointing: RA={ra_deg:.3f} deg Dec={dec_deg:.3f} deg")
1426
1427     # Annotation
1428
1429     annotated = annotate_image(img_bgr, detections, catalog,
1430                                all_matches,
1431                                (ra_deg, dec_deg))
1432
1433     cv2.imwrite(output_path, annotated)
1434
1435     print(f"Annotated image written to {output_path}")
1436
1437
1438
1439 def main() -> None:
1440
1441     processing_opts = STAR_PROCESSING_OPTIONS or processing_preset(
1442         STAR_PROCESSING_MODE)
1443
1444     base_path = STAR_IMAGE_PATH
1445
1446     image_paths: List[Path] = []
1447
1448     if base_path.is_dir():
1449
1450         image_paths = sorted(p for p in base_path.glob("IMD_*.png"))
1451
1452         if p.is_file():
1453
1454             elif base_path.exists():
1455
1456                 image_paths = [base_path]
1457
1458             elif STAR_IMAGE_FALLBACK.exists():
1459
1460                 image_paths = [STAR_IMAGE_FALLBACK]
1461
1462             if not image_paths:
```

```
1440     raise RuntimeError(
1441         f"No input images found. Tried directory/file {STAR_IMAGE_PATH} and fallback {STAR_IMAGE_FALLBACK}"
1442     )
1443
1444     print(f"Processing {len(image_paths)} image{s}")
1445     for image_path in image_paths:
1446         print(f"-> {image_path}")
1447         image_cfg = ImageInputConfig(
1448             path=str(image_path),
1449             rotate_deg=STAR_ROTATE_DEG,
1450             resize_to=STAR_RESIZE_TO,
1451         )
1452
1453     # Build mask that keeps the overlap region and removes
1454     # borders
1455
1456     mask = _load_precomputed_mask(image_path)
1457     gray_for_mask = None
1458
1459     if mask is None:
1460
1461         if STAR_AUTO_OVERLAP_MASK:
1462
1463             gray_for_mask = cv2.imread(str(image_path), cv2.IMREAD_GRAYSCALE)
1464
1465             if gray_for_mask is None:
1466
1467                 raise RuntimeError(f"Could not read image for"
1468                     f"mask: {image_path}")
1469
1470             mask = _quadrilateral_mask_from_image(
1471
1472                 gray_for_mask,
1473                 thresh_percentile=STAR_AUTO_MASK_PERCENTILE,
1474                 close_kernel=STAR_AUTO_MASK_CLOSE_KERNEL,
1475                 erode_px=STAR_AUTO_MASK_ERODE,
1476             )
1477
1478             if STAR_SAVE_DEBUG_INTERMEDIATES:
1479
1480                 _write_debug_images("mask_build", {"auto_mask": mask,
1481                     "input_for_mask": gray_for_mask})
1482
1483         elif STAR_MASK_BORDERS_PX:
```

```
1469         top, bottom, left, right = STAR_MASK_BORDERS_PX
1470
1471         gray_for_mask = cv2.imread(str(image_path), cv2.
1472                                     IMREAD_GRAYSCALE)
1473
1474         if gray_for_mask is None:
1475             raise RuntimeError(f"Could not read image for
1476                                 mask: {image_path}")
1477
1478         H, W = gray_for_mask.shape
1479
1480         mask = np.zeros((H, W), dtype=np.uint8)
1481
1482         mask[top : H - bottom, left : W - right] = 1
1483
1484
1485         # Optional Jupiter/glare mask; combine and save for preview
1486
1487         if STAR_SAVE_JUPITER_MASK:
1488
1489             if gray_for_mask is None:
1490
1491                 gray_for_mask = cv2.imread(str(image_path), cv2.
1492                                             IMREAD_GRAYSCALE)
1493
1494             if gray_for_mask is not None:
1495
1496                 glare_mask = _compute_jupiter_mask(gray_for_mask)
1497
1498                 # Save glare mask for preview (even if None -> save
1499
1500                     all zeros)
1501
1502                 STAR_DEBUG_DIR.mkdir(parents=True, exist_ok=True)
1503
1504                 glare_path = STAR_DEBUG_DIR / f"{image_path.stem}
1505
1506                     _jupiter_mask.png"
1507
1508             if glare_mask is None:
1509
1510                 blank = np.zeros_like(gray_for_mask, dtype=np.
1511
1512                     uint8)
1513
1514                 cv2.imwrite(str(glare_path), blank)
1515
1516             else:
1517
1518                 cv2.imwrite(str(glare_path), glare_mask * 255)
1519
1520                 # Apply to detection mask (AND with inverse)
1521
1522                 if mask is None:
1523
1524                     mask = np.ones_like(glare_mask, dtype=np.
1525
1526                     uint8)
1527
1528                 mask = (mask.astype(np.uint8) & (1 - glare_mask
1529
1530                     ).astype(np.uint8)).astype(np.uint8)
```

```
1495     if STAR_SAVE_DEBUG_INTERMEDIATES:
1496
1497         overlay = _make_overlay(glare_mask * 255,
1498                                  _to_uint8_debug(gray_for_mask))
1499
1500         cv2.imwrite(str(STAR_DEBUG_DIR / f"{image_path.stem}_jupiter_mask_overlay.
1501                         png"), overlay)
1502
1503
1504     output_path = image_path.with_name(f"{image_path.stem}
1505                                         _annotated.png")
1506
1507     run_star_tracker(
1508
1509         image_path=str(image_path),
1510         catalog_path=str(STAR_CATALOG_PATH),
1511         max_catalog_mag=STAR_MAX_CATALOG_MAG,
1512         output_path=str(output_path),
1513         processing_options=processing_opts,
1514         image_config=image_cfg,
1515         camera_params_override=STAR_CAMERA_OVERRIDE,
1516         denoise_config=STAR_DENOISE_SETTINGS if APPLY_DENOISING
1517
1518             else None,
1519             detection_mask=mask,
1520         )
1521
1522
1523 if __name__ == "__main__":
1524     main()
```

Listing A.2: Star tracker script